

CONNECT TO GOOGLE DRIVE

```
In [ ]: from google.colab import drive  
drive.mount('/content/gdrive')  
%cd /content
```

Mounted at /content/gdrive
/content

IMPORT DATA FROM KAGGLE

```
In [ ]: !mkdir ~/.kaggle  
!touch ~/.kaggle/kaggle.json  
token = {"username": "<YOUR USERNAME>", "key": "<YOUR KEY>"}  
import json  
with open('/root/.kaggle/kaggle.json', 'w') as file:  
    json.dump(token, file)  
!chmod 600 ~/.kaggle/kaggle.json
```

UNZIP IMPORTED DATA

```
In [ ]: !kaggle datasets download -d tarunbisht11/yolo-animal-detection-small  
!unzip yolo-animal-detection-small.zip
```

Downloading yolo-animal-detection-small.zip to /content
82% 25.0M/30.4M [00:01<00:00, 15.6MB/s]
100% 30.4M/30.4M [00:01<00:00, 21.5MB/s]
Archive: yolo-animal-detection-small.zip
 inflating: test.csv
 inflating: test.record
 inflating: train.csv
 inflating: train.record
 inflating: yolo-animal-detection-small/test/cats_000.jpg
 inflating: yolo-animal-detection-small/test/cats_000.xml
 inflating: yolo-animal-detection-small/test/cats_007.jpg
 inflating: yolo-animal-detection-small/test/cats_007.xml
 inflating: yolo-animal-detection-small/test/cats_050.jpg
 inflating: yolo-animal-detection-small/test/cats_050.xml
 inflating: yolo-animal-detection-small/test/cats_072.jpg
 inflating: yolo-animal-detection-small/test/cats_072.xml
 inflating: yolo-animal-detection-small/test/cats_076.jpg
 inflating: yolo-animal-detection-small/test/cats_076.xml
 inflating: yolo-animal-detection-small/test/cats_079.jpg
 inflating: yolo-animal-detection-small/test/cats_079.xml
 inflating: yolo-animal-detection-small/test/cats_088.jpg
 inflating: yolo-animal-detection-small/test/cats_088.xml
 inflating: yolo-animal-detection-small/test/cats_091.jpg
 inflating: yolo-animal-detection-small/test/cats_091.xml
 inflating: yolo-animal-detection-small/test/cats_099.jpg
 inflating: yolo-animal-detection-small/test/cats_099.xml
 inflating: yolo-animal-detection-small/test/cats_100.jpg
 inflating: yolo-animal-detection-small/test/cats_100.xml
 inflating: yolo-animal-detection-small/test/cats_and_monkeys_000.jpg
 inflating: yolo-animal-detection-small/test/cats_and_monkeys_000.xml
 inflating: yolo-animal-detection-small/test/cats_and_monkeys_002.jpg
 inflating: yolo-animal-detection-small/test/cats_and_monkeys_002.xml
 inflating: yolo-animal-detection-small/test/cats_and_monkeys_029.jpg


```
inflating: yolo-animal-detection-small/train/monkeys_084.xml
inflating: yolo-animal-detection-small/train/monkeys_086.jpg
inflating: yolo-animal-detection-small/train/monkeys_086.xml
inflating: yolo-animal-detection-small/train/monkeys_087.jpg
inflating: yolo-animal-detection-small/train/monkeys_087.xml
inflating: yolo-animal-detection-small/train/monkeys_089.jpg
inflating: yolo-animal-detection-small/train/monkeys_089.xml
inflating: yolo-animal-detection-small/train/monkeys_090.jpg
inflating: yolo-animal-detection-small/train/monkeys_090.xml
inflating: yolo-animal-detection-small/train/monkeys_091.jpg
inflating: yolo-animal-detection-small/train/monkeys_091.xml
inflating: yolo-animal-detection-small/train/monkeys_092.jpg
inflating: yolo-animal-detection-small/train/monkeys_092.xml
inflating: yolo-animal-detection-small/train/monkeys_094.jpg
inflating: yolo-animal-detection-small/train/monkeys_094.xml
inflating: yolo-animal-detection-small/train/monkeys_095.jpg
inflating: yolo-animal-detection-small/train/monkeys_095.xml
inflating: yolo-animal-detection-small/train/monkeys_097.jpg
inflating: yolo-animal-detection-small/train/monkeys_097.xml
inflating: yolo-animal-detection-small/train/monkeys_098.jpg
inflating: yolo-animal-detection-small/train/monkeys_098.xml
inflating: yolo-animal-detection-small/train/monkeys_099.jpg
inflating: yolo-animal-detection-small/train/monkeys_099.xml
inflating: yolo-animal-detection-small/train/monkeys_100.jpg
inflating: yolo-animal-detection-small/train/monkeys_100.xml
```

REMOVE UNNECESSARY FILES

```
In [ ]: !rm -r /content/sample_data
!rm /content/yolo-animal-detection-small.zip
!rm /content/train.record
!rm /content/test.record
```

IMPORT LIBRARIES

```
In [ ]: import os
import math
import imagesize
import cv2 as cv
import numpy as np
import pandas as pd
import random as rn
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
```

READ TRAINING AND TESTING CSV FILE

```
In [ ]: train_df = pd.read_csv('/content/train.csv')
train_df = pd.DataFrame(train_df)
test_df = pd.read_csv('/content/test.csv')
test_df = pd.DataFrame(test_df)
train_df.head()
```

```
Out[ ]:    filename  width  height  class  xmin  ymin  xmax  ymax
0  cats_001.jpg     474     266   cat     132      1     347     264
```

1	cats_002.jpg	474	474	cat	176	44	467	433
2	cats_003.jpg	474	314	cat	53	1	397	314
3	cats_004.jpg	474	355	cat	1	1	393	335
4	cats_005.jpg	474	316	cat	80	1	407	316

CONVERT FILENAME TO FULLPATHS

```
In [ ]: train_df['filename'] = '/content/yolo-animal-detection-small/train/' + train_df['filename']
test_df['filename'] = '/content/yolo-animal-detection-small/test/' + test_df['filename']
train_df.head()
```

	filename	width	height	class	xmin	ymin	xmax	ymax
0	/content/yolo-animal-detection-small/train/cat...	474	266	cat	132	1	347	264
1	/content/yolo-animal-detection-small/train/cat...	474	474	cat	176	44	467	433
2	/content/yolo-animal-detection-small/train/cat...	474	314	cat	53	1	397	314
3	/content/yolo-animal-detection-small/train/cat...	474	355	cat	1	1	393	335
4	/content/yolo-animal-detection-small/train/cat...	474	316	cat	80	1	407	316

FILTER BOUNDING BOX ON TRAINING DATAFRAME WITH HIGHEST AREA

```
In [ ]: train_df = train_df.sample(frac=1)
train_df.reset_index(drop=True, inplace=True)
filenames = []
cats = []
xmins = []
ymins = []
xmaxs = []
ymaxs = []
for i in range(train_df.shape[0]):
    img_path = train_df['filename'][i]
    bboxes = train_df[train_df['filename'] == img_path]
    bboxes.reset_index(drop=True, inplace=True)
    filenames.append(img_path)
    max_area = 0
    index = 0
    for j in range(bboxes.shape[0]):
        xmin = bboxes['xmin'][j]
        ymin = bboxes['ymin'][j]
        xmax = bboxes['xmax'][j]
        ymax = bboxes['ymax'][j]
        area = (xmax - xmin) * (ymax - ymin)
        if area > max_area:
            max_area = area
            index = j
    cats.append(bboxes['class'][index])
    xmins.append(bboxes['xmin'][index])
    ymins.append(bboxes['ymin'][index])
    xmaxs.append(bboxes['xmax'][index])
    ymaxs.append(bboxes['ymax'][index])
train_df = pd.DataFrame({'filename': filenames, 'class': cats, 'xmin': xmins, 'ymin': ymins, 'xmax': xmaxs, 'ymax': ymaxs})
```

```

train_df = train_df.drop_duplicates(subset='filename', keep='first', ignore_index=True)
train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 469 entries, 0 to 468
Data columns (total 6 columns):
 #   Column   Non-Null Count Dtype  
---  --  
 0   filename  469 non-null   object 
 1   class     469 non-null   object 
 2   xmin      469 non-null   int64  
 3   ymin      469 non-null   int64  
 4   xmax      469 non-null   int64  
 5   ymax      469 non-null   int64  
dtypes: int64(4), object(2)
memory usage: 22.1+ KB

```

FILTER BOUNDING BOX ON TESTING DATFRAME WITH HIGHEST AREA

```

In [ ]: test_df = test_df.sample(frac=1)
test_df.reset_index(drop=True, inplace=True)
filenames = []
cats = []
xmins = []
ymins = []
xmaxs = []
ymaxs = []
for i in range(test_df.shape[0]):
    img_path = test_df['filename'][i]
    bboxes = test_df[test_df['filename'] == img_path]
    bboxes.reset_index(drop=True, inplace=True)
    filenames.append(img_path)
    max_area = 0
    index = 0
    for j in range(bboxes.shape[0]):
        xmin = bboxes['xmin'][j]
        ymin = bboxes['ymin'][j]
        xmax = bboxes['xmax'][j]
        ymax = bboxes['ymax'][j]
        area = (xmax - xmin) * (ymax - ymin)
        if area > max_area:
            max_area = area
            index = j
    cats.append(bboxes['class'][index])
    xmins.append(bboxes['xmin'][index])
    ymins.append(bboxes['ymin'][index])
    xmaxs.append(bboxes['xmax'][index])
    ymaxs.append(bboxes['ymax'][index])
test_df = pd.DataFrame({'filename': filenames, 'class': cats, 'xmin': xmins, 'ymin': ymins, 'xmax': xmaxs, 'ymax': ymaxs})
test_df = test_df.drop_duplicates(subset='filename', keep='first', ignore_index=True)
test_df.info()

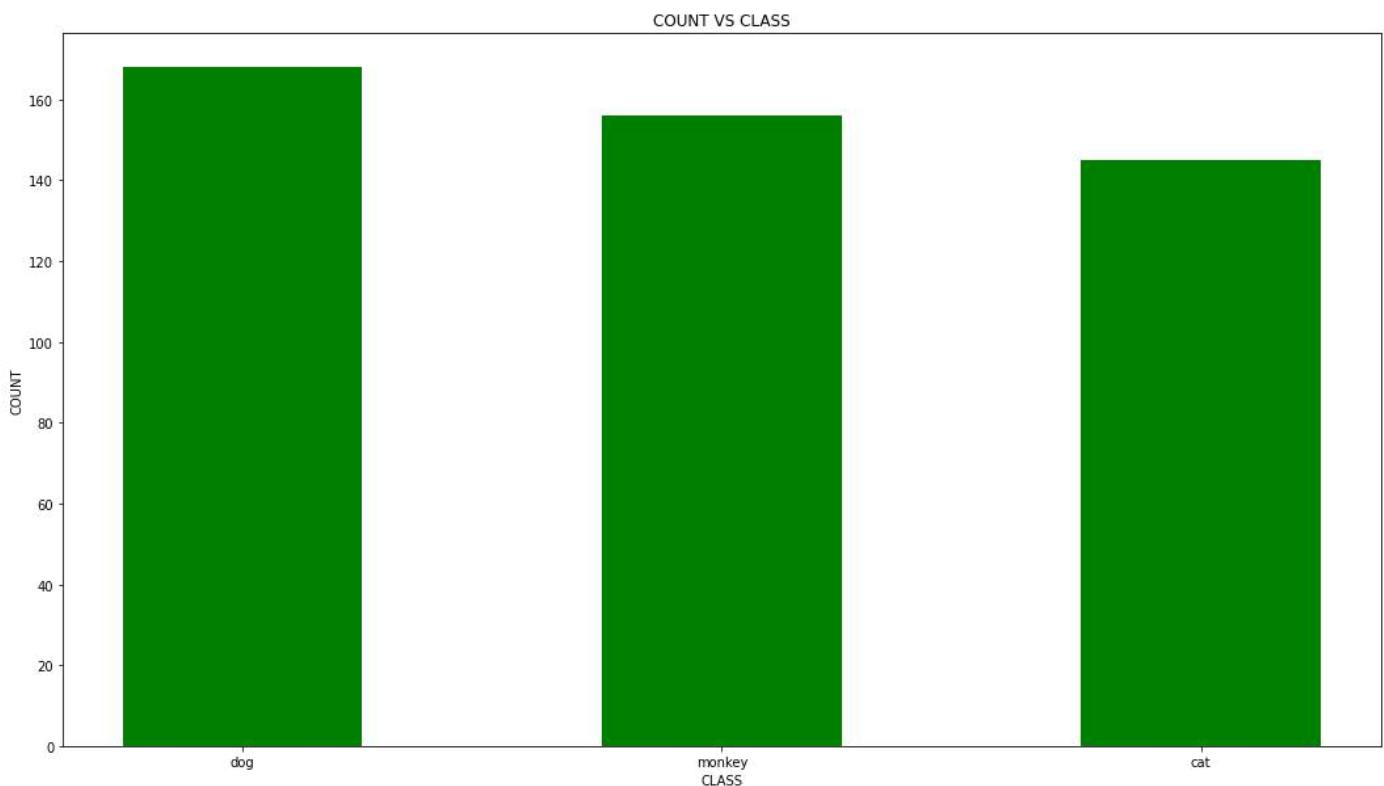
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 51 entries, 0 to 50
Data columns (total 6 columns):
 #   Column   Non-Null Count Dtype  
---  --  
 0   filename  51 non-null   object 
 1   class     51 non-null   object 
 2   xmin      51 non-null   int64  
 3   ymin      51 non-null   int64  
 4   xmax      51 non-null   int64  
 5   ymax      51 non-null   int64  

```

```
dtypes: int64(4), object(2)
memory usage: 2.5+ KB
```

PLOT COUNT VS CLASS IN TRAINING DATA

```
In [ ]: class_dict = dict(train_df['class'].value_counts())
cat = list(class_dict.keys())
count = list(class_dict.values())
plt.figure(figsize=(18, 10))
plt.bar(cat, count, color='green', width=0.5)
plt.xlabel('CLASS')
plt.ylabel('COUNT')
plt.title('COUNT VS CLASS')
plt.show()
```



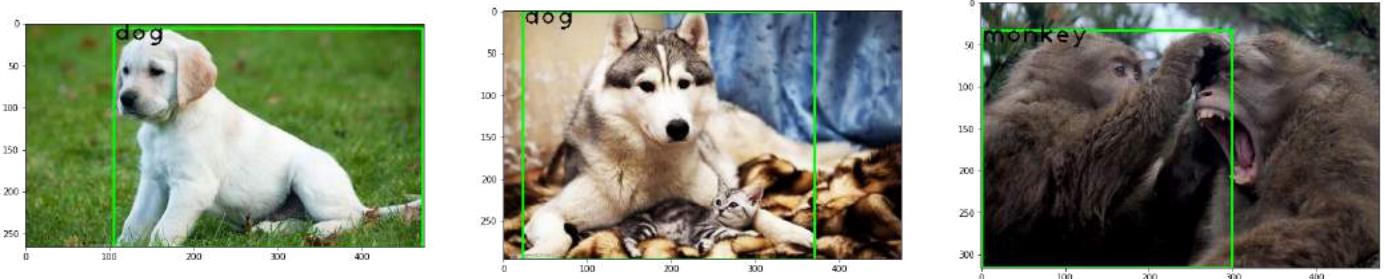
GET STATISTICAL INFORMATION FROM TRAINING DATA

```
In [ ]: train_df.describe()
```

```
Out[ ]:   xmin    ymin    xmax    ymax
count  469.000000  469.000000  469.000000  469.000000
mean   95.076759  44.582090  381.622601  338.820896
std    90.105888  51.569928  91.120848  119.823722
min    1.000000  1.000000  98.000000  100.000000
25%   18.000000  7.000000  336.000000  266.000000
50%   69.000000  27.000000  407.000000  311.000000
75%   150.000000  65.000000  459.000000  355.000000
max   396.000000  392.000000  474.000000  746.000000
```

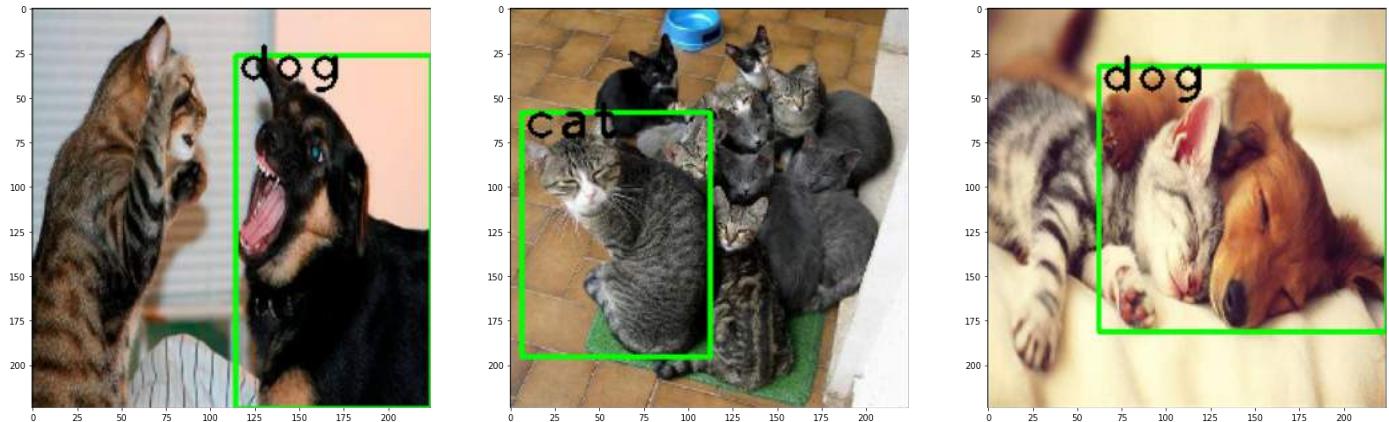
DISPLAY IMAGES AND BOUNDING BOXES

```
In [ ]: def plot_images_and_bboxes (dataframe):
    rows = 1
    cols = 3
    fig, ax = plt.subplots(rows, cols, figsize=(28, 13))
    for i in range(cols):
        index = rn.randint(0, dataframe.shape[0] - 1)
        img_path = dataframe['filename'][index]
        img = cv.imread(img_path)
        img = cv.cvtColor(img, cv.COLOR_BGR2RGB)
        cat = dataframe['class'][index]
        xmin = int(dataframe['xmin'][index])
        ymin = int(dataframe['ymin'][index])
        xmax = int(dataframe['xmax'][index])
        ymax = int(dataframe['ymax'][index])
        img = cv.rectangle(img, (xmin, ymin), (xmax, ymax), (0, 255, 0), 2)
        img = cv.putText(img, f'{cat}', (xmin, ymin + 15), cv.FONT_HERSHEY_PLAIN, 2, (0,
        ax[i].imshow(img)
    plt.show()
plot_images_and_bboxes (train_df)
```



DISPLAY RESIZED IMAGES AND CORRESPONDING BOUNDING BOXES

```
In [ ]: def plot_resized_images_and_bboxes (dataframe, img_h, img_w):
    rows = 1
    cols = 3
    fig, ax = plt.subplots(rows, cols, figsize=(28, 13))
    for i in range(cols):
        index = rn.randint(0, dataframe.shape[0] - 1)
        img_path = dataframe['filename'][index]
        img = cv.imread(img_path)
        img = cv.cvtColor(img, cv.COLOR_BGR2RGB)
        h, w, _ = img.shape
        img = cv.resize(img, (img_w, img_h), interpolation=cv.INTER_CUBIC)
        h_ratio = h / img_h
        w_ratio = w / img_w
        cat = dataframe['class'][index]
        xmin = int(dataframe['xmin'][index] / w_ratio)
        ymin = int(dataframe['ymin'][index] / h_ratio)
        xmax = int(dataframe['xmax'][index] / w_ratio)
        ymax = int(dataframe['ymax'][index] / h_ratio)
        img = cv.rectangle(img, (xmin, ymin), (xmax, ymax), (0, 255, 0), 2)
        img = cv.putText(img, f'{cat}', (xmin, ymin + 15), cv.FONT_HERSHEY_PLAIN, 2, (0,
        ax[i].imshow(img)
    plt.show()
H, W, C = 224, 224, 3
plot_resized_images_and_bboxes (train_df, H, W)
```



LABEL ENCODE CLASSES

```
In [ ]: encoder = LabelEncoder()
train_df['class'] = encoder.fit_transform(train_df['class'])
test_df['class'] = encoder.transform(test_df['class'])
encoder.classes_
Out[ ]: array(['cat', 'dog', 'monkey'], dtype=object)
```

SPLIT TESTING DATA INTO VALIDATION AND TESTING

```
In [ ]: val_df, test_df = train_test_split(test_df, test_size=0.5, random_state=24)
val_df.reset_index(drop=True, inplace=True)
test_df.reset_index(drop=True, inplace=True)
train_df.shape, val_df.shape, test_df.shape
Out[ ]: ((469, 6), (25, 6), (26, 6))
```

CREATE DATA PIPELINE

```
In [ ]: class Pipeline(tf.keras.utils.Sequence):
    def __init__(self, dataframe, xcol, ycol, batch_size, img_h, img_w, img_c, classes):
        super(Pipeline, self).__init__()
        self.dataframe = dataframe
        self.xcol = xcol
        self.ycol = ycol
        self.batch_size = batch_size
        self.img_h = img_h
        self.img_w = img_w
        self.img_c = img_c
        self.classes = classes

    def __len__(self):
        return self.dataframe.shape[0] // self.batch_size

    def on_epoch_end(self):
        self.dataframe = self.dataframe.sample(frac=1)
        self.dataframe.reset_index(drop=True, inplace=True)

    def preprocess_img(self, img_path):
        img = tf.io.read_file(img_path)
        img = tf.io.decode_image(img, channels=self.img_c, dtype=tf.float32)
        img = tf.image.resize(img, (self.img_h, self.img_w))
```

```

    img = tf.expand_dims(img, axis=0)
    img = img / 255.
    return img

def __getitem__(self, item):
    X = []
    Y = []
    for i in range(self.batch_size):
        img_path = self.dataframe[self.xcol][i + (self.batch_size * item)]
        img = self.preprocess_img(img_path)
        X.append(img)
        img = cv.imread(img_path)
        h, w, _ = img.shape
        h_ratio = h / self.img_h
        w_ratio = w / self.img_w
        cat = self.dataframe[self.ycol[0]][i + (self.batch_size * item)]
        cat = list(tf.keras.utils.to_categorical(cat, num_classes=len(self.classes)))
        xmin = self.dataframe[self.ycol[1]][i + (self.batch_size * item)] / (w_ratio)
        ymin = self.dataframe[self.ycol[2]][i + (self.batch_size * item)] / (h_ratio)
        xmax = self.dataframe[self.ycol[3]][i + (self.batch_size * item)] / (w_ratio)
        ymax = self.dataframe[self.ycol[4]][i + (self.batch_size * item)] / (h_ratio)
        p = 1.0
        cat.extend([p, xmin, ymin, xmax, ymax])
        cat = np.array(cat, dtype='float32')
        cat = cat.reshape(-1, 5 + len(self.classes))
        Y.append(cat)
    X = tf.concat(X, axis=0)
    Y = tf.concat(Y, axis=0)
    return tf.convert_to_tensor(X, dtype=tf.float32), tf.convert_to_tensor(Y, dtype=
B = 5
N = list(encoder.classes_)
train_gen = Pipeline(train_df, 'filename', ['class', 'xmin', 'ymin', 'xmax', 'ymax'], B,
val_gen = Pipeline(val_df, 'filename', ['class', 'xmin', 'ymin', 'xmax', 'ymax'], B, H,
test_gen = Pipeline(test_df, 'filename', ['class', 'xmin', 'ymin', 'xmax', 'ymax'], B, H
train_gen.__getitem__(0)[0].shape, train_gen.__getitem__(0)[1].shape, train_gen.__getite

```

Out[]: (TensorShape([5, 224, 224, 3]), TensorShape([5, 8]), 0.0, 1.0)

CREATE CUSTOM BLOCK FOR CUSTOM MODEL

```

In [ ]: class BGBlock(tf.keras.layers.Layer):
    def __init__(self, f, k, s, p, d, pool, stage):
        super(BhavyaBlock, self).__init__(trainable=True, name=f'BGBlock{stage}')
        self.f = f
        self.k = k
        self.s = s
        self.p = p
        self.d = d
        self.pool = pool
        self.stage = stage

    def build(self, input_shape):
        setattr(self, f'conv0_stage{self.stage}', tf.keras.layers.Conv2D(filters=self.f[
            kernel_size=self
            strides=self.s[0]
            padding=self.p[0]
            dilation_rate=se
            kernel_initializer
            kernel_regularizer
            activation='relu
            name=f'conv0_st
            input_shape=input_
        setattr(self, f'leaky0_stage{self.stage}', tf.keras.layers.LeakyReLU(alpha=0.1,
```

```

        setattr(self, f'batch0_stage{self.stage}', tf.keras.layers.BatchNormalization(name=f'batch0_stage{self.stage}')
for i in range(1, len(self.f), 1):
    setattr(self, f'conv{i}_stage{self.stage}', tf.keras.layers.Conv2D(filters=kernel_size,
                                                                     strides=strides,
                                                                     padding=padding,
                                                                     dilation_rate=dilation_rate,
                                                                     kernel_initializer=kernel_initializer,
                                                                     kernel_regularizer=kernel_regularizer,
                                                                     activation=activation,
                                                                     name=f'conv{i}_stage{self.stage}'))
# vars(self)[f'leaky{i}_stage{self.stage}'] = tf.keras.layers.LeakyReLU(name=f'leaky{i}_stage{self.stage}')
setattr(self, f'batch{i}_stage{self.stage}', tf.keras.layers.BatchNormalization(name=f'batch{i}_stage{self.stage}'))
self.pool = tf.keras.layers.MaxPooling2D(pool_size=self.pool, strides=2, padding='same')

def call(self, inputs):
    x = inputs
    for i in range(len(self.f)):
        x = getattr(self, f'conv{i}_stage{self.stage}')(x)
        # x = vars(self)[f'leaky{i}_stage{self.stage}'](x)
        x = getattr(self, f'batch{i}_stage{self.stage}')(x)
    x = self.pool(x)
    return x

```

CREATE CUSTOM MODEL 1

```

In [ ]: class BGNet(tf.keras.Model):
    def __init__(self, img_h, img_w, img_c, batch_size, classes):
        super(BGNet, self).__init__(name='BGNet')
        self.img_h = img_h
        self.img_w = img_w
        self.img_c = img_c
        self.batch_size = batch_size
        self.classes = classes
        # (224, 224, 3)
        self.block0 = BhavyaBlock(f=[32, 32, 32, 32],
                                  k=[7, 5, 5, 5],
                                  s=[1, 1, 1, 1],
                                  p=['same', 'same', 'same', 'same'],
                                  d=[1, 2, 1, 2],
                                  pool=2,
                                  stage=0)
        # (112, 112, 32)
        self.block1 = BhavyaBlock(f=[64, 64, 64, 64],
                                  k=[5, 5, 5, 3],
                                  s=[1, 1, 1, 1],
                                  p=['same', 'same', 'same', 'same'],
                                  d=[2, 1, 2, 1],
                                  pool=2,
                                  stage=1)
        # (56, 56, 64)
        self.block2 = BhavyaBlock(f=[128, 128, 128, 128],
                                  k=[5, 5, 3, 3],
                                  s=[1, 1, 1, 1],
                                  p=['same', 'same', 'same', 'same'],
                                  d=[1, 2, 2, 1],
                                  pool=2,
                                  stage=2)
        # (28, 28, 128)
        self.block3 = BhavyaBlock(f=[256, 256, 256, 256],
                                  k=[5, 3, 3, 3],
                                  s=[1, 1, 1, 1],
                                  p=['same', 'same', 'same', 'same'],
                                  d=[1, 2, 2, 1],
                                  pool=2,
                                  stage=3)

```

```

        d=[2, 1, 1, 2],
        pool=2,
        stage=3)

# (14, 14, 512)
self.block4 = BhavyaBlock(f=[512, 512, 512, 512],
                           k=[3, 3, 3, 1],
                           s=[1, 1, 1, 1],
                           p=['same', 'same', 'same', 'same'],
                           d=[2, 2, 1, 1],
                           pool=2,
                           stage=4)

# (7, 7, 512)
self.block5 = BhavyaBlock(f=[1024, 1024, 1024],
                           k=[3, 3, 1],
                           s=[1, 2, 1],
                           p=['same', 'same', 'same'],
                           d=[1, 1, 2],
                           pool=2,
                           stage=5)

# (4, 4, 1024)
self.block6 = BhavyaBlock(f=[2048, 2048],
                           k=[1, 1],
                           s=[2, 1],
                           p=['same', 'same', 'same', 'same'],
                           d=[1, 1, 2, 1],
                           pool=2,
                           stage=6)

# (2, 2, 2048)
self.flatten = tf.keras.layers.Flatten(name='flatten')
self.dense1 = tf.keras.layers.Dense(1024, activation='tanh', name='dense1')
self.dropout1 = tf.keras.layers.Dropout(0.4, name='dropout1')
self.dense2 = tf.keras.layers.Dense(512, activation='relu', name='dense2')
self.dropout2 = tf.keras.layers.Dropout(0.3, name='dropout2')
self.dense3 = tf.keras.layers.Dense(256, activation='sigmoid', name='dense3')
self.dropout3 = tf.keras.layers.Dropout(0.2, name='dropout3')
self.final = tf.keras.layers.Dense(5 + len(self.classes), name='final')

def call(self, input_tensor, training=True) :
    x = self.block0(input_tensor)
    x = self.block1(x)
    x = self.block2(x)
    x = self.block3(x)
    x = self.block4(x)
    x = self.block5(x)
    x = self.block6(x)
    x = self.flatten(x)
    x = self.dense1(x)
    x = self.dropout1(x)
    x = self.dense2(x)
    x = self.dropout2(x)
    x = self.dense3(x)
    x = self.dropout3(x)
    x = self.final(x)
    return x

model1 = BhavyaNet(H, W, C, B, N)
model1.build(input_shape=(B, H, W, C))
model1.call(tf.keras.layers.Input(shape=(H, W, C), batch_size=B))
model1.summary(print_fn=tf.print, expand_nested=True, show_trainable=True)

```

Model: "BhavyaNet"

Layer (type)	Output Shape	Param #	Trainable
<hr/>			
BhavyaBlock0 (BhavyaBlock)	(5, 112, 112, 32)	82144	Y

BhavyaBlock1 (BhavyaBlock)	(5, 56, 56, 64)	294144	Y
BhavyaBlock2 (BhavyaBlock)	(5, 28, 28, 128)	911872	Y
BhavyaBlock3 (BhavyaBlock)	(5, 14, 14, 256)	2593792	Y
BhavyaBlock4 (BhavyaBlock)	(5, 7, 7, 512)	6170624	Y
BhavyaBlock5 (BhavyaBlock)	(5, 2, 2, 1024)	15219712	Y
BhavyaBlock6 (BhavyaBlock)	(5, 1, 1, 2048)	6311936	Y
flatten (Flatten)	(5, 2048)	0	Y
dense1 (Dense)	(5, 1024)	2098176	Y
dropout1 (Dropout)	(5, 1024)	0	Y
dense2 (Dense)	(5, 512)	524800	Y
dropout2 (Dropout)	(5, 512)	0	Y
dense3 (Dense)	(5, 256)	131328	Y
dropout3 (Dropout)	(5, 256)	0	Y
final (Dense)	(5, 8)	2056	Y
<hr/>			
Total params: 34,340,584			
Trainable params: 34,318,312			
Non-trainable params: 22,272			

CREATING A CUSTOM LOSS

```
In [ ]: class CustomLoss1(tf.keras.losses.Loss):
    def __init__(self, classes):
        super(CustomLoss1, self).__init__()
        self.classes = classes

    def unpack_values(self, y):
        c = tf.slice(y, [0, 0], [-1, len(self.classes)])
        p = tf.slice(y, [0, len(self.classes) + 0], [-1, 1])
        x1 = tf.slice(y, [0, len(self.classes) + 1], [-1, 1])
        y1 = tf.slice(y, [0, len(self.classes) + 2], [-1, 1])
        x2 = tf.slice(y, [0, len(self.classes) + 3], [-1, 1])
        y2 = tf.slice(y, [0, len(self.classes) + 4], [-1, 1])
        return c, p, x1, y1, x2, y2

    def call(self, y_true, y_pred):
        c_t, p_t, x1_t, y1_t, x2_t, y2_t = self.unpack_values(y_true)
        c_p, p_p, x1_p, y1_p, x2_p, y2_p = self.unpack_values(y_pred)
        c_loss = tf.keras.losses.mse(c_t, c_p)
        p_loss = tf.keras.losses.mse(p_t, p_p)
        x1_loss = tf.keras.losses.mse(x1_t, x1_p)
        y1_loss = tf.keras.losses.mse(y1_t, y1_p)
        x2_loss = tf.keras.losses.mse(x2_t, x2_p)
        y2_loss = tf.keras.losses.mse(y2_t, y2_p)
        loss = c_loss + p_loss + x1_loss + y1_loss + x2_loss + y2_loss
        return tf.reduce_mean(loss)

loss = CustomLoss1(N)
```

CREATING A CUSTOM METRIC FOR LOCALIZATION

In []:

```
class CustomMetric1(tf.keras.metrics.Metric):
    def __init__(self, name, classes):
        super(CustomMetric1, self).__init__(name=name)
        self.classes = classes
        self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
        self.examples = self.add_weight(name='total_examples', initializer='zeros')

    def unpack_values(self, y):
        c = tf.slice(y, [0, 0], [-1, len(self.classes)])
        p = tf.slice(y, [0, len(self.classes) + 0], [-1, 1])
        x1 = tf.slice(y, [0, len(self.classes) + 1], [-1, 1])
        y1 = tf.slice(y, [0, len(self.classes) + 2], [-1, 1])
        x2 = tf.slice(y, [0, len(self.classes) + 3], [-1, 1])
        y2 = tf.slice(y, [0, len(self.classes) + 4], [-1, 1])
        return c, p, x1, y1, x2, y2

    def get_iou(self, x1_t, y1_t, x2_t, y2_t, x1_p, y1_p, x2_p, y2_p):
        x1_i = tf.maximum(x1_t, x1_p)
        y1_i = tf.maximum(y1_t, y1_p)
        x2_i = tf.minimum(x2_t, x2_p)
        y2_i = tf.minimum(y2_t, y2_p)
        area_i = tf.maximum(0.0, x2_i - x1_i + 1.0) * tf.maximum(0.0, y2_i - y1_i + 1.0)
        area_t = (x2_t - x1_t + 1.0) * (y2_t - y1_t + 1.0)
        area_p = (x2_p - x1_p + 1.0) * (y2_p - y1_p + 1.0)
        area_u = area_t + area_p - area_i
        return area_i / area_u

    def update_state(self, y_true, y_pred, sample_weight=None):
        c_t, p_t, x1_t, y1_t, x2_t, y2_t = self.unpack_values(y_true)
        c_p, p_p, x1_p, y1_p, x2_p, y2_p = self.unpack_values(y_pred)
        iou = self.get_iou(x1_t, y1_t, x2_t, y2_t, x1_p, y1_p, x2_p, y2_p)
        self.total_metric.assign_add(tf.reduce_mean(iou))
        self.examples.assign_add(1.0)

    def result(self):
        return self.total_metric / self.examples

    def reset_state(self):
        self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
        self.examples = self.add_weight(name='total_examples', initializer='zeros')
metric1 = CustomMetric1('iou_metric', N)
```

CREATING A CUSTOM METRIC FOR CONFIDENCE

In []:

```
class CustomMetric2(tf.keras.metrics.Metric):
    def __init__(self, name, classes):
        super(CustomMetric2, self).__init__(name=name)
        self.classes = classes
        self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
        self.examples = self.add_weight(name='total_examples', initializer='zeros')

    def unpack_values(self, y):
        c = tf.slice(y, [0, 0], [-1, len(self.classes)])
        p = tf.slice(y, [0, len(self.classes) + 0], [-1, 1])
        x1 = tf.slice(y, [0, len(self.classes) + 1], [-1, 1])
        y1 = tf.slice(y, [0, len(self.classes) + 2], [-1, 1])
        x2 = tf.slice(y, [0, len(self.classes) + 3], [-1, 1])
        y2 = tf.slice(y, [0, len(self.classes) + 4], [-1, 1])
        return c, p, x1, y1, x2, y2
```

```

def update_state(self, y_true, y_pred, sample_weight=None) :
    c_t, p_t, x1_t, y1_t, x2_t, y2_t = self.unpack_values(y_true)
    c_p, p_p, x1_p, y1_p, x2_p, y2_p = self.unpack_values(y_pred)
    b_acc = tf.keras.metrics.binary_accuracy(p_t, p_p)
    self.total_metric.assign_add(tf.reduce_mean(b_acc))
    self.examples.assign_add(1.0)

def result(self) :
    return self.total_metric / self.examples

def reset_state(self) :
    self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
    self.examples = self.add_weight(name='total_examples', initializer='zeros')
metric2 = CustomMetric2('p_metric', N)

```

CREATING A CUSTOM METRIC FOR CLASSIFICATION

```

In [ ]: class CustomMetric3(tf.keras.metrics.Metric) :
    def __init__(self, name, classes) :
        super(CustomMetric3, self). __init__(name=name)
        self.classes = classes
        self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
        self.examples = self.add_weight(name='total_examples', initializer='zeros')

    def unpack_values(self, y) :
        c = tf.slice(y, [0, 0], [-1, len(self.classes)])
        p = tf.slice(y, [0, len(self.classes) + 0], [-1, 1])
        x1 = tf.slice(y, [0, len(self.classes) + 1], [-1, 1])
        y1 = tf.slice(y, [0, len(self.classes) + 2], [-1, 1])
        x2 = tf.slice(y, [0, len(self.classes) + 3], [-1, 1])
        y2 = tf.slice(y, [0, len(self.classes) + 4], [-1, 1])
        return c, p, x1, y1, x2, y2

    def update_state(self, y_true, y_pred, sample_weight=None) :
        c_t, p_t, x1_t, y1_t, x2_t, y2_t = self.unpack_values(y_true)
        c_p, p_p, x1_p, y1_p, x2_p, y2_p = self.unpack_values(y_pred)
        c_acc = tf.keras.metrics.categorical_accuracy(c_t, c_p)
        self.total_metric.assign_add(tf.reduce_mean(c_acc))
        self.examples.assign_add(1.0)

    def result(self) :
        return self.total_metric / self.examples

    def reset_state(self) :
        self.total_metric = self.add_weight(name='total_metric', initializer='zeros')
        self.examples = self.add_weight(name='total_examples', initializer='zeros')
metric3 = CustomMetric3('c_metric', N)

```

CREATE A CUSTOM CALLBACK TO PRINT RESIZED IMAGES AND BOUNDING BOXES AFTER EVERY EPOCH

```

In [ ]: class CustomCallback1(tf.keras.callbacks.Callback) :
    def __init__(self, dataframe, xcol, ycol, cols, img_h, img_w, img_c, classes) :
        super(CustomCallback1, self). __init__()
        self.dataframe = dataframe
        self.xcol = xcol
        self.ycol = ycol
        self.rows = 1
        self.cols = cols

```

```

    self.img_h = img_h
    self.img_w = img_w
    self.img_c = img_c
    self.classes = classes

    def get_iou(self, x1_t, y1_t, x2_t, y2_t, x1_p, y1_p, x2_p, y2_p):
        x1_i = max(x1_t, x1_p)
        y1_i = max(y1_t, y1_p)
        x2_i = min(x2_t, x2_p)
        y2_i = min(y2_t, y2_p)
        area_i = max(0, x2_i - x1_i) * max(0, y2_i - y1_i)
        area_t = (x2_t - x1_t) * (y2_t - y1_t)
        area_p = (x2_p - x1_p) * (y2_p - y1_p)
        area_u = area_t + area_p - area_i
        return area_i / area_u

    def predict(self, image_path):
        img = tf.io.read_file(image_path)
        img = tf.io.decode_image(img, channels=self.img_c)
        img = tf.image.resize(img, (self.img_h, self.img_w))
        img = tf.expand_dims(img, axis=0)
        img = img / 255.0
        pred = self.model.predict(img)
        return pred

    def on_epoch_end(self, epoch, logs=None):
        fig, ax = plt.subplots(self.rows, self.cols, figsize=(28, 13))
        for i in range(self.cols):
            index = rn.randint(0, self.dataframe.shape[0] - 1)
            img_path = self.dataframe[self.xcol][index]
            img = cv.imread(img_path)
            h, w, _ = img.shape
            w_ratio = w / self.img_w
            h_ratio = h / self.img_h
            cat = self.dataframe[self.ycol[0]][index]
            cat = self.classes[cat]
            x1_t = int(self.dataframe[self.ycol[1]][index] / w_ratio)
            y1_t = int(self.dataframe[self.ycol[2]][index] / h_ratio)
            x2_t = int(self.dataframe[self.ycol[3]][index] / w_ratio)
            y2_t = int(self.dataframe[self.ycol[4]][index] / h_ratio)
            img = cv.cvtColor(img, cv.COLOR_BGR2RGB)
            img = cv.resize(img, (self.img_w, self.img_h))
            img = cv.rectangle(img, (x1_t, y1_t), (x2_t, y2_t), (0, 255, 0), 1)
            # img = cv.putText(img, cat, (x1, y1 + 10), cv.FONT_HERSHEY_PLAIN, 1, (0, 255, 0))
            pred = self.predict(img_path)
            c = np.argmax(pred[0][0: len(self.classes)], axis=0)
            conf = round(pred[0][0: len(self.classes)][c], 3)
            c = self.classes[c]
            p = round(pred[0][len(self.classes) + 0], 3)
            x1_p = int(pred[0][len(self.classes) + 1] * self.img_w)
            y1_p = int(pred[0][len(self.classes) + 2] * self.img_h)
            x2_p = int(pred[0][len(self.classes) + 3] * self.img_w)
            y2_p = int(pred[0][len(self.classes) + 4] * self.img_h)
            img = cv.rectangle(img, (x1_p, y1_p), (x2_p, y2_p), (255, 0, 255), 1)
            # img = cv.putText(img, f'{c}: {p}', (x1, y1 + 10), cv.FONT_HERSHEY_PLAIN, 1, (0, 255, 0))
            ax[i].imshow(img)
            ax[i].set_ylabel(f'TRUE LABEL: {cat}')
            ax[i].set_xlabel(f'PRED LABEL: {c}, C: {conf}, P: {p}')
            ax[i].set_title(f'IOU: {self.get_iou(x1_t, y1_t, x2_t, y2_t, x1_p, y1_p, x2_p, y2_p)}')
        plt.show()

cb1 = CustomCallback1(val_df, 'filename', ['class', 'xmin', ' ymin', 'xmax', 'ymax'], 4,

```

TEST PREDICTION RANGE

```
In [ ]: modell.predict(np.random.rand(1, H, W, C))

Out[ ]: array([[nan, nan, nan, nan, nan, nan, nan, nan]], dtype=float32)
```

DEFINE NECESSARY OBJECTS AND FUNCTIONS

```
In [ ]: e = 1000
es = tf.keras.callbacks.EarlyStopping(monitor='val_loss', mode='min', patience=5, restore_best_weights=True)
rlrop = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', mode='min', patience=3, factor=0.5)
metric = [metric1, metric2, metric3]
callback = [es, rlrop, cb1]

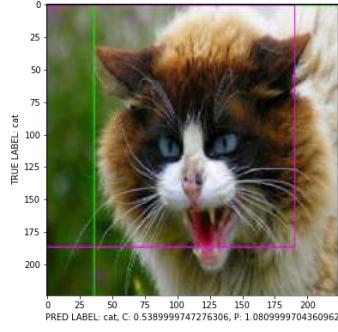
def plot_metric(f, m):
    plt.figure(figsize=(14, 8))
    plt.plot(f.history[m])
    plt.plot(f.history['val_' + m])
    plt.title('model ' + m)
    plt.ylabel(m)
    plt.xlabel('epoch')
    plt.legend(['train', 'val'], loc='upper left')
    plt.show()
```

START MODEL 1 TRAINING

```
In [ ]: modell.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
                      loss=loss,
                      metrics=metric)
fit1 = modell.fit(train_gen,
                   epochs=e,
                   batch_size=B,
                   callbacks=callback,
                   validation_data=val_gen,
                   validation_batch_size=B)
```

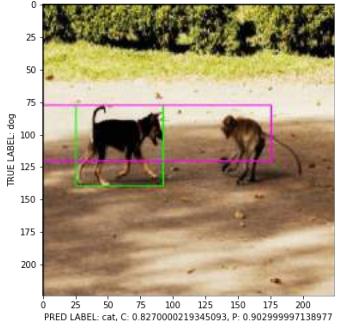
Epoch 1/100
93/93 [=====] - ETA: 0s - loss: 141421.8125 - iou_metric: 0.338
8 - p_metric: 0.6903 - c_metric: 0.3054

IOU: 0.5009619084263178



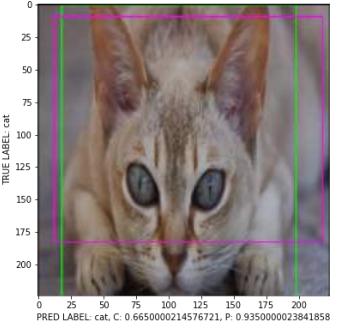
TRUE LABEL: cat
PRED LABEL: cat, C: 0.5389999747276306, P: 1.0809999704360962

IOU: 0.31220199393151277



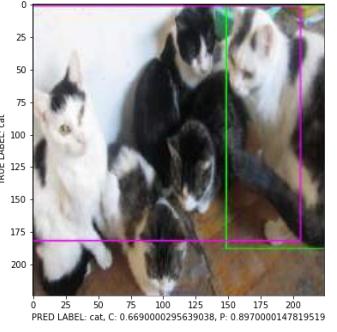
TRUE LABEL: dog
PRED LABEL: dog, C: 0.8270000219345093, P: 0.902999997138977

IOU: 0.6948101209335534



TRUE LABEL: cat
PRED LABEL: cat, C: 0.6650000214576721, P: 0.935000023841858

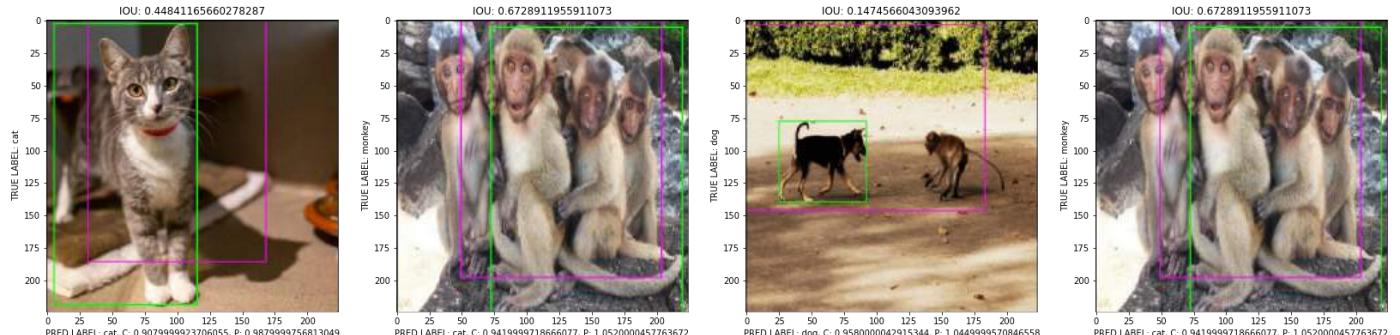
IOU: 0.23174918684495843



TRUE LABEL: cat
PRED LABEL: cat, C: 0.6690000295639038, P: 0.897000147819519

93/93 [=====] - 33s 227ms/step - loss: 141421.8125 - iou_metric: 0.3388 - p_metric: 0.6903 - c_metric: 0.3054 - val_loss: 140997.9062 - val_iou_metric: 0.6628 - val_p_metric: 1.0000 - val_c_metric: 0.2800 - lr: 1.0000e-04

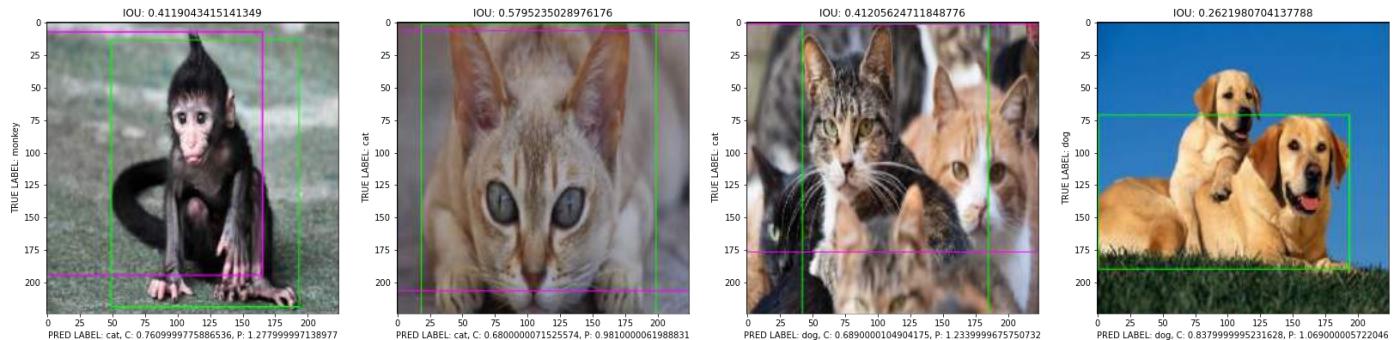
Epoch 2/100
93/93 [=====] - ETA: 0s - loss: 140605.2969 - iou_metric: 0.3939 - p_metric: 0.7753 - c_metric: 0.3011



93/93 [=====] - 21s 226ms/step - loss: 140605.2969 - iou_metric: 0.3939 - p_metric: 0.7753 - c_metric: 0.3011 - val_loss: 140193.7969 - val_iou_metric: 0.6701 - val_p_metric: 1.0000 - val_c_metric: 0.2200 - lr: 1.0000e-04

Epoch 3/100

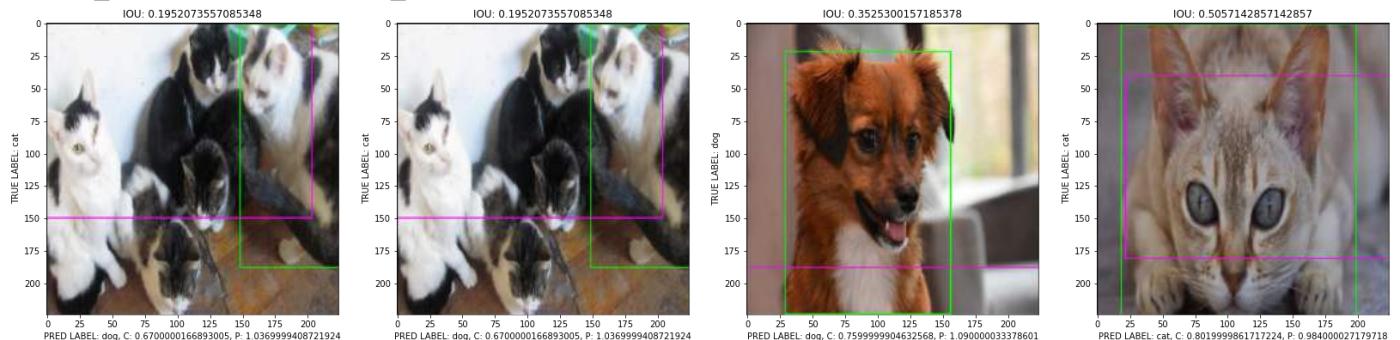
93/93 [=====] - ETA: 0s - loss: 139769.5312 - iou_metric: 0.4167 - p_metric: 0.8093 - c_metric: 0.3075



93/93 [=====] - 20s 215ms/step - loss: 139769.5312 - iou_metric: 0.4167 - p_metric: 0.8093 - c_metric: 0.3075 - val_loss: 139336.7656 - val_iou_metric: 0.6697 - val_p_metric: 1.0000 - val_c_metric: 0.2533 - lr: 1.0000e-04

Epoch 4/100

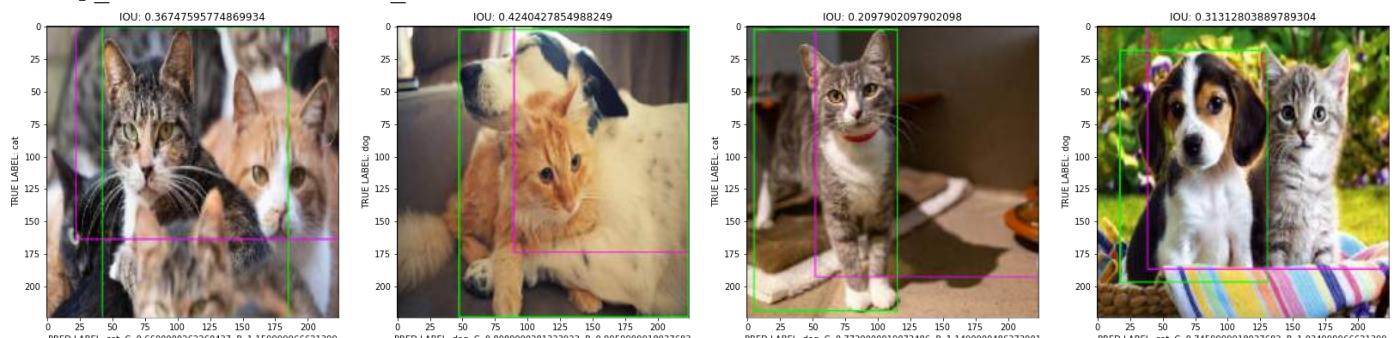
93/93 [=====] - ETA: 0s - loss: 138906.0938 - iou_metric: 0.4314 - p_metric: 0.8355 - c_metric: 0.3177



93/93 [=====] - 19s 203ms/step - loss: 138906.0938 - iou_metric: 0.4314 - p_metric: 0.8355 - c_metric: 0.3177 - val_loss: 138465.8594 - val_iou_metric: 0.6736 - val_p_metric: 1.0000 - val_c_metric: 0.2400 - lr: 1.0000e-04

Epoch 5/100

93/93 [=====] - ETA: 0s - loss: 138061.9375 - iou_metric: 0.4448 - p_metric: 0.8516 - c_metric: 0.3247



93/93 [=====] - 19s 203ms/step - loss: 138061.9375 - iou_metric: 0.4448 - p_metric: 0.8516 - c_metric: 0.3247 - val_loss: 137623.5938 - val_iou_metric: 0.6705 - val_p_metric: 1.0000 - val_c_metric: 0.2800 - lr: 1.0000e-04

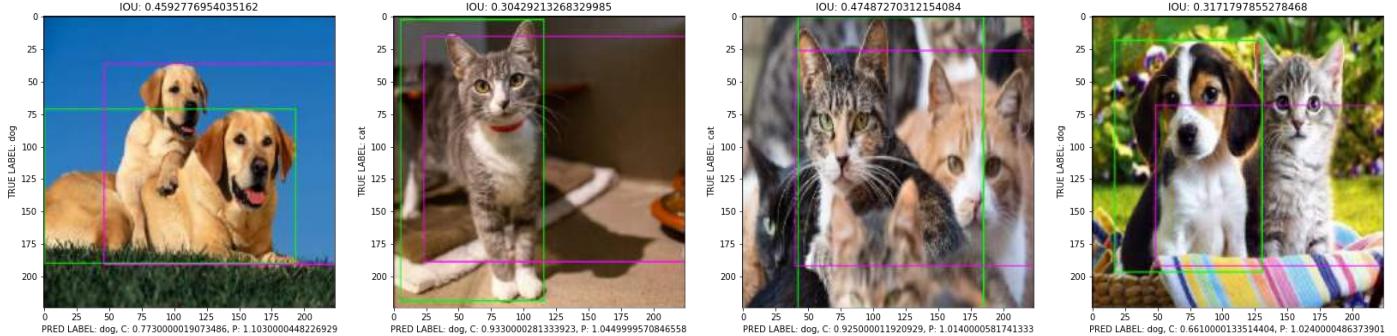
Epoch 6/100



93/93 [=====] - 19s 209ms/step - loss: 137213.0312 - iou_metric: 0.4557 - p_metric: 0.8638 - c_metric: 0.3237 - val_loss: 136784.3750 - val_iou_metric: 0.6765 - val_p_metric: 1.0000 - val_c_metric: 0.3000 - lr: 1.0000e-04

Epoch 7/100

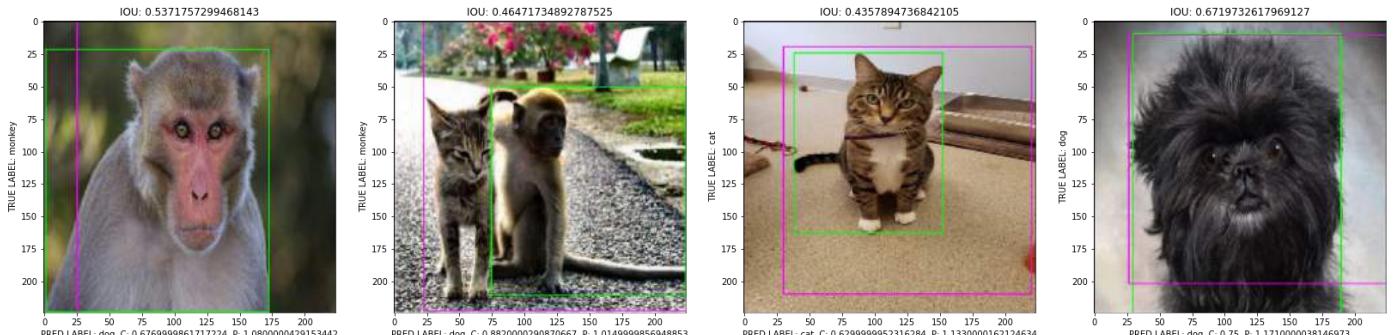
93/93 [=====] - ETA: 0s - loss: 136348.3750 - iou_metric: 0.466
 3 - p_metric: 0.8722 - c_metric: 0.3269



93/93 [=====] - 19s 205ms/step - loss: 136348.3750 - iou_metric: 0.4663 - p_metric: 0.8722 - c_metric: 0.3269 - val_loss: 135912.1406 - val_iou_metric: 0.6739 - val_p_metric: 1.0000 - val_c_metric: 0.2971 - lr: 1.0000e-04

Epoch 8/100

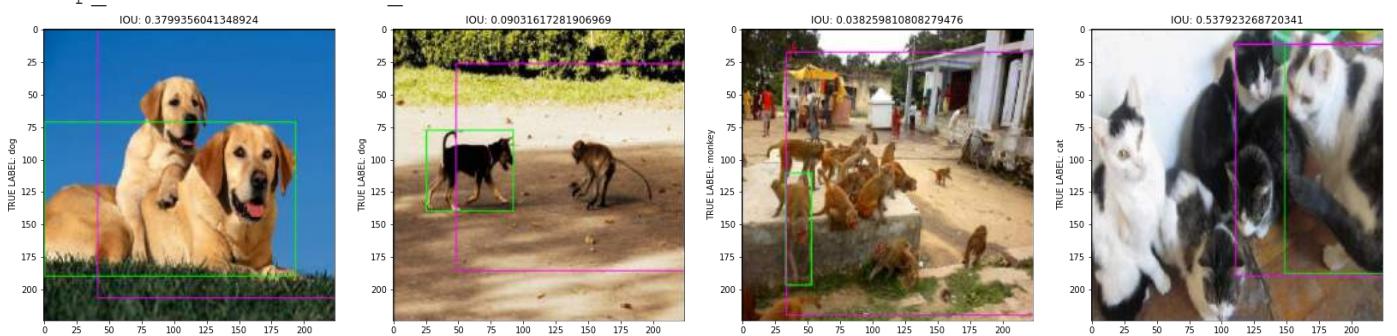
93/93 [=====] - ETA: 0s - loss: 135480.9531 - iou_metric: 0.475
 8 - p_metric: 0.8815 - c_metric: 0.3347



93/93 [=====] - 19s 206ms/step - loss: 135480.9531 - iou_metric: 0.4758 - p_metric: 0.8815 - c_metric: 0.3347 - val_loss: 135040.7031 - val_iou_metric: 0.6757 - val_p_metric: 1.0000 - val_c_metric: 0.3050 - lr: 1.0000e-04

Epoch 9/100

93/93 [=====] - ETA: 0s - loss: 134615.7969 - iou_metric: 0.485
 6 - p_metric: 0.8903 - c_metric: 0.3419



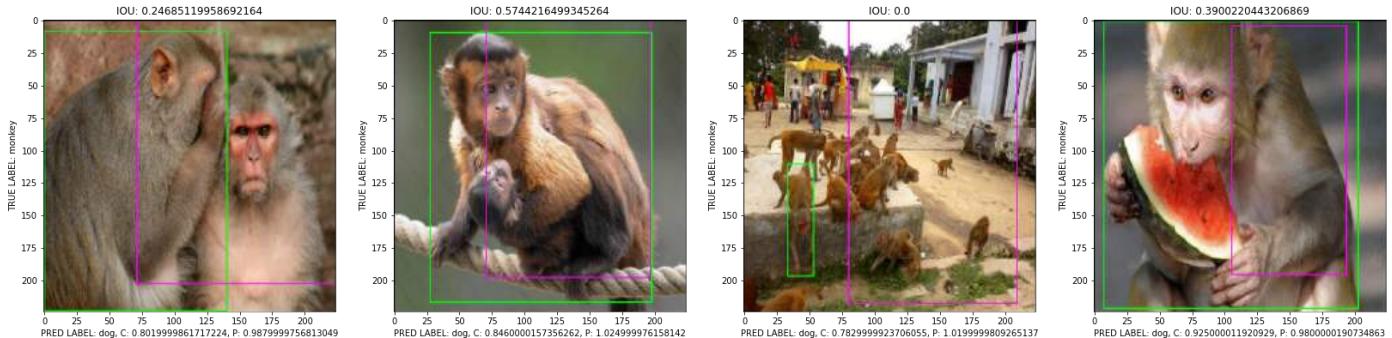
93/93 [=====] - 19s 206ms/step - loss: 134615.7969 - iou_metric: 0.4856 - p_metric: 0.8903 - c_metric: 0.3419 - val_loss: 134160.8750 - val_iou_metric:

c: 0.6781 - val_p_metric: 1.0000 - val_c_metric: 0.3200 - lr: 1.0000e-04

Epoch 10/100

93/93 [=====] - ETA: 0s - loss: 133706.9219 - iou_metric: 0.495

0 - p_metric: 0.8985 - c_metric: 0.3404



93/93 [=====] - 19s 206ms/step - loss: 133706.9219 - iou_metric

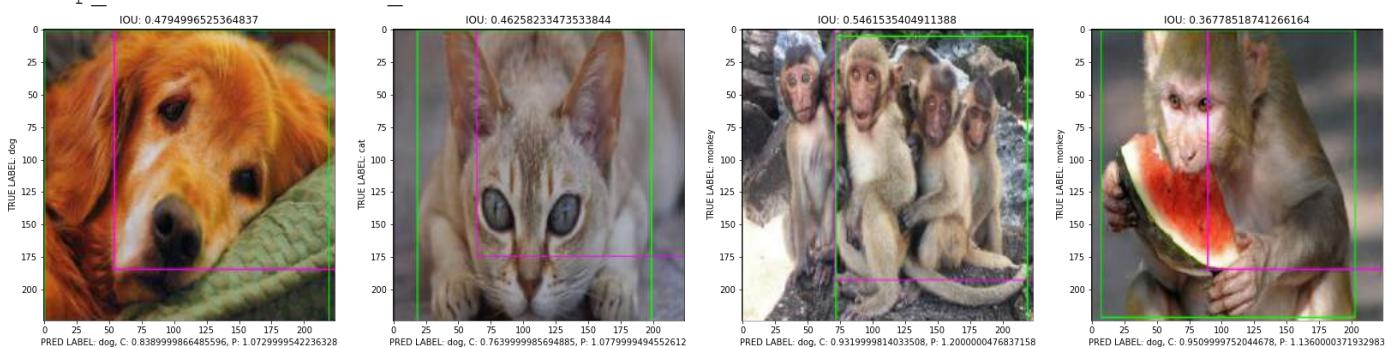
c: 0.4950 - p_metric: 0.8985 - c_metric: 0.3404 - val_loss: 133244.1406 - val_iou_metric

c: 0.6795 - val_p_metric: 1.0000 - val_c_metric: 0.3320 - lr: 1.0000e-04

Epoch 11/100

93/93 [=====] - ETA: 0s - loss: 132804.3125 - iou_metric: 0.502

7 - p_metric: 0.9054 - c_metric: 0.3429



93/93 [=====] - 19s 208ms/step - loss: 132804.3125 - iou_metric

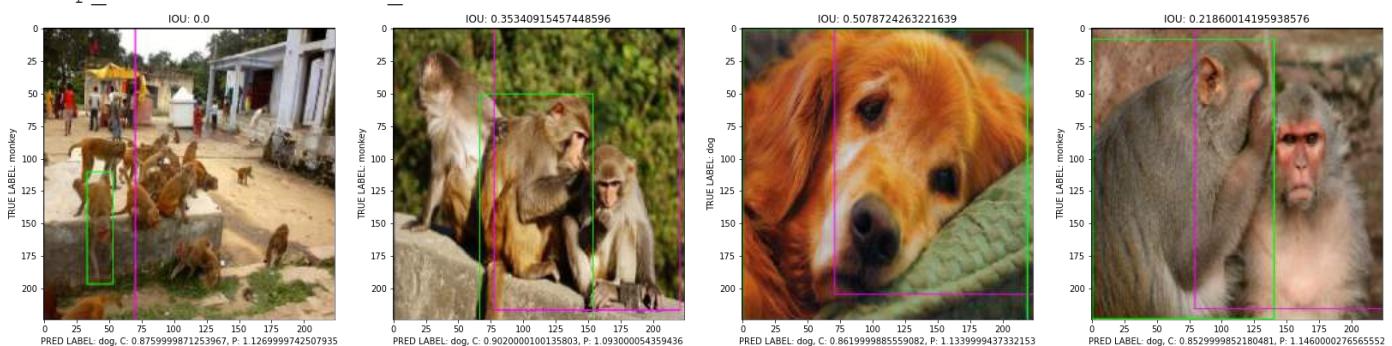
c: 0.5027 - p_metric: 0.9054 - c_metric: 0.3429 - val_loss: 132356.6875 - val_iou_metric

c: 0.6802 - val_p_metric: 1.0000 - val_c_metric: 0.3491 - lr: 1.0000e-04

Epoch 12/100

93/93 [=====] - ETA: 0s - loss: 131909.8438 - iou_metric: 0.509

6 - p_metric: 0.9118 - c_metric: 0.3446



93/93 [=====] - 19s 209ms/step - loss: 131909.8438 - iou_metric

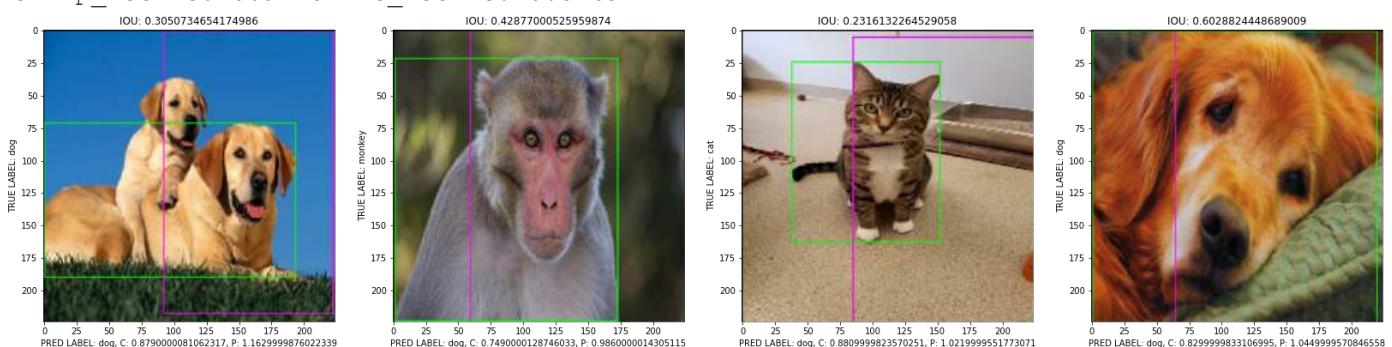
c: 0.5096 - p_metric: 0.9118 - c_metric: 0.3446 - val_loss: 131444.8438 - val_iou_metric

c: 0.6839 - val_p_metric: 1.0000 - val_c_metric: 0.3500 - lr: 1.0000e-04

Epoch 13/100

93/93 [=====] - ETA: 0s - loss: 130983.0469 - iou_metric: 0.517

3 - p_metric: 0.9175 - c_metric: 0.3469



93/93 [=====] - 19s 208ms/step - loss: 130983.0469 - iou_metric
c: 0.5173 - p_metric: 0.9175 - c_metric: 0.3469 - val_loss: 130521.6562 - val_iou_metric
c: 0.6871 - val_p_metric: 1.0000 - val_c_metric: 0.3538 - lr: 1.0000e-04

Epoch 14/100

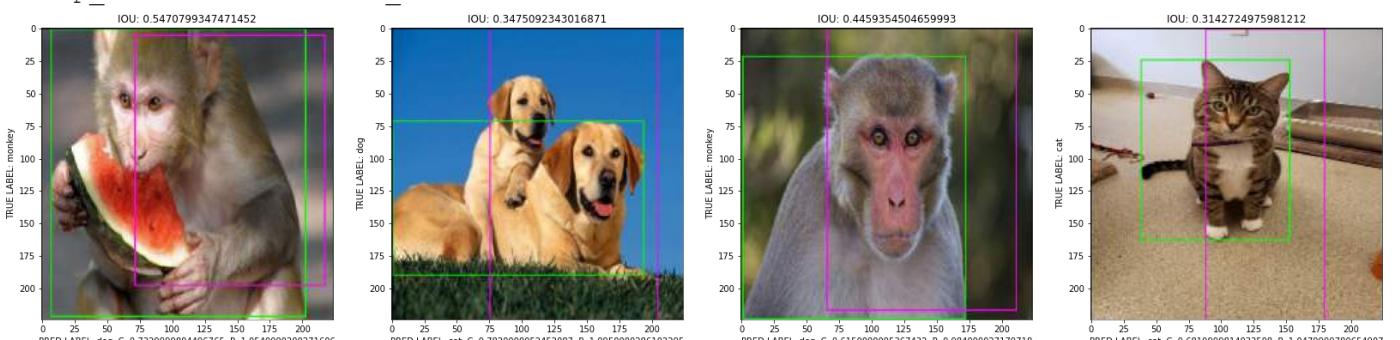
93/93 [=====] - ETA: 0s - loss: 130060.1641 - iou_metric: 0.524
6 - p_metric: 0.9226 - c_metric: 0.3464



93/93 [=====] - 20s 210ms/step - loss: 130060.1641 - iou_metric
c: 0.5246 - p_metric: 0.9226 - c_metric: 0.3464 - val_loss: 129593.5312 - val_iou_metric
c: 0.6878 - val_p_metric: 1.0000 - val_c_metric: 0.3514 - lr: 1.0000e-04

Epoch 15/100

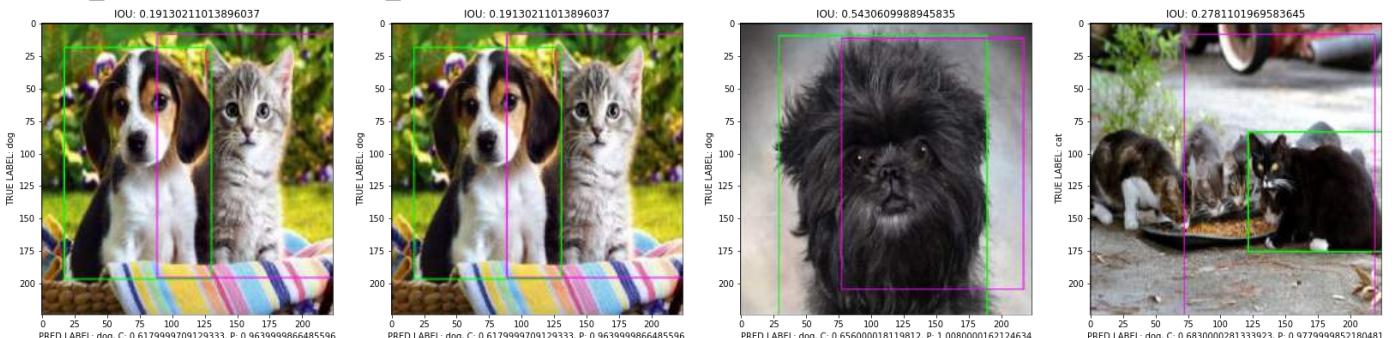
93/93 [=====] - ETA: 0s - loss: 129121.2578 - iou_metric: 0.532
3 - p_metric: 0.9273 - c_metric: 0.3495



93/93 [=====] - 19s 207ms/step - loss: 129121.2578 - iou_metric
c: 0.5323 - p_metric: 0.9273 - c_metric: 0.3495 - val_loss: 128634.8203 - val_iou_metric
c: 0.6904 - val_p_metric: 1.0000 - val_c_metric: 0.3493 - lr: 1.0000e-04

Epoch 16/100

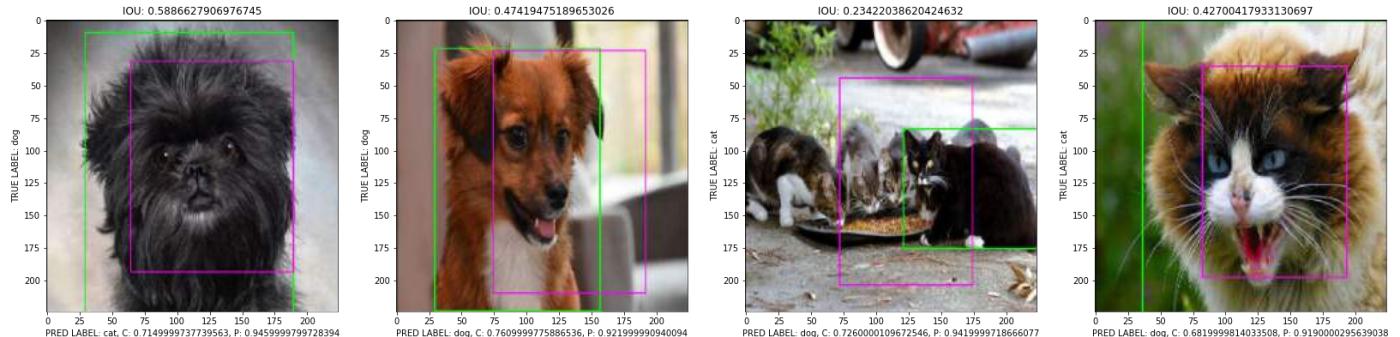
93/93 [=====] - ETA: 0s - loss: 128169.5391 - iou_metric: 0.539
4 - p_metric: 0.9319 - c_metric: 0.3504



93/93 [=====] - 19s 207ms/step - loss: 128169.5391 - iou_metric
c: 0.5394 - p_metric: 0.9319 - c_metric: 0.3504 - val_loss: 127682.6562 - val_iou_metric
c: 0.6943 - val_p_metric: 1.0000 - val_c_metric: 0.3700 - lr: 1.0000e-04

Epoch 17/100

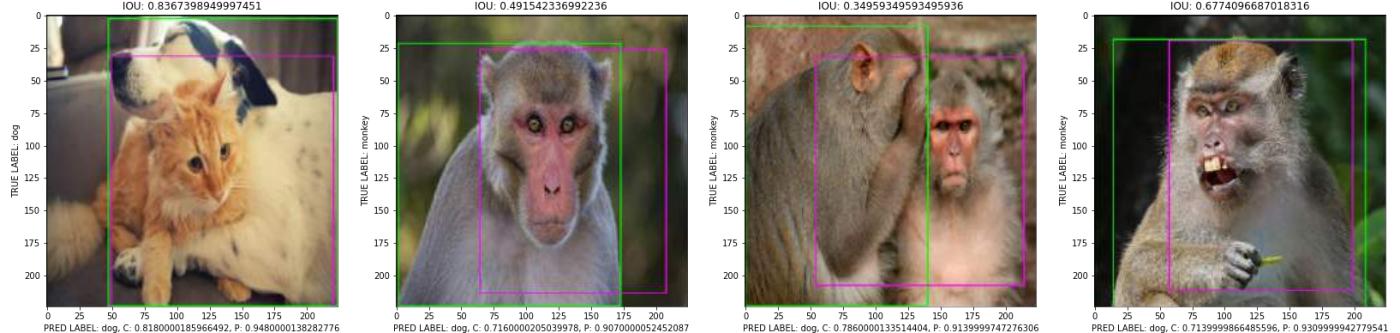
93/93 [=====] - ETA: 0s - loss: 127229.0703 - iou_metric: 0.546
0 - p_metric: 0.9356 - c_metric: 0.3523



93/93 [=====] - 19s 208ms/step - loss: 127229.0703 - iou_metric: 0.5460 - p_metric: 0.9356 - c_metric: 0.3523 - val_loss: 126760.2188 - val_iou_metric: 0.6969 - val_p_metric: 1.0000 - val_c_metric: 0.3788 - lr: 1.0000e-04

Epoch 18/100

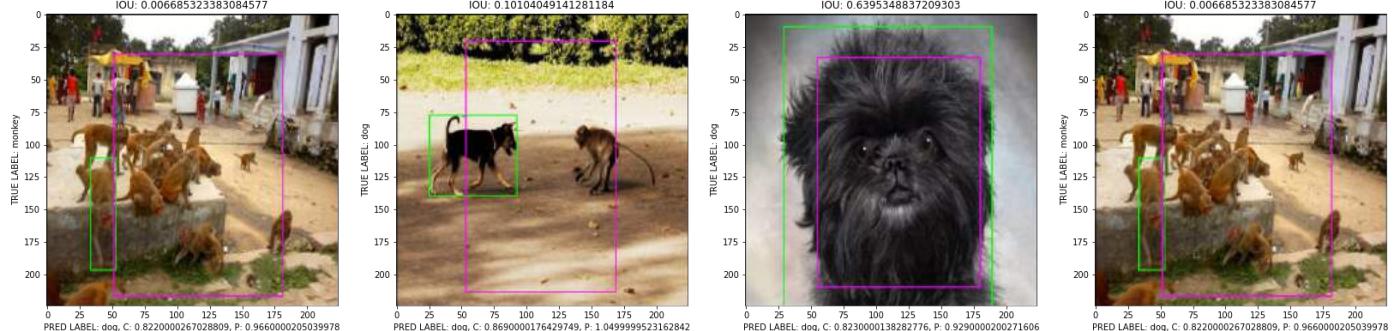
93/93 [=====] - ETA: 0s - loss: 126300.7656 - iou_metric: 0.5524 - p_metric: 0.9392 - c_metric: 0.3538



93/93 [=====] - 20s 210ms/step - loss: 126300.7656 - iou_metric: 0.5524 - p_metric: 0.9392 - c_metric: 0.3538 - val_loss: 125836.5312 - val_iou_metric: 0.7002 - val_p_metric: 1.0000 - val_c_metric: 0.3800 - lr: 1.0000e-04

Epoch 19/100

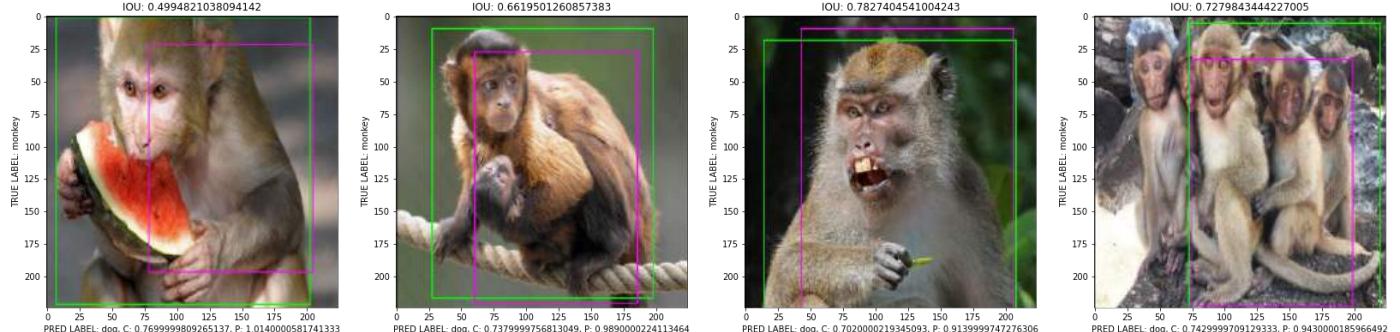
93/93 [=====] - ETA: 0s - loss: 125373.3359 - iou_metric: 0.5587 - p_metric: 0.9422 - c_metric: 0.3548



93/93 [=====] - 20s 214ms/step - loss: 125373.3359 - iou_metric: 0.5587 - p_metric: 0.9422 - c_metric: 0.3548 - val_loss: 124898.0000 - val_iou_metric: 0.7006 - val_p_metric: 1.0000 - val_c_metric: 0.3895 - lr: 1.0000e-04

Epoch 20/100

93/93 [=====] - ETA: 0s - loss: 124438.2891 - iou_metric: 0.5647 - p_metric: 0.9451 - c_metric: 0.3554



93/93 [=====] - 20s 217ms/step - loss: 124438.2891 - iou_metric: 0.5647 - p_metric: 0.9451 - c_metric: 0.3554 - val_loss: 123963.6562 - val_iou_metric: 0.7026 - val_p_metric: 1.0000 - val_c_metric: 0.3960 - lr: 1.0000e-04

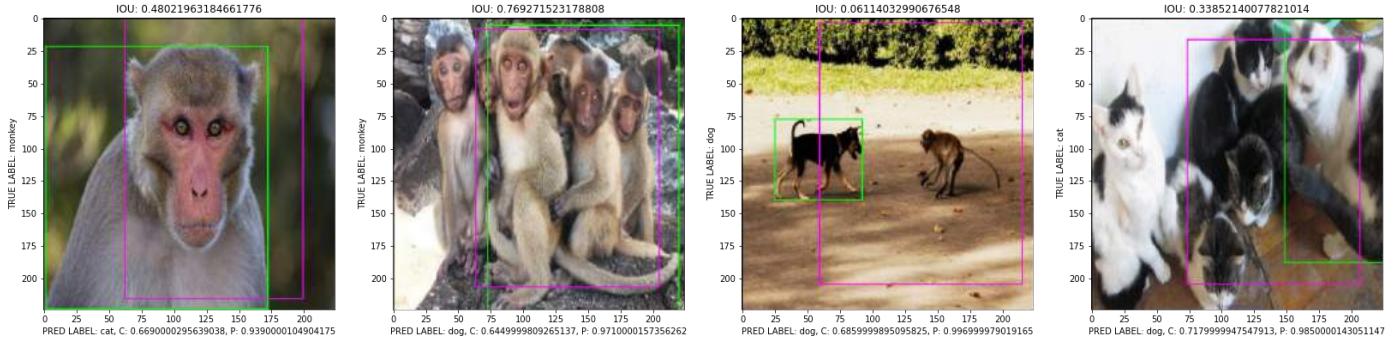
Epoch 21/100

93/93 [=====] - ETA: 0s - loss: 123507.1016 - iou_metric: 0.570
0 - p_metric: 0.9477 - c_metric: 0.3570

93/93 [=====] - 20s 218ms/step - loss: 123507.1016 - iou_metric: 0.5700 - p_metric: 0.9477 - c_metric: 0.3570 - val_loss: 123043.7578 - val_iou_metric: 0.7042 - val_p_metric: 1.0000 - val_c_metric: 0.4000 - lr: 1.0000e-04

Epoch 22/100

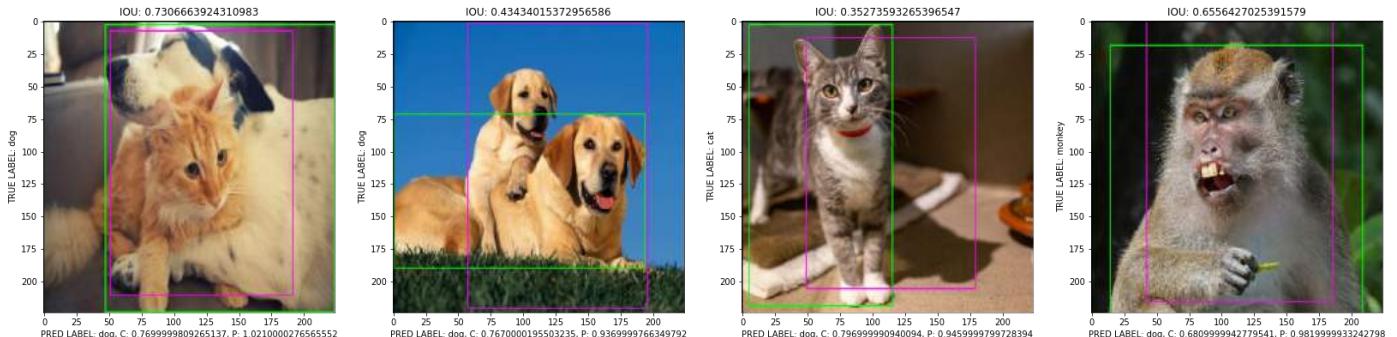
93/93 [=====] - ETA: 0s - loss: 122585.9531 - iou_metric: 0.575
4 - p_metric: 0.9500 - c_metric: 0.3577



93/93 [=====] - 20s 218ms/step - loss: 122585.9531 - iou_metric: 0.5754 - p_metric: 0.9500 - c_metric: 0.3577 - val_loss: 122117.6094 - val_iou_metric: 0.7065 - val_p_metric: 1.0000 - val_c_metric: 0.4091 - lr: 1.0000e-04

Epoch 23/100

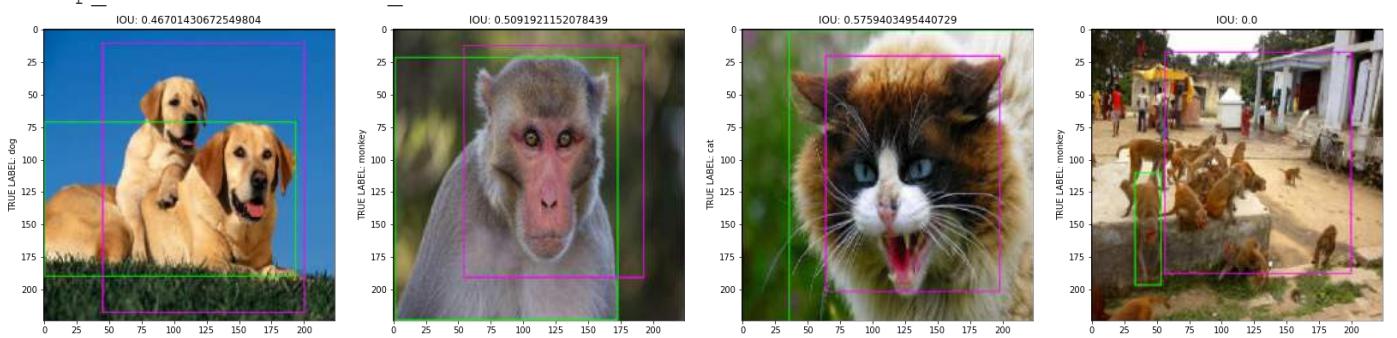
93/93 [=====] - ETA: 0s - loss: 121655.0547 - iou_metric: 0.580
4 - p_metric: 0.9522 - c_metric: 0.3583



93/93 [=====] - 20s 219ms/step - loss: 121655.0547 - iou_metric: 0.5804 - p_metric: 0.9522 - c_metric: 0.3583 - val_loss: 121185.1094 - val_iou_metric: 0.7089 - val_p_metric: 1.0000 - val_c_metric: 0.4104 - lr: 1.0000e-04

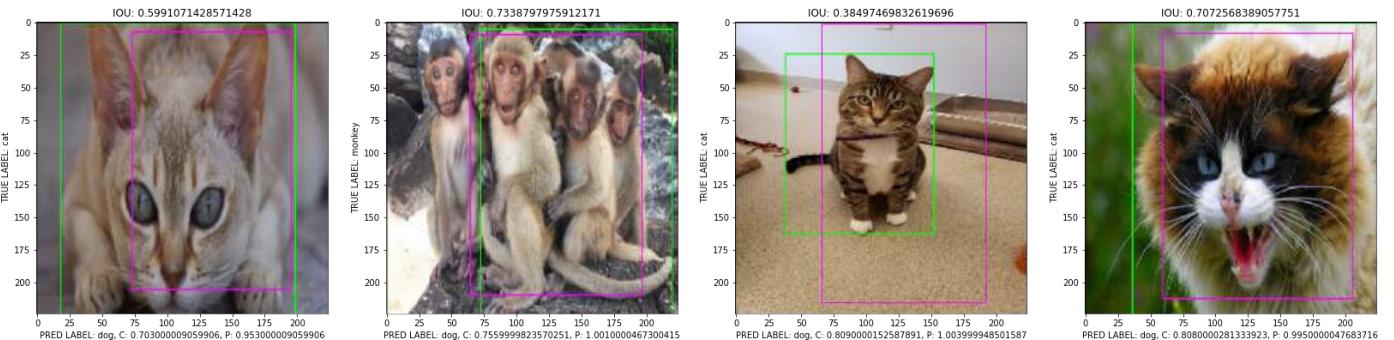
Epoch 24/100

93/93 [=====] - ETA: 0s - loss: 120726.8359 - iou_metric: 0.585
0 - p_metric: 0.9542 - c_metric: 0.3577



93/93 [=====] - 20s 219ms/step - loss: 120726.8359 - iou_metric: 0.5850 - p_metric: 0.9542 - c_metric: 0.3577 - val_loss: 120255.7266 - val_iou_metric:

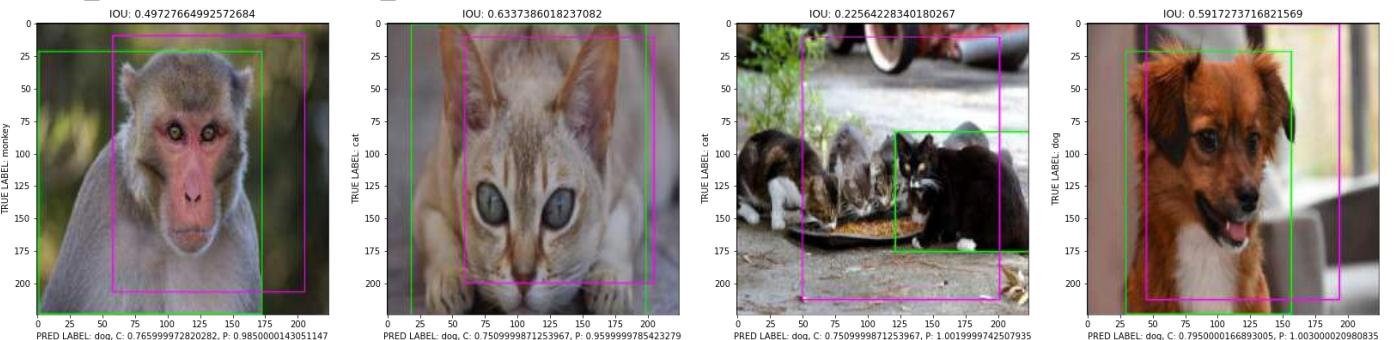
c: 0.7104 - val_p_metric: 1.0000 - val_c_metric: 0.4150 - lr: 1.0000e-04
Epoch 25/100
93/93 [=====] - ETA: 0s - loss: 119805.1641 - iou_metric: 0.589
4 - p_metric: 0.9560 - c_metric: 0.3583



TRUE LABEL: cat
IOU: 0.5991071428571428
PRED LABEL: dog, C: 0.70300009059906, P: 0.95300009059906
TRUE LABEL: monkey
IOU: 0.7338797975912171
PRED LABEL: dog, C: 0.755999823570251, P: 1.0010000467300415
TRUE LABEL: cat
IOU: 0.38497469832619696
PRED LABEL: dog, C: 0.809000012587891, P: 1.003999948501587
TRUE LABEL: cat
IOU: 0.7072568389057751
PRED LABEL: dog, C: 0.808000028133392, P: 0.995000047683716

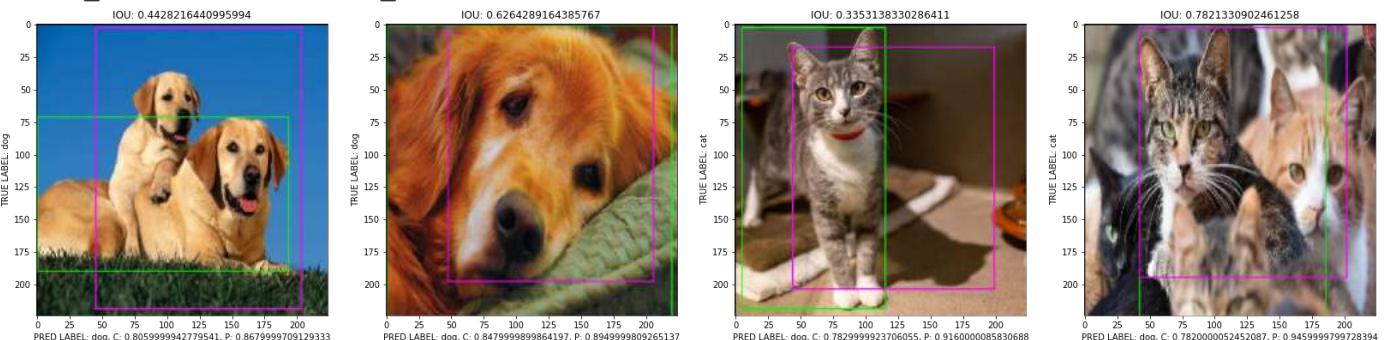
93/93 [=====] - 21s 222ms/step - loss: 119805.1641 - iou_metric: 0.5894 - p_metric: 0.9560 - c_metric: 0.3583 - val_loss: 119347.2266 - val_iou_metric: 0.7120 - val_p_metric: 1.0000 - val_c_metric: 0.4224 - lr: 1.0000e-04
Epoch 26/100

93/93 [=====] - ETA: 0s - loss: 118891.8281 - iou_metric: 0.593
5 - p_metric: 0.9577 - c_metric: 0.3605



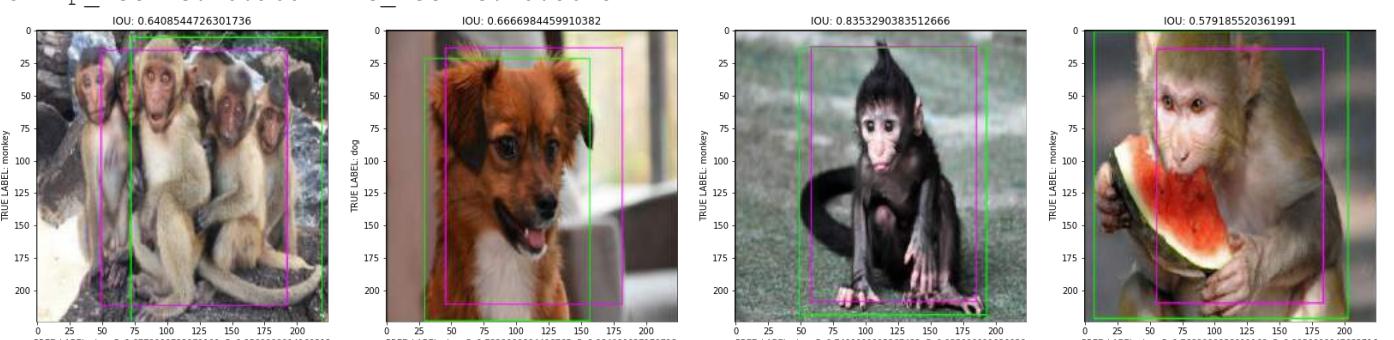
93/93 [=====] - 21s 221ms/step - loss: 118891.8281 - iou_metric: 0.5935 - p_metric: 0.9577 - c_metric: 0.3605 - val_loss: 118422.5000 - val_iou_metric: 0.7137 - val_p_metric: 1.0000 - val_c_metric: 0.4215 - lr: 1.0000e-04
Epoch 27/100

93/93 [=====] - ETA: 0s - loss: 117971.8828 - iou_metric: 0.597
4 - p_metric: 0.9592 - c_metric: 0.3629



93/93 [=====] - 20s 217ms/step - loss: 117971.8828 - iou_metric: 0.5974 - p_metric: 0.9592 - c_metric: 0.3629 - val_loss: 117510.2031 - val_iou_metric: 0.7154 - val_p_metric: 1.0000 - val_c_metric: 0.4193 - lr: 1.0000e-04
Epoch 28/100

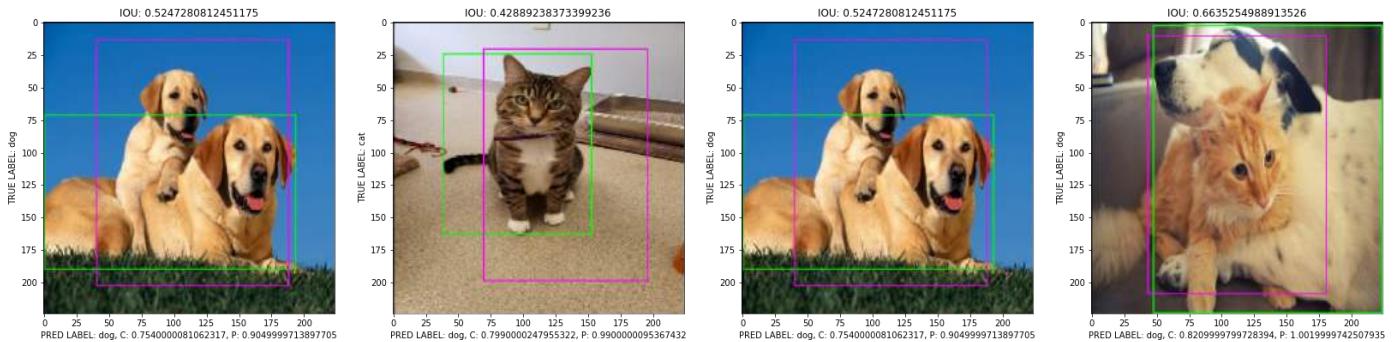
93/93 [=====] - ETA: 0s - loss: 117051.7812 - iou_metric: 0.601
3 - p_metric: 0.9607 - c_metric: 0.3643



93/93 [=====] - 20s 219ms/step - loss: 117051.7812 - iou_metric
c: 0.6013 - p_metric: 0.9607 - c_metric: 0.3643 - val_loss: 116588.5000 - val_iou_metric
c: 0.7163 - val_p_metric: 1.0000 - val_c_metric: 0.4157 - lr: 1.0000e-04

Epoch 29/100

93/93 [=====] - ETA: 0s - loss: 116129.0859 - iou_metric: 0.604
8 - p_metric: 0.9620 - c_metric: 0.3665



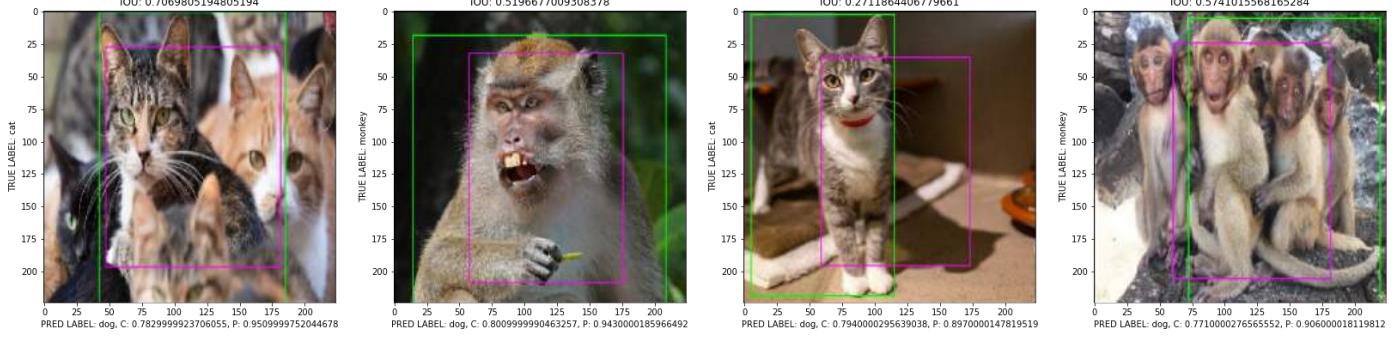
93/93 [=====] - 20s 218ms/step - loss: 116129.0859 - iou_metric

c: 0.6048 - p_metric: 0.9620 - c_metric: 0.3665 - val_loss: 115664.3906 - val_iou_metric

c: 0.7180 - val_p_metric: 1.0000 - val_c_metric: 0.4097 - lr: 1.0000e-04

Epoch 30/100

93/93 [=====] - ETA: 0s - loss: 115215.8203 - iou_metric: 0.608
3 - p_metric: 0.9633 - c_metric: 0.3673



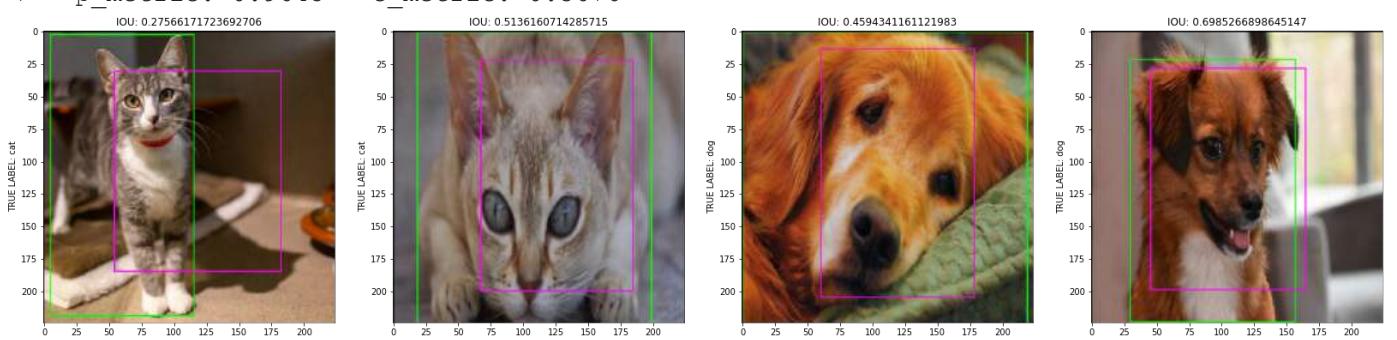
93/93 [=====] - 20s 215ms/step - loss: 115215.8203 - iou_metric

c: 0.6083 - p_metric: 0.9633 - c_metric: 0.3673 - val_loss: 114756.8516 - val_iou_metric

c: 0.7189 - val_p_metric: 1.0000 - val_c_metric: 0.4133 - lr: 1.0000e-04

Epoch 31/100

93/93 [=====] - ETA: 0s - loss: 114304.3672 - iou_metric: 0.611
7 - p_metric: 0.9645 - c_metric: 0.3676



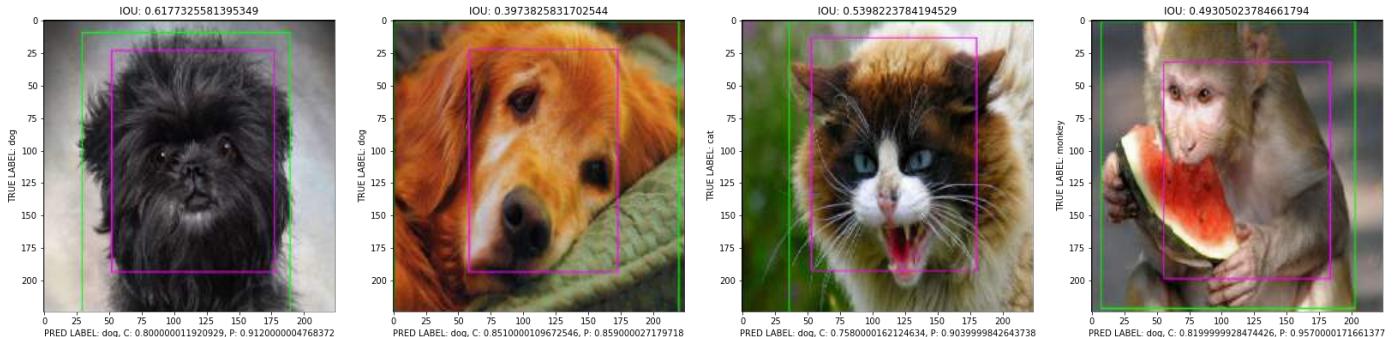
93/93 [=====] - 20s 217ms/step - loss: 114304.3672 - iou_metric

c: 0.6117 - p_metric: 0.9645 - c_metric: 0.3676 - val_loss: 113848.5781 - val_iou_metric

c: 0.7199 - val_p_metric: 1.0000 - val_c_metric: 0.4142 - lr: 1.0000e-04

Epoch 32/100

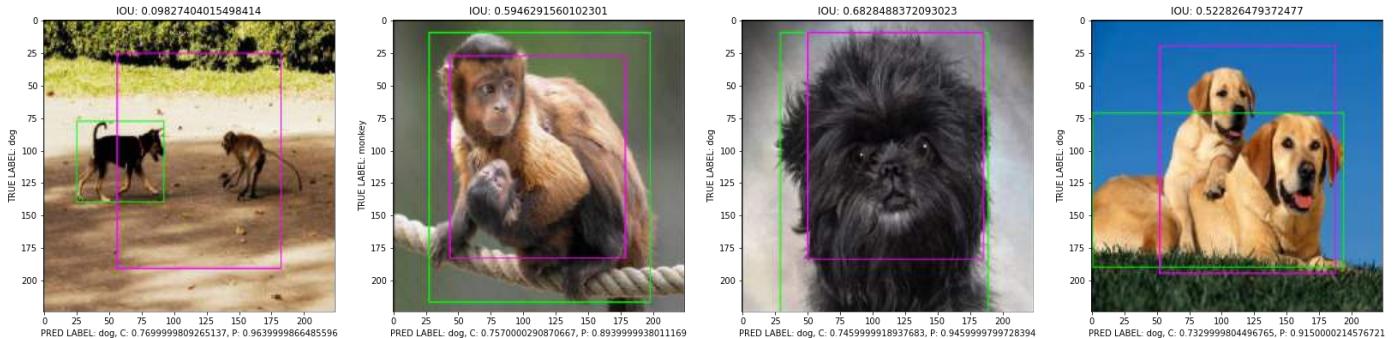
93/93 [=====] - ETA: 0s - loss: 113402.2500 - iou_metric: 0.614
9 - p_metric: 0.9656 - c_metric: 0.3688



93/93 [=====] - 20s 220ms/step - loss: 113402.2500 - iou_metric: 0.6149 - p_metric: 0.9656 - c_metric: 0.3688 - val_loss: 112951.6562 - val_iou_metric: 0.7204 - val_p_metric: 1.0000 - val_c_metric: 0.4163 - lr: 1.0000e-04

Epoch 33/100

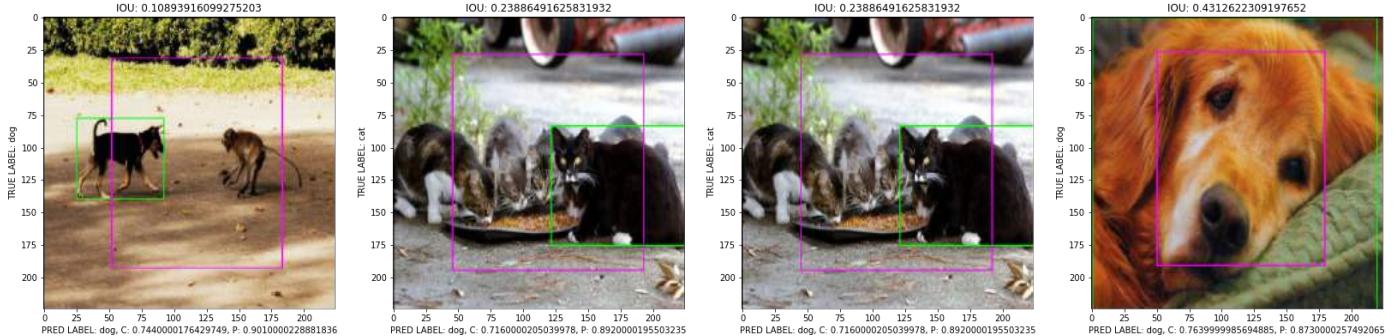
93/93 [=====] - ETA: 0s - loss: 112512.7344 - iou_metric: 0.6179 - p_metric: 0.9666 - c_metric: 0.3698



93/93 [=====] - 20s 218ms/step - loss: 112512.7344 - iou_metric: 0.6179 - p_metric: 0.9666 - c_metric: 0.3698 - val_loss: 112062.7500 - val_iou_metric: 0.7214 - val_p_metric: 1.0000 - val_c_metric: 0.4182 - lr: 1.0000e-04

Epoch 34/100

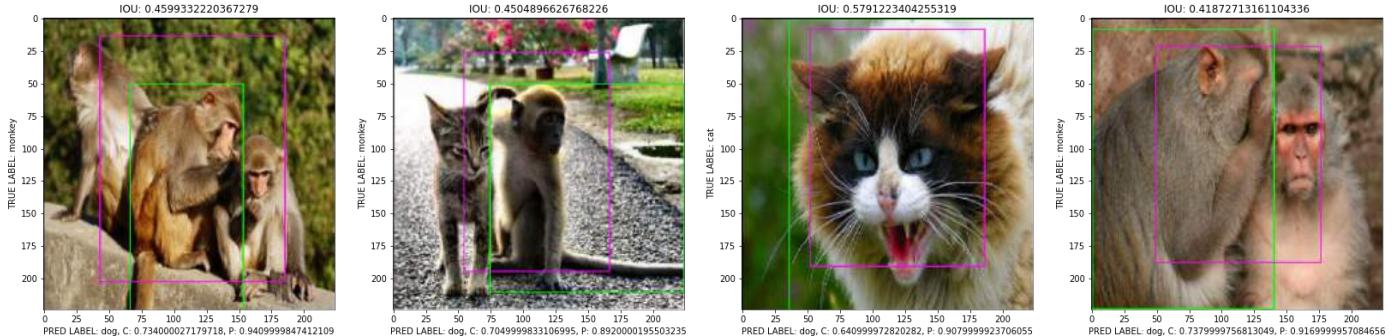
93/93 [=====] - ETA: 0s - loss: 111622.8750 - iou_metric: 0.6207 - p_metric: 0.9676 - c_metric: 0.3704



93/93 [=====] - 21s 221ms/step - loss: 111622.8750 - iou_metric: 0.6207 - p_metric: 0.9676 - c_metric: 0.3704 - val_loss: 111175.3828 - val_iou_metric: 0.7222 - val_p_metric: 1.0000 - val_c_metric: 0.4224 - lr: 1.0000e-04

Epoch 35/100

93/93 [=====] - ETA: 0s - loss: 110721.2812 - iou_metric: 0.6236 - p_metric: 0.9685 - c_metric: 0.3714

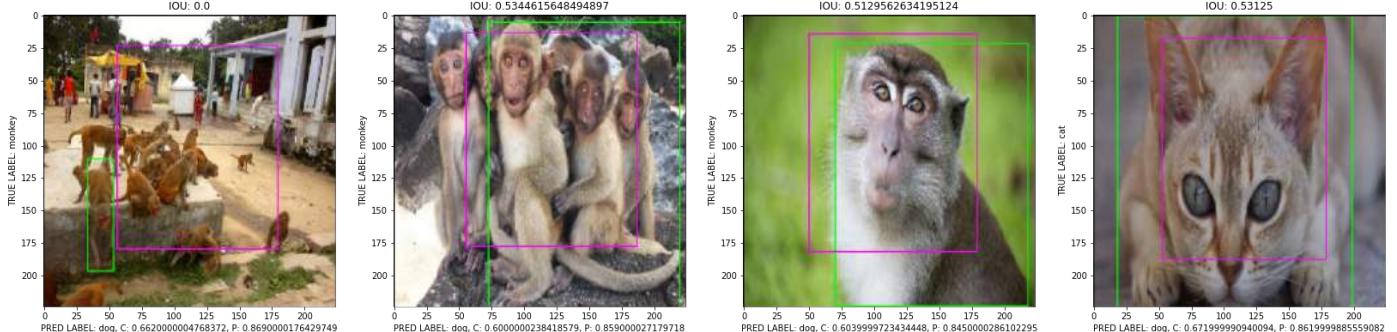


93/93 [=====] - 21s 223ms/step - loss: 110721.2812 - iou_metric: 0.6236 - p_metric: 0.9685 - c_metric: 0.3714 - val_loss: 110259.0234 - val_iou_metric: 0.7223 - val_p_metric: 1.0000 - val_c_metric: 0.4251 - lr: 1.0000e-04

Epoch 36/100

93/93 [=====] - ETA: 0s - loss: 109816.3594 - iou_metric: 0.626

0 - p_metric: 0.9694 - c_metric: 0.3726



93/93 [=====] - 20s 219ms/step - loss: 109816.3594 - iou_metric

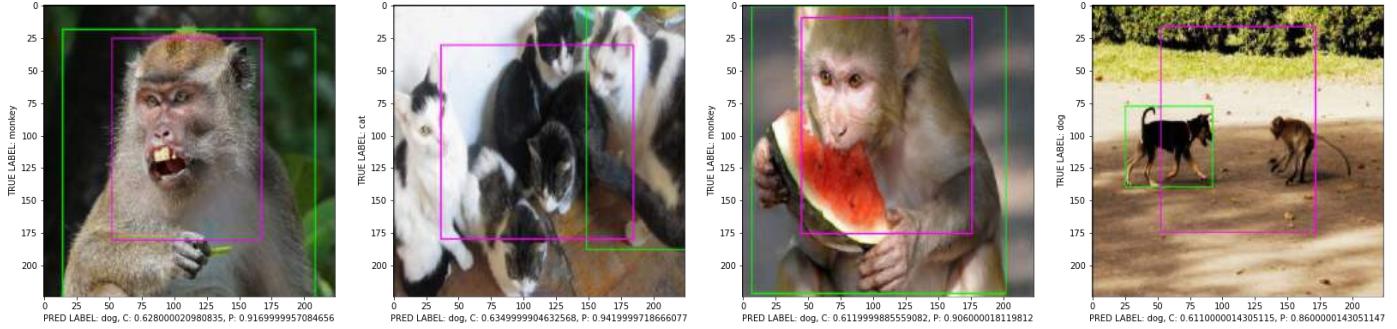
c: 0.6260 - p_metric: 0.9694 - c_metric: 0.3726 - val_loss: 109366.1797 - val_iou_metric

c: 0.7232 - val_p_metric: 1.0000 - val_c_metric: 0.4256 - lr: 1.0000e-04

Epoch 37/100

93/93 [=====] - ETA: 0s - loss: 108923.0000 - iou_metric: 0.628

5 - p_metric: 0.9702 - c_metric: 0.3726



93/93 [=====] - 20s 218ms/step - loss: 108923.0000 - iou_metric

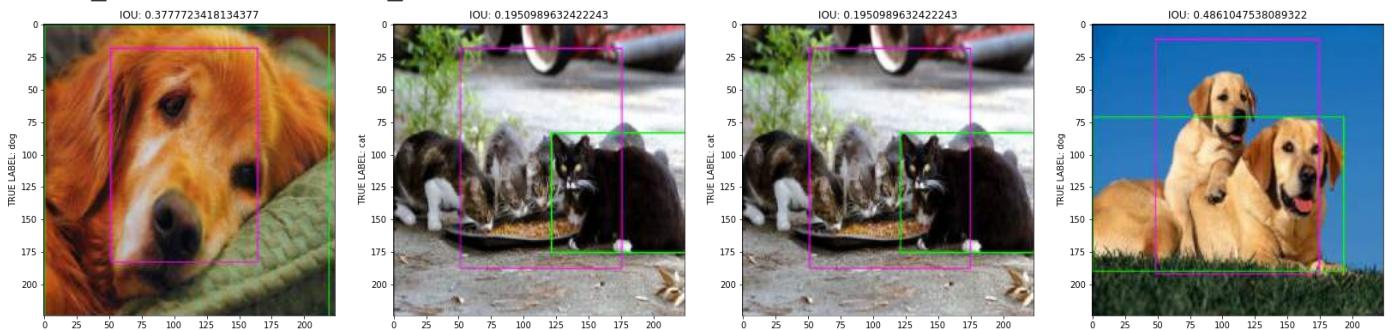
c: 0.6285 - p_metric: 0.9702 - c_metric: 0.3726 - val_loss: 108470.9609 - val_iou_metric

c: 0.7238 - val_p_metric: 1.0000 - val_c_metric: 0.4227 - lr: 1.0000e-04

Epoch 38/100

93/93 [=====] - ETA: 0s - loss: 108021.2188 - iou_metric: 0.630

8 - p_metric: 0.9710 - c_metric: 0.3739



93/93 [=====] - 20s 218ms/step - loss: 108021.2188 - iou_metric

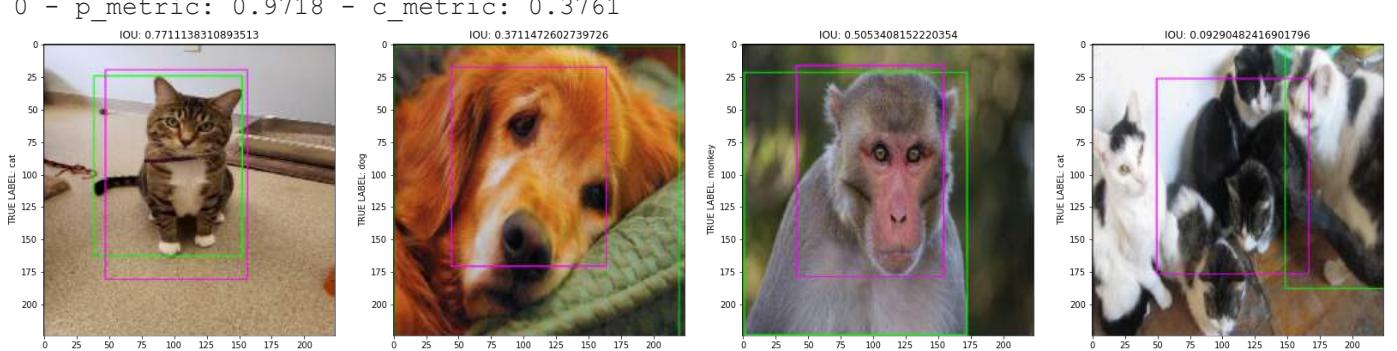
c: 0.6308 - p_metric: 0.9710 - c_metric: 0.3739 - val_loss: 107568.1406 - val_iou_metric

c: 0.7243 - val_p_metric: 1.0000 - val_c_metric: 0.4221 - lr: 1.0000e-04

Epoch 39/100

93/93 [=====] - ETA: 0s - loss: 107133.5391 - iou_metric: 0.633

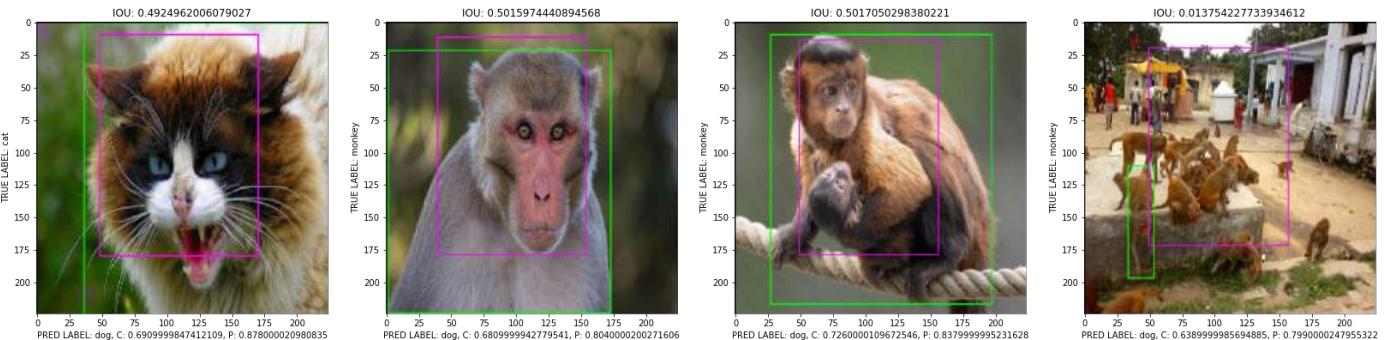
0 - p_metric: 0.9718 - c_metric: 0.3761



93/93 [=====] - 20s 219ms/step - loss: 107133.5391 - iou_metric

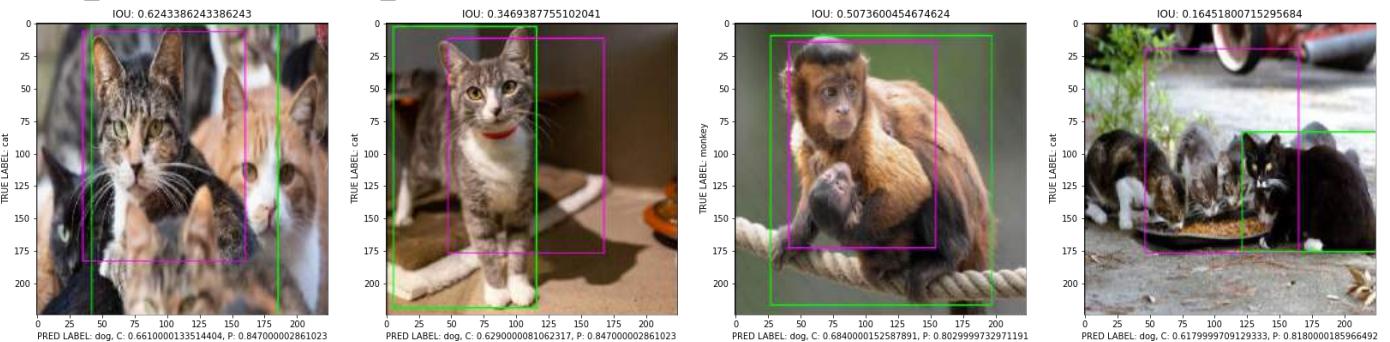
c: 0.6330 - p_metric: 0.9718 - c_metric: 0.3761 - val_loss: 106691.2500 - val_iou_metric

c: 0.7249 - val_p_metric: 1.0000 - val_c_metric: 0.4226 - lr: 1.0000e-04
Epoch 40/100
93/93 [=====] - ETA: 0s - loss: 106266.7266 - iou_metric: 0.635
2 - p_metric: 0.9725 - c_metric: 0.3768



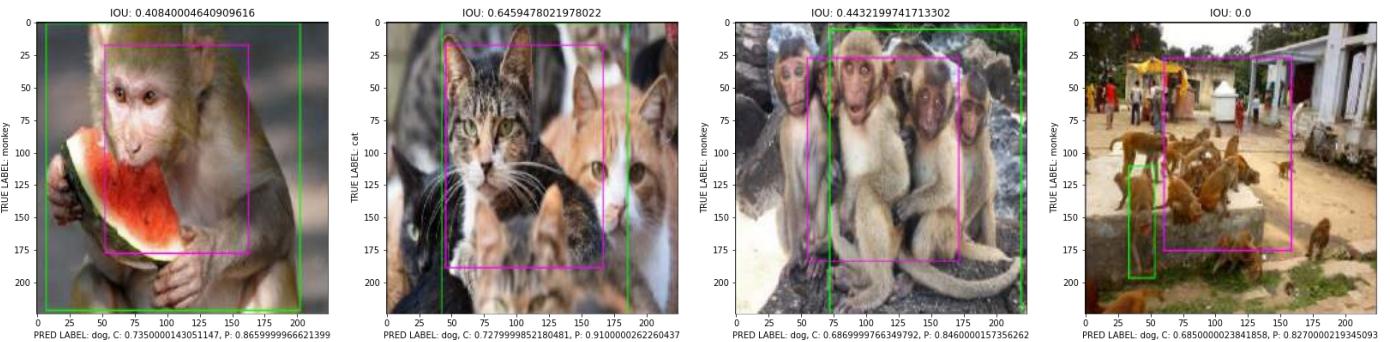
93/93 [=====] - 20s 216ms/step - loss: 106266.7266 - iou_metric: 0.6352 - p_metric: 0.9725 - c_metric: 0.3768 - val_loss: 105830.9297 - val_iou_metric: 0.7262 - val_p_metric: 1.0000 - val_c_metric: 0.4220 - lr: 1.0000e-04
Epoch 41/100

93/93 [=====] - ETA: 0s - loss: 105395.3125 - iou_metric: 0.637
3 - p_metric: 0.9731 - c_metric: 0.3776



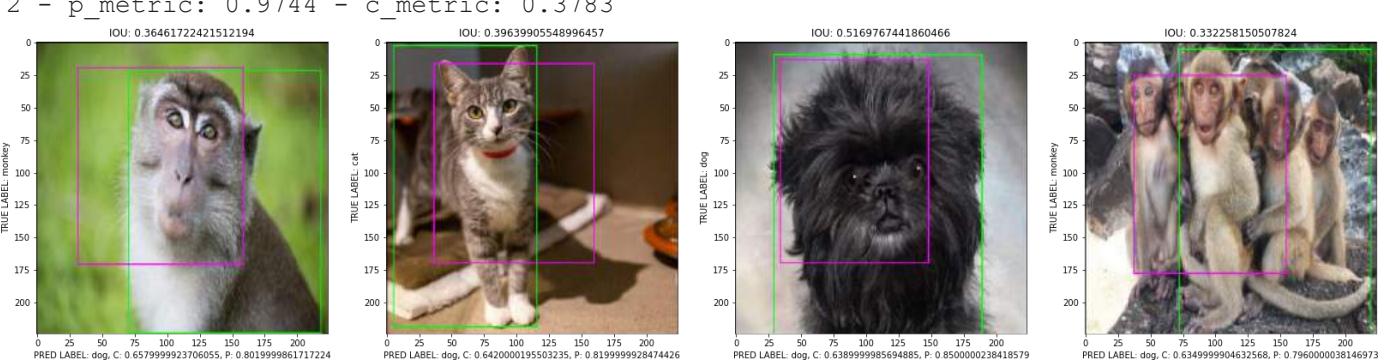
93/93 [=====] - 20s 218ms/step - loss: 105395.3125 - iou_metric: 0.6373 - p_metric: 0.9731 - c_metric: 0.3776 - val_loss: 104954.9922 - val_iou_metric: 0.7266 - val_p_metric: 1.0000 - val_c_metric: 0.4244 - lr: 1.0000e-04
Epoch 42/100

93/93 [=====] - ETA: 0s - loss: 104533.1250 - iou_metric: 0.639
3 - p_metric: 0.9738 - c_metric: 0.3775



93/93 [=====] - 20s 218ms/step - loss: 104533.1250 - iou_metric: 0.6393 - p_metric: 0.9738 - c_metric: 0.3775 - val_loss: 104099.8203 - val_iou_metric: 0.7270 - val_p_metric: 1.0000 - val_c_metric: 0.4219 - lr: 1.0000e-04
Epoch 43/100

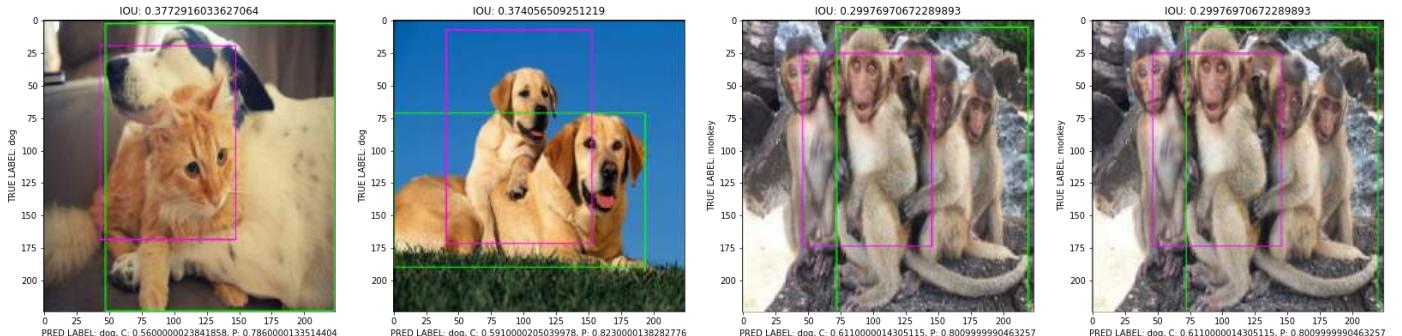
93/93 [=====] - ETA: 0s - loss: 103682.8672 - iou_metric: 0.641
2 - p_metric: 0.9744 - c_metric: 0.3783



93/93 [=====] - 20s 217ms/step - loss: 103682.8672 - iou_metric
c: 0.6412 - p_metric: 0.9744 - c_metric: 0.3783 - val_loss: 103253.1484 - val_iou_metric
c: 0.7278 - val_p_metric: 1.0000 - val_c_metric: 0.4223 - lr: 1.0000e-04

Epoch 44/100

93/93 [=====] - ETA: 0s - loss: 102830.0547 - iou_metric: 0.643
1 - p_metric: 0.9750 - c_metric: 0.3795

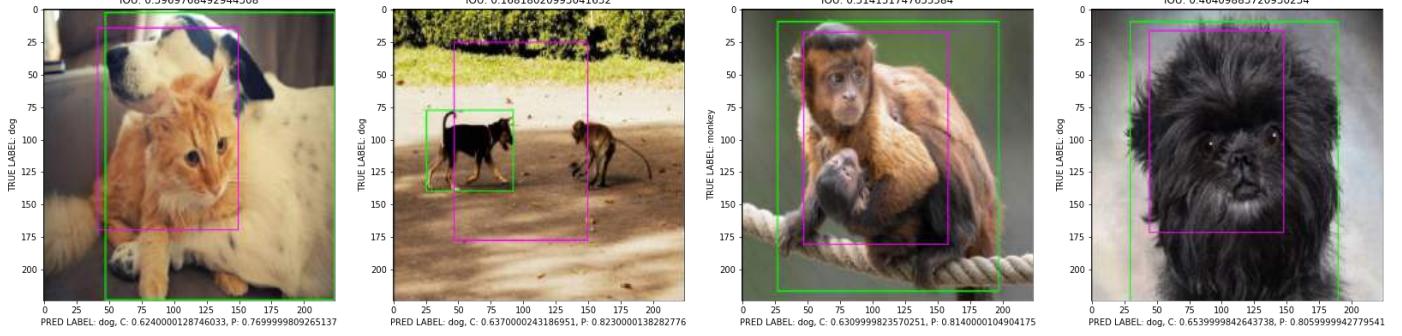


93/93 [=====] - 20s 218ms/step - loss: 102830.0547 - iou_metric
c: 0.6431 - p_metric: 0.9750 - c_metric: 0.3795 - val_loss: 102398.3516 - val_iou_metric

c: 0.7285 - val_p_metric: 1.0000 - val_c_metric: 0.4218 - lr: 1.0000e-04

Epoch 45/100

93/93 [=====] - ETA: 0s - loss: 101984.8750 - iou_metric: 0.644
8 - p_metric: 0.9755 - c_metric: 0.3803

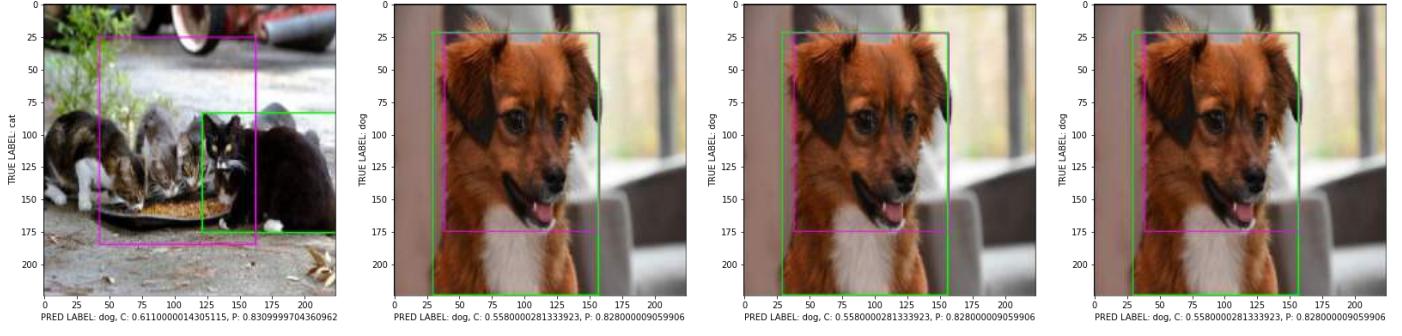


93/93 [=====] - 20s 219ms/step - loss: 101984.8750 - iou_metric
c: 0.6448 - p_metric: 0.9755 - c_metric: 0.3803 - val_loss: 101578.8516 - val_iou_metric

c: 0.7292 - val_p_metric: 1.0000 - val_c_metric: 0.4204 - lr: 1.0000e-04

Epoch 46/100

93/93 [=====] - ETA: 0s - loss: 101168.4219 - iou_metric: 0.646
5 - p_metric: 0.9761 - c_metric: 0.3809

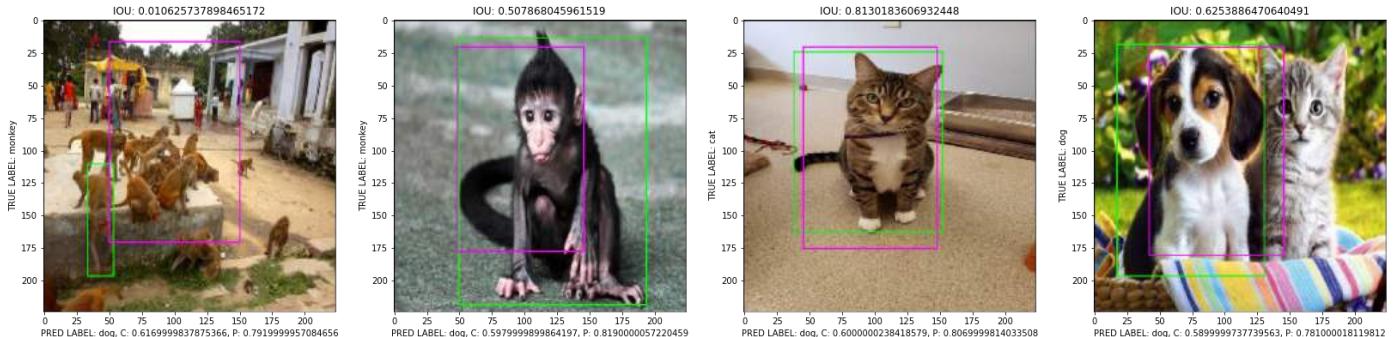


93/93 [=====] - 20s 218ms/step - loss: 101168.4219 - iou_metric
c: 0.6465 - p_metric: 0.9761 - c_metric: 0.3809 - val_loss: 100742.0312 - val_iou_metric

c: 0.7298 - val_p_metric: 1.0000 - val_c_metric: 0.4243 - lr: 1.0000e-04

Epoch 47/100

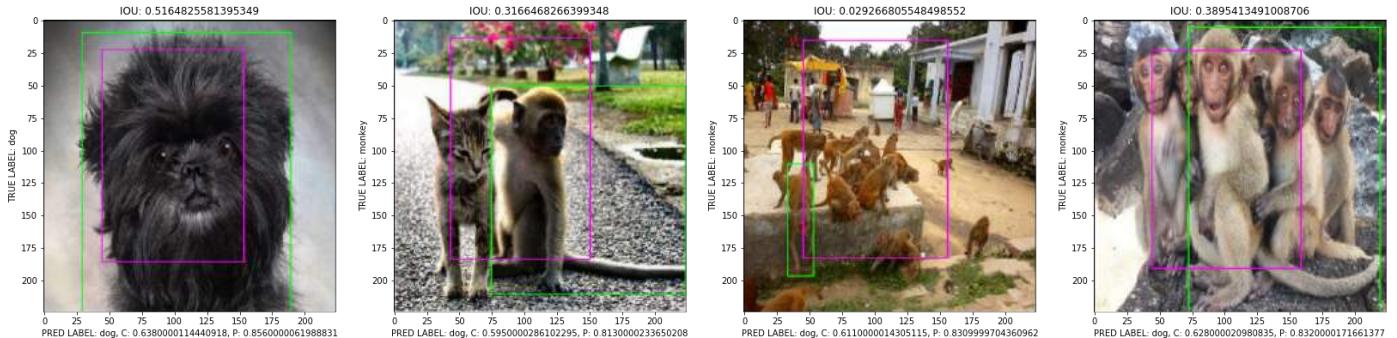
93/93 [=====] - ETA: 0s - loss: 100331.7188 - iou_metric: 0.648
1 - p_metric: 0.9766 - c_metric: 0.3819



93/93 [=====] - 20s 219ms/step - loss: 100331.7188 - iou_metric: 0.6481 - p_metric: 0.9766 - c_metric: 0.3819 - val_loss: 99913.2891 - val_iou_metric: 0.7306 - val_p_metric: 1.0000 - val_c_metric: 0.4264 - lr: 1.0000e-04

Epoch 48/100

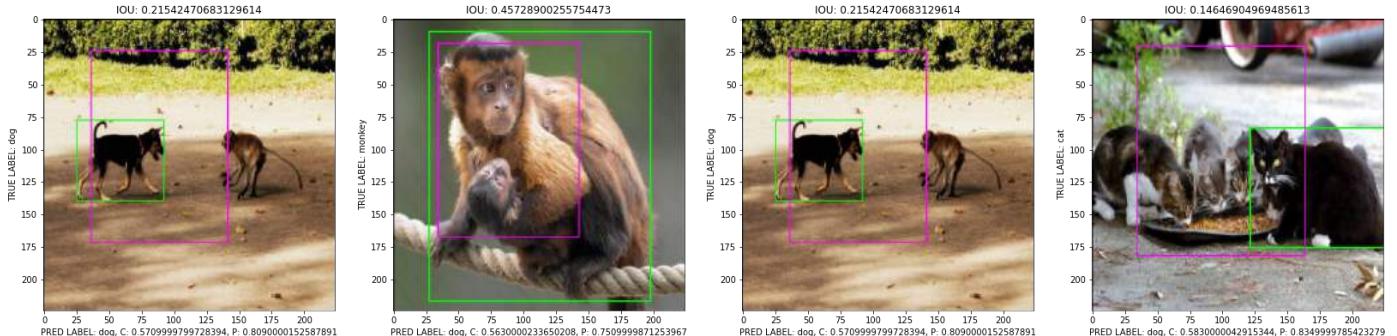
93/93 [=====] - ETA: 0s - loss: 99494.4766 - iou_metric: 0.6496 - p_metric: 0.9771 - c_metric: 0.3820



93/93 [=====] - 20s 218ms/step - loss: 99494.4766 - iou_metric: 0.6496 - p_metric: 0.9771 - c_metric: 0.3820 - val_loss: 99068.9062 - val_iou_metric: 0.7312 - val_p_metric: 1.0000 - val_c_metric: 0.4283 - lr: 1.0000e-04

Epoch 49/100

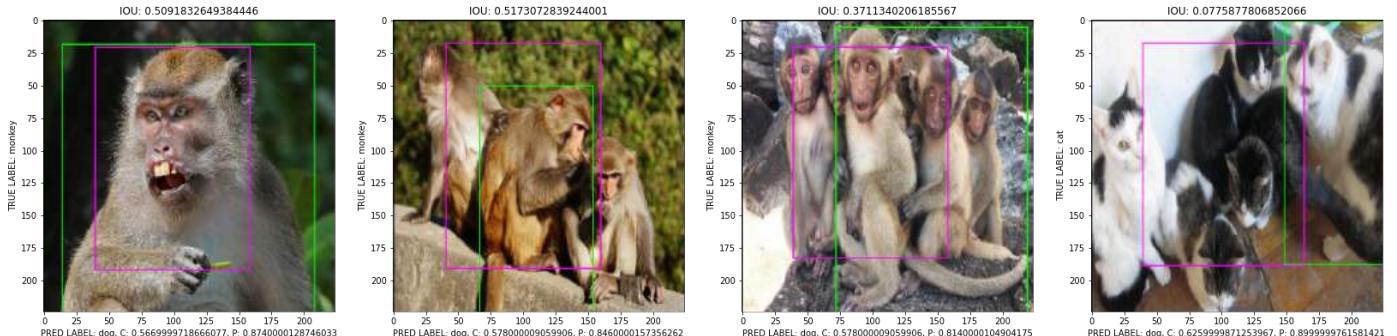
93/93 [=====] - ETA: 0s - loss: 98658.8125 - iou_metric: 0.6513 - p_metric: 0.9775 - c_metric: 0.3831



93/93 [=====] - 20s 217ms/step - loss: 98658.8125 - iou_metric: 0.6513 - p_metric: 0.9775 - c_metric: 0.3831 - val_loss: 98237.2812 - val_iou_metric: 0.7320 - val_p_metric: 1.0000 - val_c_metric: 0.4302 - lr: 1.0000e-04

Epoch 50/100

93/93 [=====] - ETA: 0s - loss: 97829.7500 - iou_metric: 0.6528 - p_metric: 0.9780 - c_metric: 0.3847

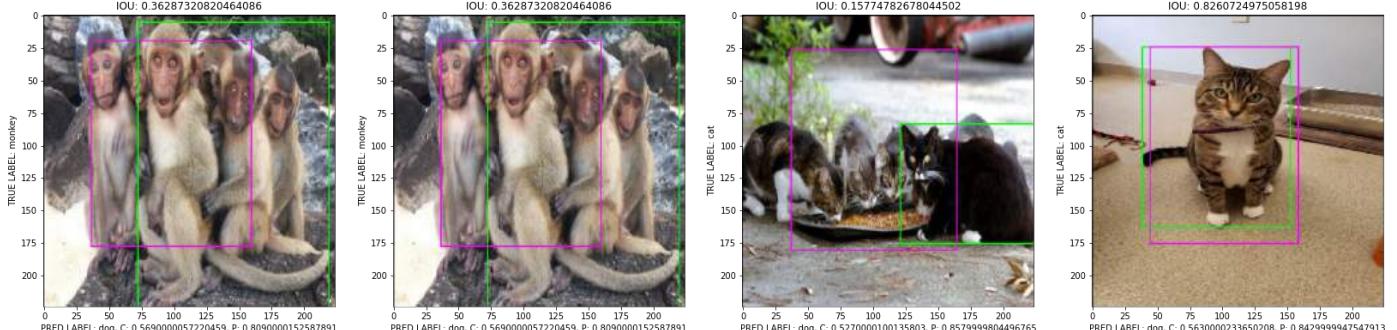


93/93 [=====] - 20s 216ms/step - loss: 97829.7500 - iou_metric: 0.6528 - p_metric: 0.9780 - c_metric: 0.3847 - val_loss: 97410.7500 - val_iou_metric: 0.7327 - val_p_metric: 1.0000 - val_c_metric: 0.4304 - lr: 1.0000e-04

Epoch 51/100

93/93 [=====] - ETA: 0s - loss: 97003.3906 - iou_metric: 0.6542

- p_metric: 0.9784 - c_metric: 0.3851



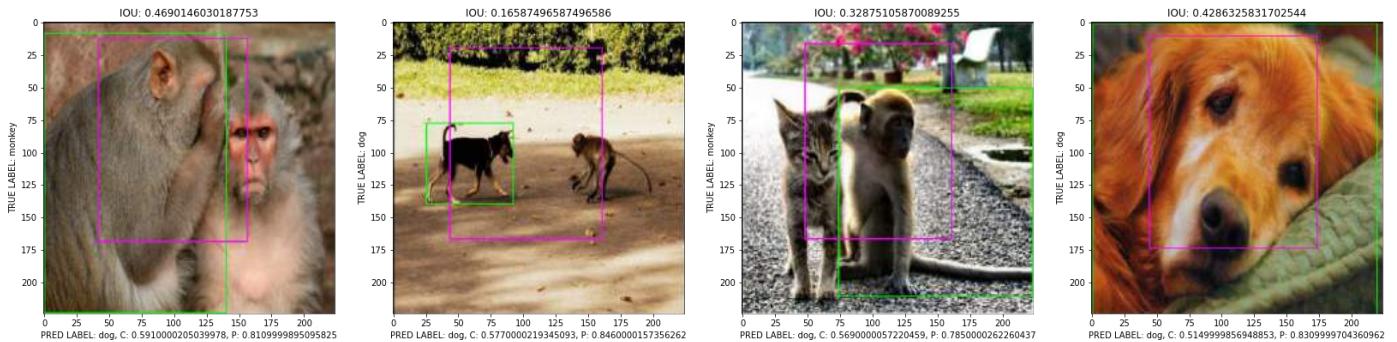
93/93 [=====] - 20s 215ms/step - loss: 97003.3906 - iou_metric:

0.6542 - p_metric: 0.9784 - c_metric: 0.3851 - val_loss: 96582.3438 - val_iou_metric: 0.7333 - val_p_metric: 1.0000 - val_c_metric: 0.4290 - lr: 1.0000e-04

Epoch 52/100

93/93 [=====] - ETA: 0s - loss: 96173.3828 - iou_metric: 0.6557

- p_metric: 0.9788 - c_metric: 0.3851



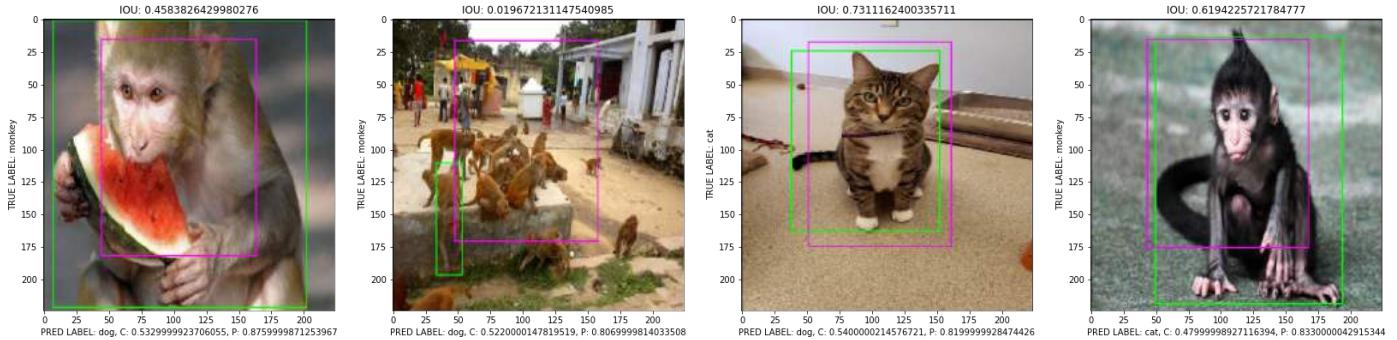
93/93 [=====] - 20s 216ms/step - loss: 96173.3828 - iou_metric:

0.6557 - p_metric: 0.9788 - c_metric: 0.3851 - val_loss: 95755.6406 - val_iou_metric: 0.7339 - val_p_metric: 1.0000 - val_c_metric: 0.4308 - lr: 1.0000e-04

Epoch 53/100

93/93 [=====] - ETA: 0s - loss: 95359.4141 - iou_metric: 0.6570

- p_metric: 0.9792 - c_metric: 0.3858



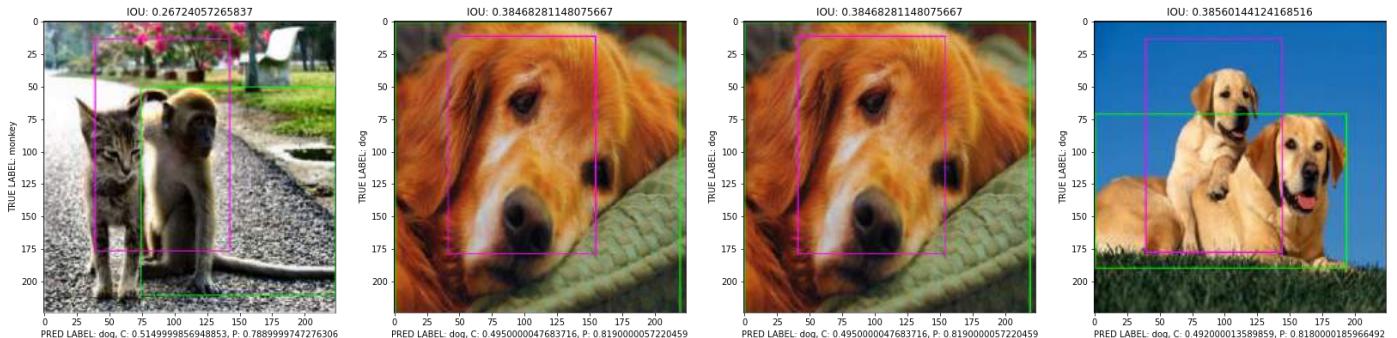
93/93 [=====] - 20s 218ms/step - loss: 95359.4141 - iou_metric:

0.6570 - p_metric: 0.9792 - c_metric: 0.3858 - val_loss: 94957.9375 - val_iou_metric: 0.7345 - val_p_metric: 1.0000 - val_c_metric: 0.4317 - lr: 1.0000e-04

Epoch 54/100

93/93 [=====] - ETA: 0s - loss: 94554.4844 - iou_metric: 0.6584

- p_metric: 0.9796 - c_metric: 0.3858



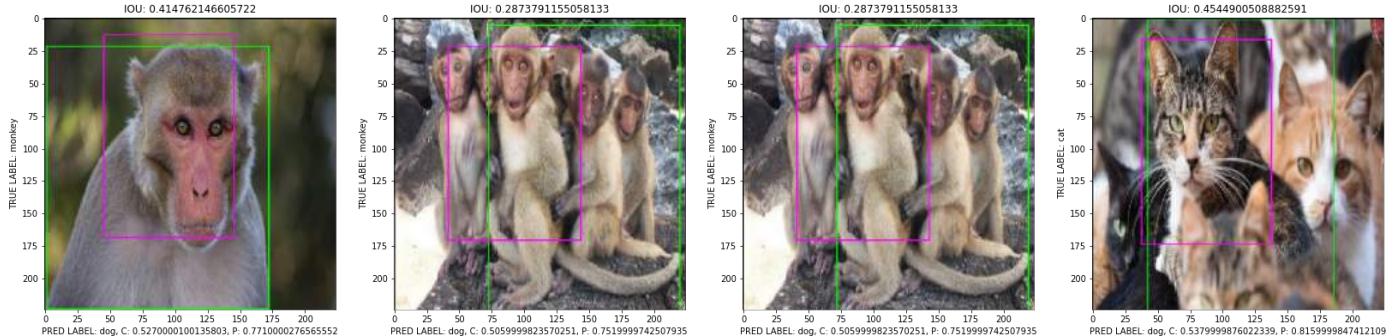
93/93 [=====] - 20s 219ms/step - loss: 94554.4844 - iou_metric:

0.6584 - p_metric: 0.9796 - c_metric: 0.3858 - val_loss: 94147.6484 - val_iou_metric: 0.

7348 - val_p_metric: 1.0000 - val_c_metric: 0.4356 - lr: 1.0000e-04

Epoch 55/100

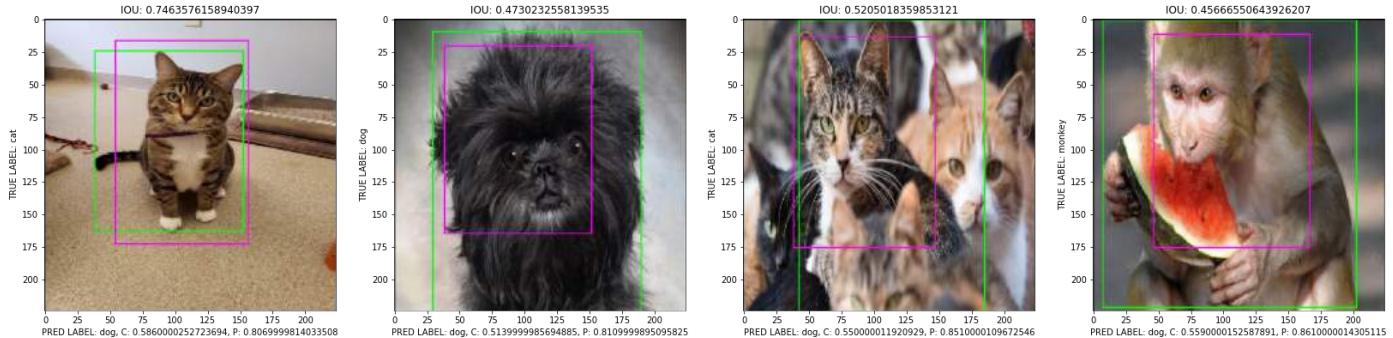
93/93 [=====] - ETA: 0s - loss: 93752.5234 - iou_metric: 0.6596 - p_metric: 0.9800 - c_metric: 0.3865



93/93 [=====] - 20s 217ms/step - loss: 93752.5234 - iou_metric: 0.6596 - p_metric: 0.9800 - c_metric: 0.3865 - val_loss: 93351.7578 - val_iou_metric: 0.7353 - val_p_metric: 1.0000 - val_c_metric: 0.4356 - lr: 1.0000e-04

Epoch 56/100

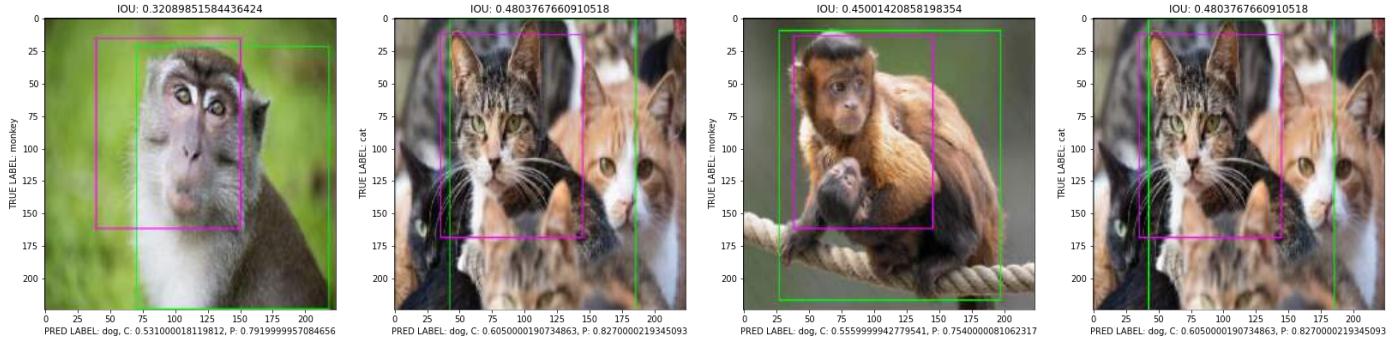
93/93 [=====] - ETA: 0s - loss: 92951.1562 - iou_metric: 0.6608 - p_metric: 0.9803 - c_metric: 0.3871



93/93 [=====] - 20s 218ms/step - loss: 92951.1562 - iou_metric: 0.6608 - p_metric: 0.9803 - c_metric: 0.3871 - val_loss: 92542.7031 - val_iou_metric: 0.7357 - val_p_metric: 1.0000 - val_c_metric: 0.4379 - lr: 1.0000e-04

Epoch 57/100

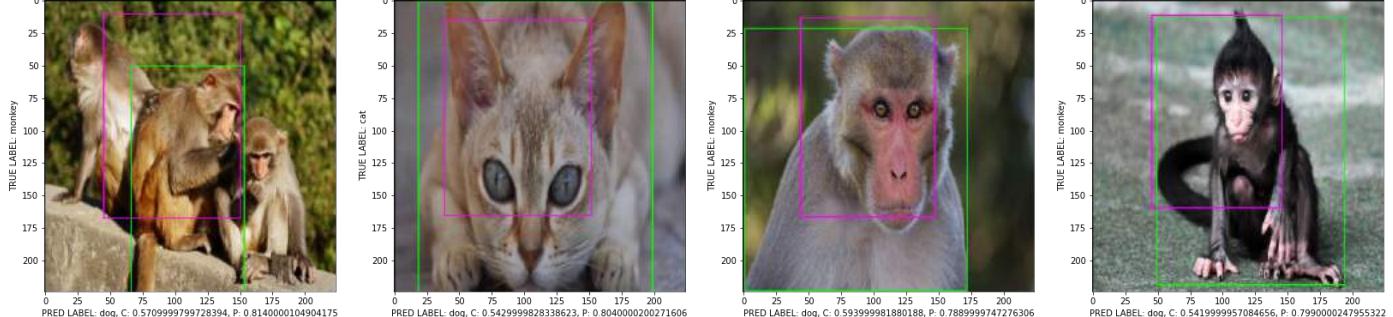
93/93 [=====] - ETA: 0s - loss: 92151.8984 - iou_metric: 0.6619 - p_metric: 0.9807 - c_metric: 0.3877



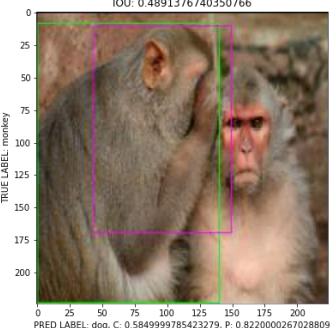
93/93 [=====] - 21s 223ms/step - loss: 92151.8984 - iou_metric: 0.6619 - p_metric: 0.9807 - c_metric: 0.3877 - val_loss: 91753.3984 - val_iou_metric: 0.7362 - val_p_metric: 1.0000 - val_c_metric: 0.4386 - lr: 1.0000e-04

Epoch 58/100

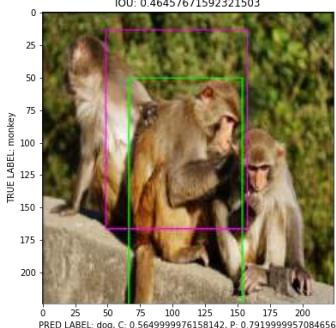
93/93 [=====] - ETA: 0s - loss: 91363.4062 - iou_metric: 0.6631 - p_metric: 0.9810 - c_metric: 0.3886



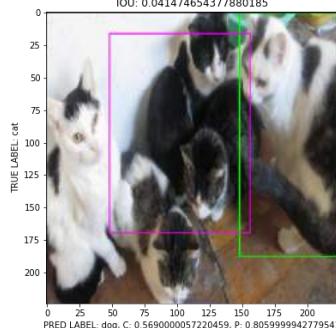
93/93 [=====] - 20s 219ms/step - loss: 91363.4062 - iou_metric: 0.6631 - p_metric: 0.9810 - c_metric: 0.3886 - val_loss: 90960.3594 - val_iou_metric: 0.7365 - val_p_metric: 1.0000 - val_c_metric: 0.4421 - lr: 1.0000e-04
 Epoch 59/100
 93/93 [=====] - ETA: 0s - loss: 90578.1641 - iou_metric: 0.6641 - p_metric: 0.9813 - c_metric: 0.3889



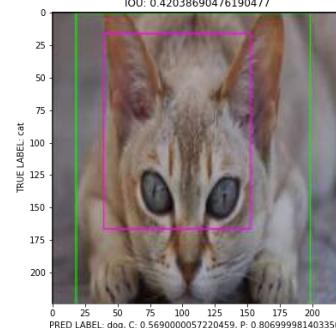
TRUE LABEL: monkey
PRED LABEL: dog, C: 0.5849999785423279, P: 0.8220000267028809



TRUE LABEL: monkey
PRED LABEL: dog, C: 0.564999976158142, P: 0.7919999570846541

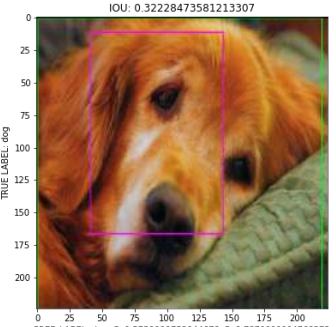


TRUE LABEL: cat
PRED LABEL: dog, C: 0.5690000057220459, P: 0.805999942779541

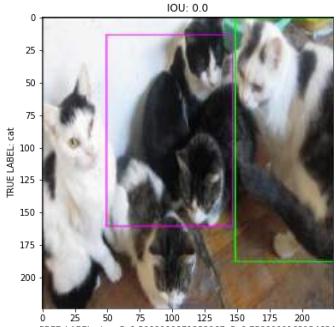


TRUE LABEL: cat
PRED LABEL: dog, C: 0.5690000057220459, P: 0.8069999814033508

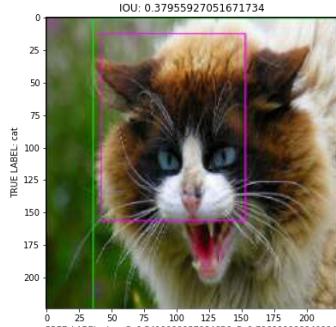
93/93 [=====] - 21s 220ms/step - loss: 90578.1641 - iou_metric: 0.6641 - p_metric: 0.9813 - c_metric: 0.3889 - val_loss: 90192.3125 - val_iou_metric: 0.7370 - val_p_metric: 1.0000 - val_c_metric: 0.4407 - lr: 1.0000e-04
 Epoch 60/100
 93/93 [=====] - ETA: 0s - loss: 89801.3047 - iou_metric: 0.6652 - p_metric: 0.9816 - c_metric: 0.3896



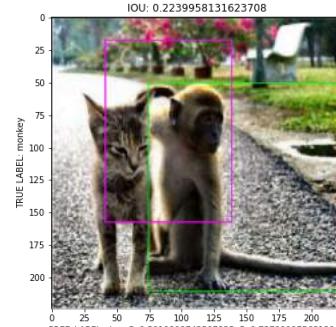
TRUE LABEL: dog
PRED LABEL: dog, C: 0.5759999752044678, P: 0.787000004768372



TRUE LABEL: cat
PRED LABEL: dog, C: 0.500999871253967, P: 0.7580000162124634

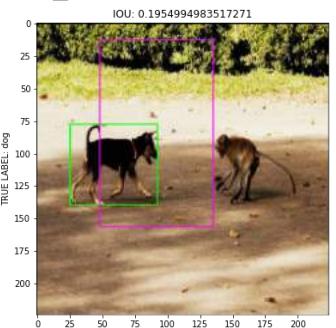


TRUE LABEL: cat
PRED LABEL: dog, C: 0.541999957084656, P: 0.796999990440094

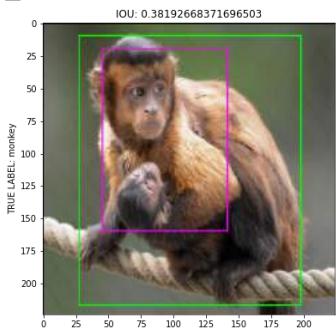


TRUE LABEL: monkey
PRED LABEL: dog, C: 0.5019999742507935, P: 0.7379999756813049

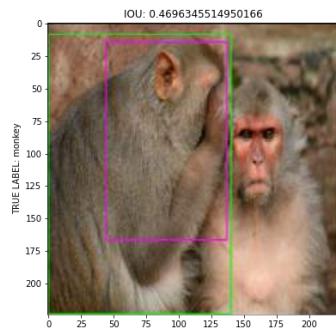
93/93 [=====] - 20s 216ms/step - loss: 89801.3047 - iou_metric: 0.6652 - p_metric: 0.9816 - c_metric: 0.3896 - val_loss: 89399.6875 - val_iou_metric: 0.7374 - val_p_metric: 1.0000 - val_c_metric: 0.4400 - lr: 1.0000e-04
 Epoch 61/100
 93/93 [=====] - ETA: 0s - loss: 89011.0938 - iou_metric: 0.6662 - p_metric: 0.9819 - c_metric: 0.3898



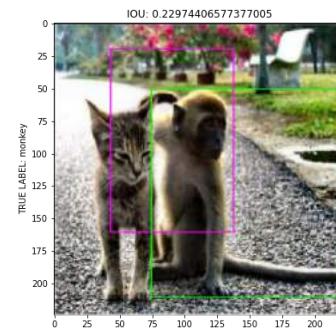
TRUE LABEL: dog
PRED LABEL: dog, C: 0.5379999876022339, P: 0.7459999918937683



TRUE LABEL: monkey
PRED LABEL: dog, C: 0.5910000205039978, P: 0.7519999742507935

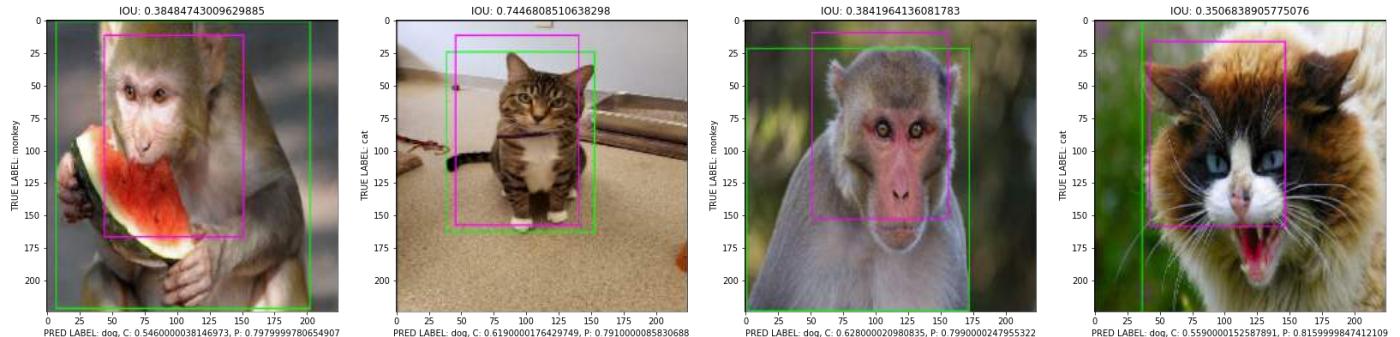


TRUE LABEL: monkey
PRED LABEL: dog, C: 0.601999980926514, P: 0.787000004768372



TRUE LABEL: monkey
PRED LABEL: dog, C: 0.5759999752044678, P: 0.753000020980835

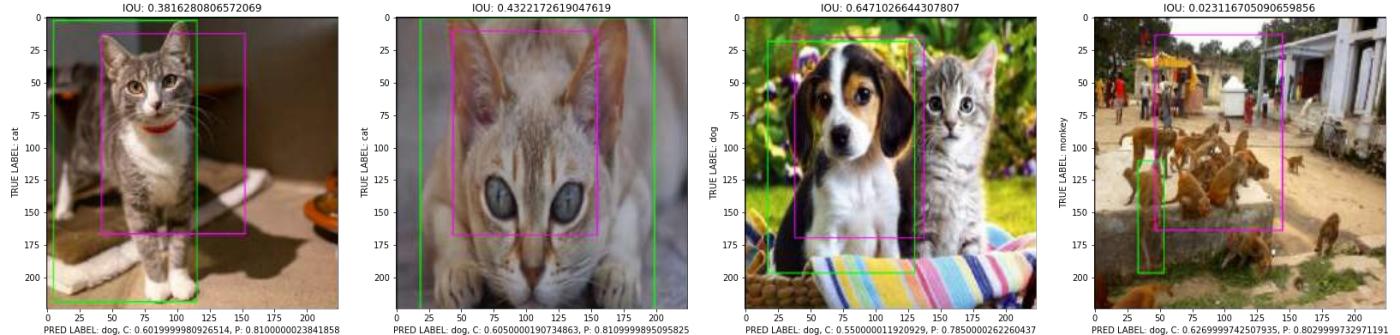
93/93 [=====] - 21s 221ms/step - loss: 89011.0938 - iou_metric: 0.6662 - p_metric: 0.9819 - c_metric: 0.3898 - val_loss: 88615.4219 - val_iou_metric: 0.7377 - val_p_metric: 1.0000 - val_c_metric: 0.4413 - lr: 1.0000e-04
 Epoch 62/100
 93/93 [=====] - ETA: 0s - loss: 88232.1719 - iou_metric: 0.6672 - p_metric: 0.9822 - c_metric: 0.3899



93/93 [=====] - 20s 217ms/step - loss: 88232.1719 - iou_metric: 0.6672 - p_metric: 0.9822 - c_metric: 0.3899 - val_loss: 87843.5391 - val_iou_metric: 0.7381 - val_p_metric: 1.0000 - val_c_metric: 0.4419 - lr: 1.0000e-04

Epoch 63/100

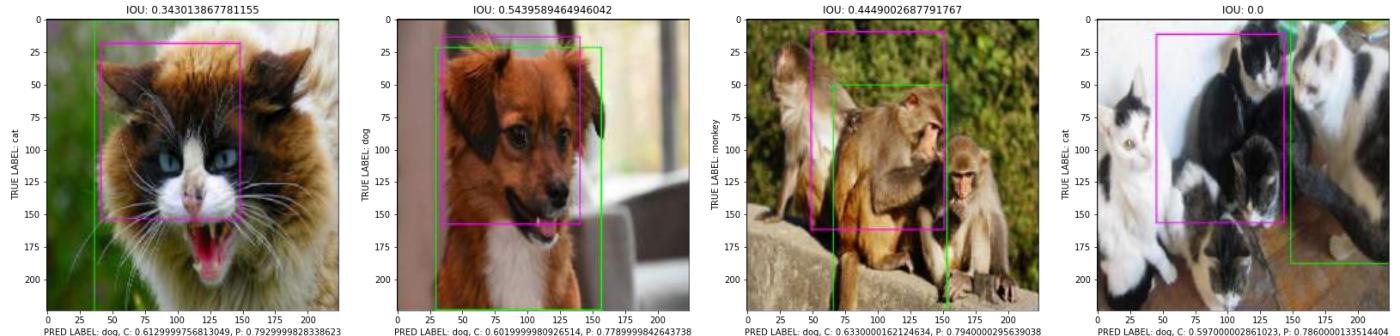
93/93 [=====] - ETA: 0s - loss: 87463.5859 - iou_metric: 0.6682 - p_metric: 0.9825 - c_metric: 0.3907



93/93 [=====] - 20s 219ms/step - loss: 87463.5859 - iou_metric: 0.6682 - p_metric: 0.9825 - c_metric: 0.3907 - val_loss: 87073.4688 - val_iou_metric: 0.7385 - val_p_metric: 1.0000 - val_c_metric: 0.4419 - lr: 1.0000e-04

Epoch 64/100

93/93 [=====] - ETA: 0s - loss: 86688.3359 - iou_metric: 0.6691 - p_metric: 0.9828 - c_metric: 0.3911



93/93 [=====] - 20s 216ms/step - loss: 86688.3359 - iou_metric: 0.6691 - p_metric: 0.9828 - c_metric: 0.3911 - val_loss: 86298.7266 - val_iou_metric: 0.7388 - val_p_metric: 1.0000 - val_c_metric: 0.4413 - lr: 1.0000e-04

Epoch 65/100

93/93 [=====] - ETA: 0s - loss: 85925.8750 - iou_metric: 0.6701 - p_metric: 0.9831 - c_metric: 0.3919



93/93 [=====] - 20s 217ms/step - loss: 85925.8750 - iou_metric: 0.6701 - p_metric: 0.9831 - c_metric: 0.3919 - val_loss: 85541.4375 - val_iou_metric: 0.7391 - val_p_metric: 1.0000 - val_c_metric: 0.4412 - lr: 1.0000e-04

Epoch 66/100



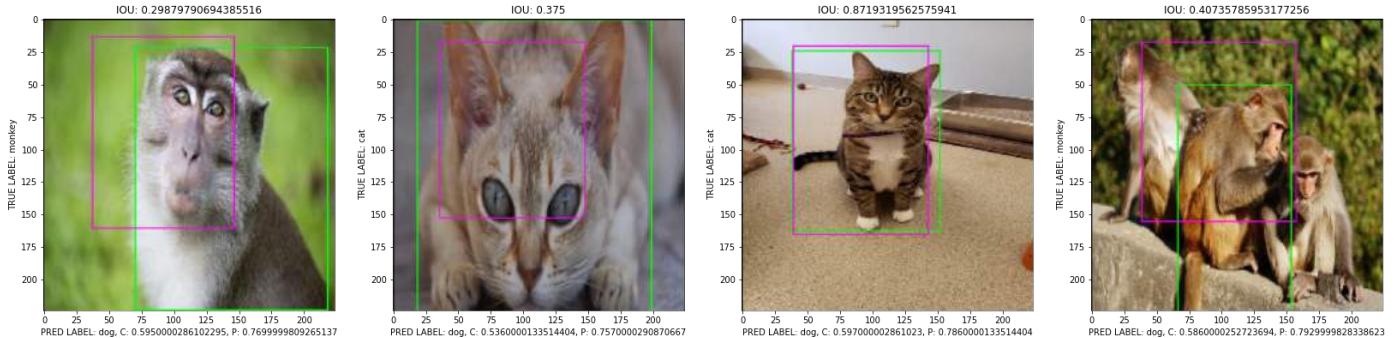
93/93 [=====] - 20s 217ms/step - loss: 85166.2969 - iou_metric: 0.6710 - p_metric: 0.9833 - c_metric: 0.3921 - val_loss: 84776.0625 - val_iou_metric: 0.7396 - val_p_metric: 1.0000 - val_c_metric: 0.4412 - lr: 1.0000e-04
Epoch 67/100

93/93 [=====] - ETA: 0s - loss: 84403.5000 - iou_metric: 0.6719
- p_metric: 0.9836 - c_metric: 0.3925



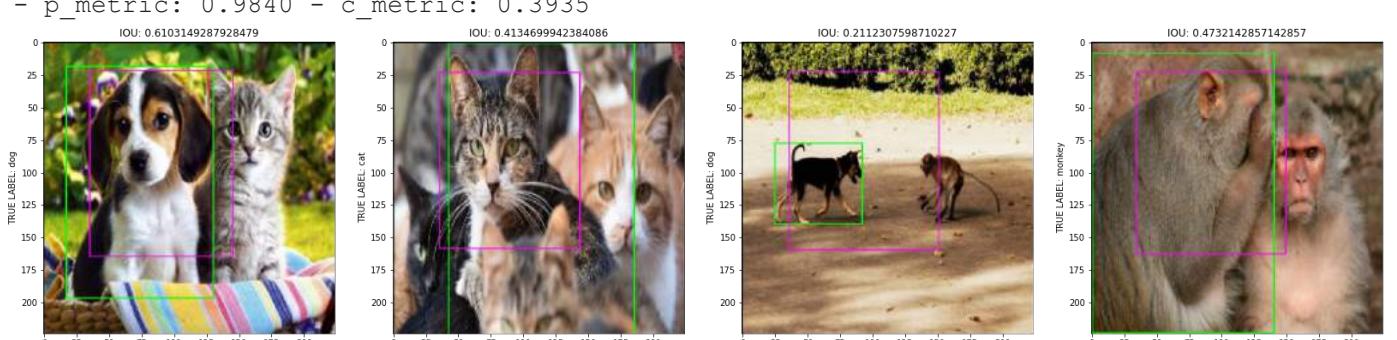
93/93 [=====] - 20s 216ms/step - loss: 84403.5000 - iou_metric: 0.6719 - p_metric: 0.9836 - c_metric: 0.3925 - val_loss: 84023.9375 - val_iou_metric: 0.7399 - val_p_metric: 1.0000 - val_c_metric: 0.4412 - lr: 1.0000e-04
Epoch 68/100

93/93 [=====] - ETA: 0s - loss: 83659.8281 - iou_metric: 0.6728 - p_metric: 0.9838 - c_metric: 0.3934



93/93 [=====] - 21s 222ms/step - loss: 83659.8281 - iou_metric: 0.6728 - p_metric: 0.9838 - c_metric: 0.3934 - val_loss: 83284.7031 - val_iou_metric: 0.7404 - val_p_metric: 1.0000 - val_c_metric: 0.4418 - lr: 1.0000e-04
Epoch 69/100

93/93 [=====] - ETA: 0s - loss: 82910.3750 - iou_metric: 0.6737 - p_metric: 0.9840 - c_metric: 0.3935

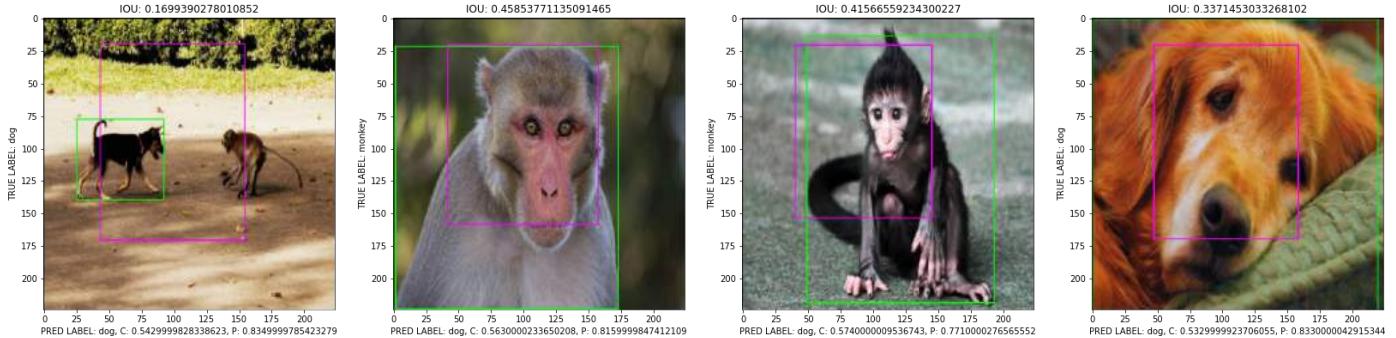


93/93 [=====] - 19s 206ms/step - loss: 82910.3750 - iou_metric: 0.6737 - p_metric: 0.9840 - c_metric: 0.3935 - val_loss: 82530.4375 - val_iou_metric: 0.

7410 - val_p_metric: 1.0000 - val_c_metric: 0.4400 - lr: 1.0000e-04

Epoch 70/100

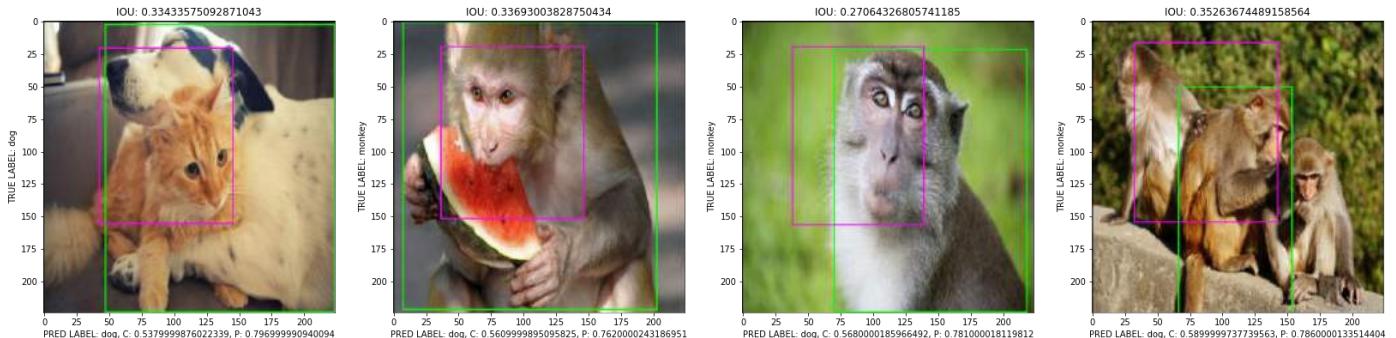
93/93 [=====] - ETA: 0s - loss: 82161.9766 - iou_metric: 0.6745
- p_metric: 0.9843 - c_metric: 0.3936



93/93 [=====] - 19s 198ms/step - loss: 82161.9766 - iou_metric: 0.6745 - p_metric: 0.9843 - c_metric: 0.3936 - val_loss: 81788.2266 - val_iou_metric: 0.7413 - val_p_metric: 1.0000 - val_c_metric: 0.4394 - lr: 1.0000e-04

Epoch 71/100

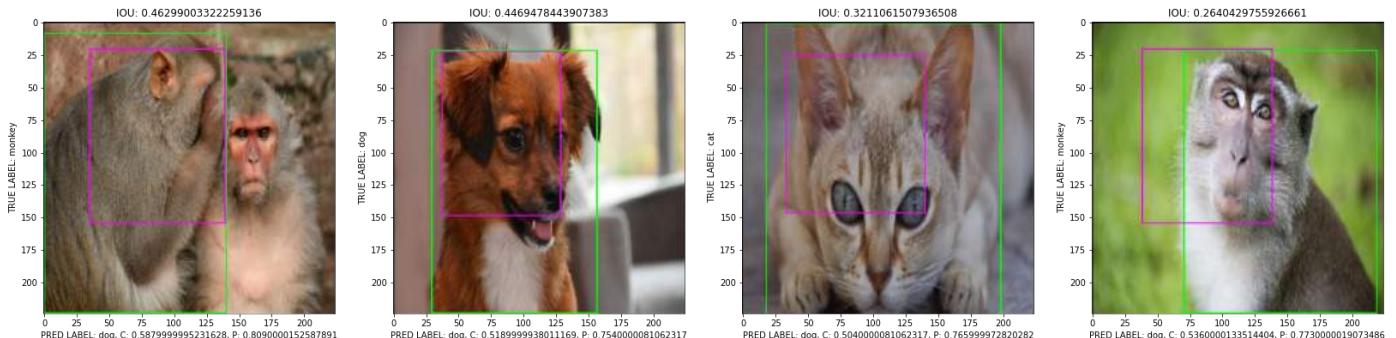
93/93 [=====] - ETA: 0s - loss: 81434.1250 - iou_metric: 0.6754
- p_metric: 0.9845 - c_metric: 0.3944



93/93 [=====] - 18s 198ms/step - loss: 81434.1250 - iou_metric: 0.6754 - p_metric: 0.9845 - c_metric: 0.3944 - val_loss: 81064.9531 - val_iou_metric: 0.7416 - val_p_metric: 1.0000 - val_c_metric: 0.4394 - lr: 1.0000e-04

Epoch 72/100

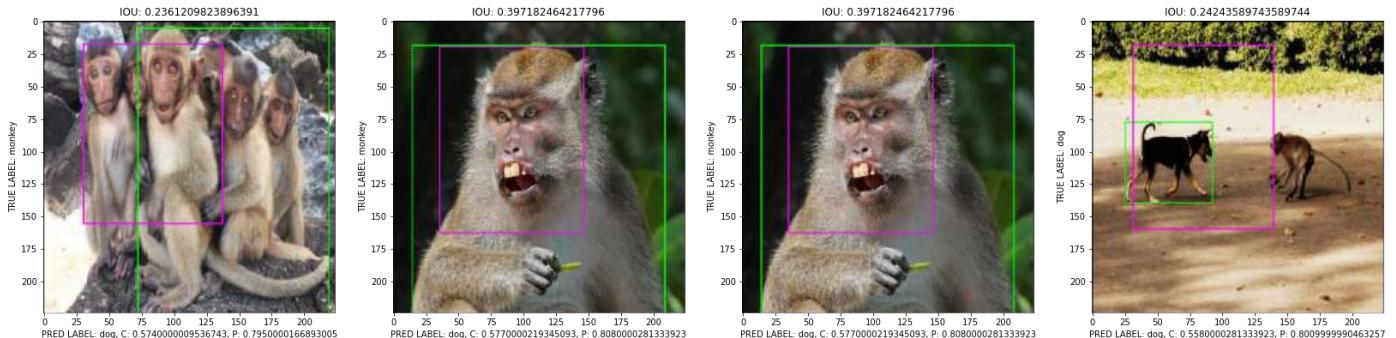
93/93 [=====] - ETA: 0s - loss: 80700.7969 - iou_metric: 0.6762
- p_metric: 0.9847 - c_metric: 0.3950



93/93 [=====] - 18s 198ms/step - loss: 80700.7969 - iou_metric: 0.6762 - p_metric: 0.9847 - c_metric: 0.3950 - val_loss: 80335.3672 - val_iou_metric: 0.7418 - val_p_metric: 1.0000 - val_c_metric: 0.4406 - lr: 1.0000e-04

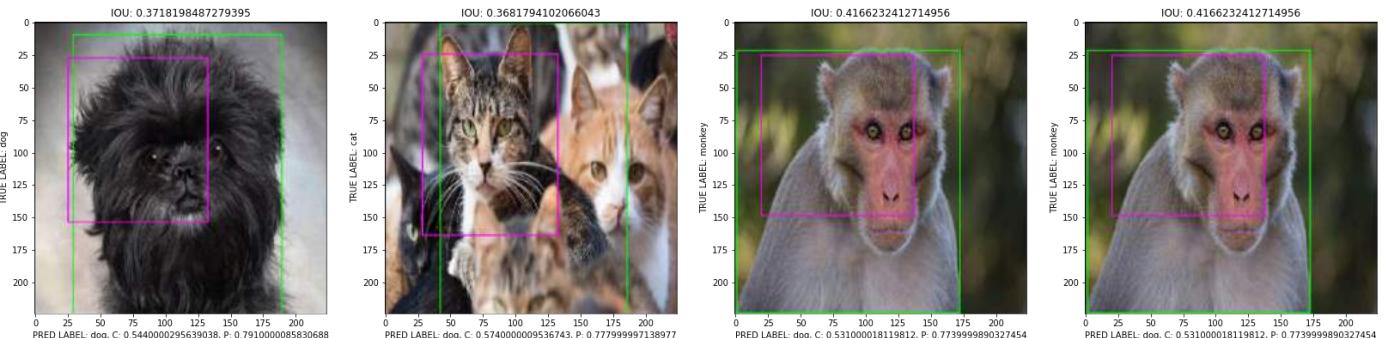
Epoch 73/100

93/93 [=====] - ETA: 0s - loss: 79976.9219 - iou_metric: 0.6771
- p_metric: 0.9849 - c_metric: 0.3955



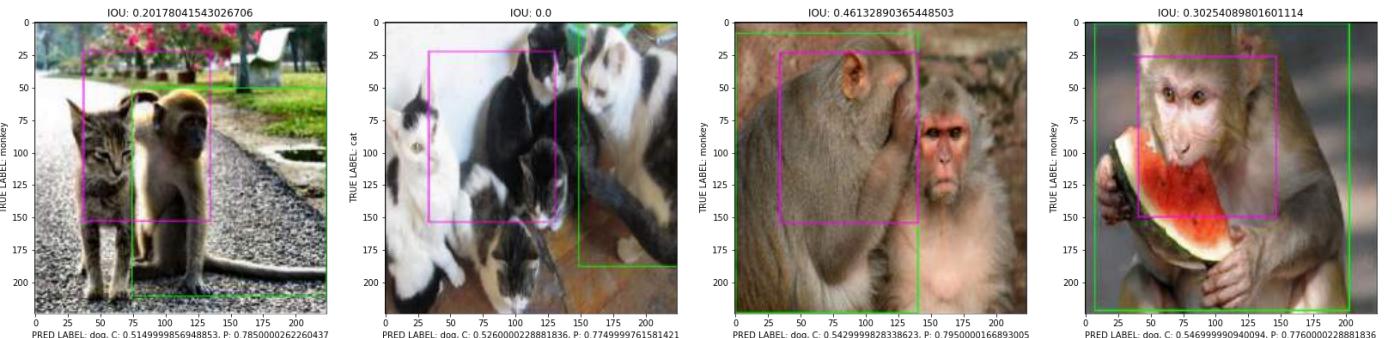
93/93 [=====] - 18s 197ms/step - loss: 79976.9219 - iou_metric: 0.6771 - p_metric: 0.9849 - c_metric: 0.3955 - val_loss: 79606.9141 - val_iou_metric: 0.7421 - val_p_metric: 1.0000 - val_c_metric: 0.4411 - lr: 1.0000e-04
Epoch 74/100

93/93 [=====] - ETA: 0s - loss: 79247.9141 - iou_metric: 0.6778 - p_metric: 0.9851 - c_metric: 0.3957



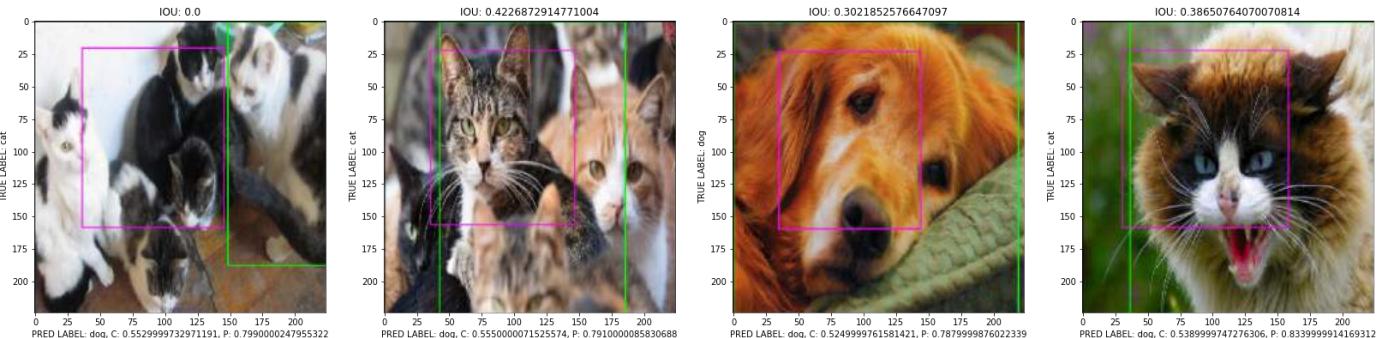
93/93 [=====] - 18s 198ms/step - loss: 79247.9141 - iou_metric: 0.6778 - p_metric: 0.9851 - c_metric: 0.3957 - val_loss: 78880.5547 - val_iou_metric: 0.7424 - val_p_metric: 1.0000 - val_c_metric: 0.4422 - lr: 1.0000e-04
Epoch 75/100

93/93 [=====] - ETA: 0s - loss: 78527.9531 - iou_metric: 0.6786 - p_metric: 0.9853 - c_metric: 0.3966



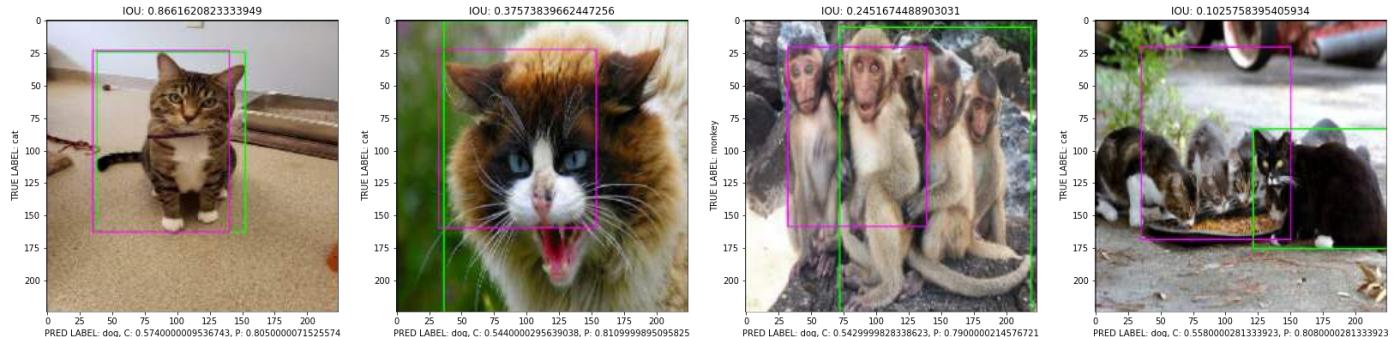
93/93 [=====] - 18s 197ms/step - loss: 78527.9531 - iou_metric: 0.6786 - p_metric: 0.9853 - c_metric: 0.3966 - val_loss: 78169.0000 - val_iou_metric: 0.7426 - val_p_metric: 1.0000 - val_c_metric: 0.4421 - lr: 1.0000e-04
Epoch 76/100

93/93 [=====] - ETA: 0s - loss: 77818.1875 - iou_metric: 0.6793 - p_metric: 0.9855 - c_metric: 0.3973



93/93 [=====] - 19s 198ms/step - loss: 77818.1875 - iou_metric: 0.6793 - p_metric: 0.9855 - c_metric: 0.3973 - val_loss: 77459.3672 - val_iou_metric: 0.7430 - val_p_metric: 1.0000 - val_c_metric: 0.4437 - lr: 1.0000e-04
Epoch 77/100

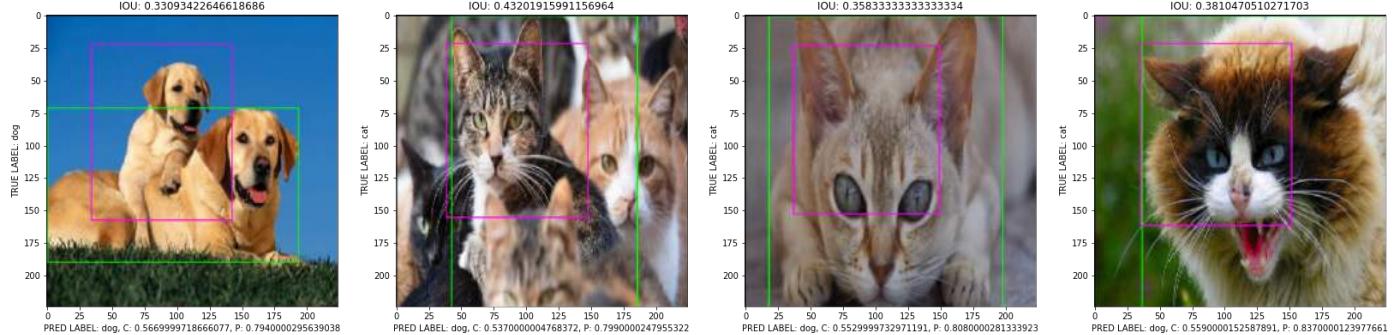
93/93 [=====] - ETA: 0s - loss: 77110.3672 - iou_metric: 0.6800 - p_metric: 0.9857 - c_metric: 0.3981



93/93 [=====] - 19s 199ms/step - loss: 77110.3672 - iou_metric: 0.6800 - p_metric: 0.9857 - c_metric: 0.3981 - val_loss: 76759.6328 - val_iou_metric: 0.7433 - val_p_metric: 1.0000 - val_c_metric: 0.4447 - lr: 1.0000e-04

Epoch 78/100

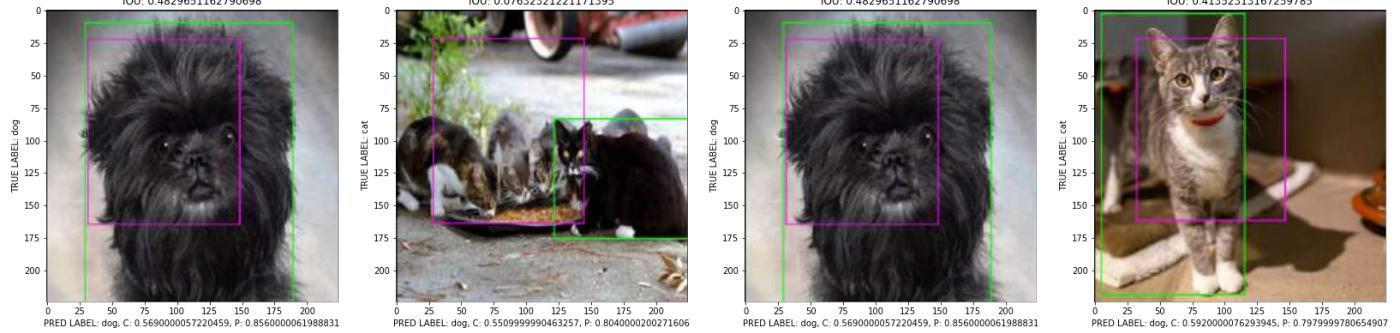
93/93 [=====] - ETA: 0s - loss: 76407.9062 - iou_metric: 0.6806 - p_metric: 0.9859 - c_metric: 0.3986



93/93 [=====] - 19s 199ms/step - loss: 76407.9062 - iou_metric: 0.6806 - p_metric: 0.9859 - c_metric: 0.3986 - val_loss: 76049.9609 - val_iou_metric: 0.7434 - val_p_metric: 1.0000 - val_c_metric: 0.4436 - lr: 1.0000e-04

Epoch 79/100

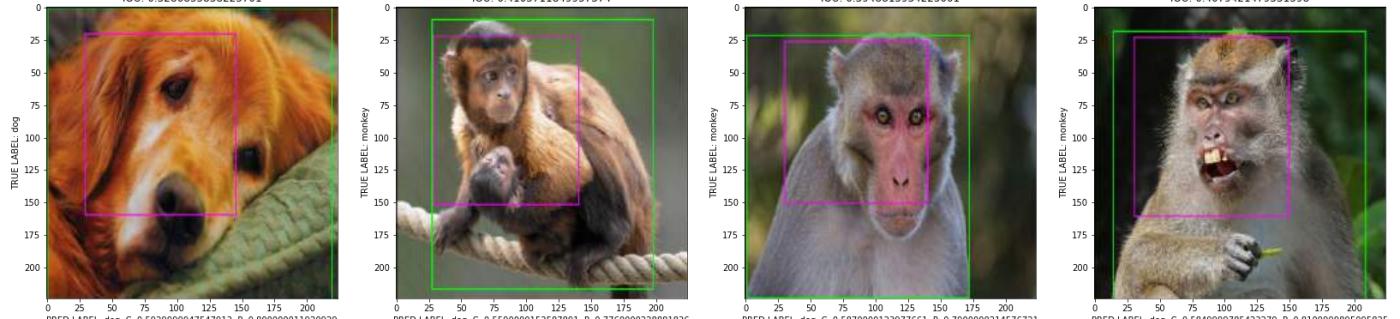
93/93 [=====] - ETA: 0s - loss: 75704.1719 - iou_metric: 0.6814 - p_metric: 0.9861 - c_metric: 0.3992



93/93 [=====] - 19s 201ms/step - loss: 75704.1719 - iou_metric: 0.6814 - p_metric: 0.9861 - c_metric: 0.3992 - val_loss: 75355.4766 - val_iou_metric: 0.7437 - val_p_metric: 1.0000 - val_c_metric: 0.4441 - lr: 1.0000e-04

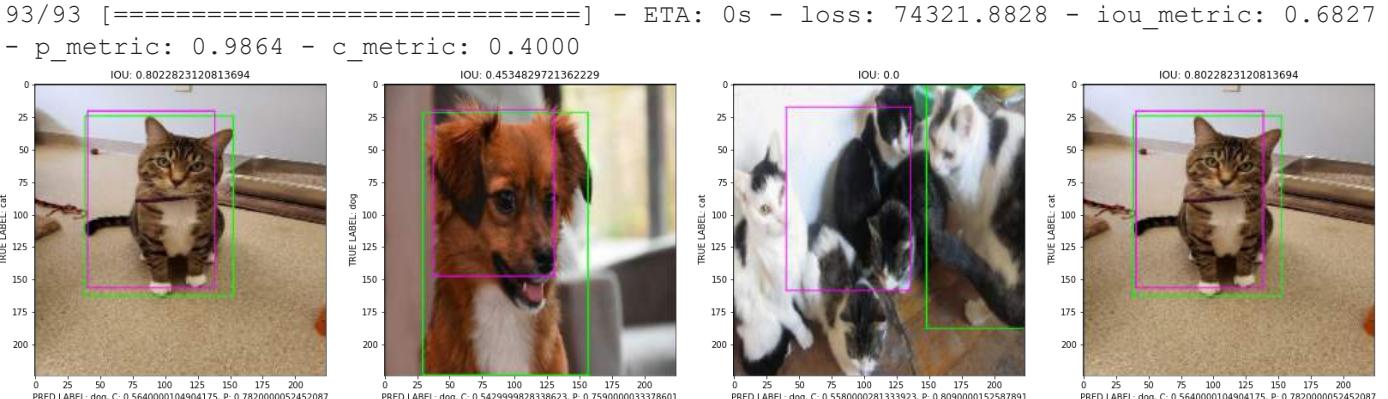
Epoch 80/100

93/93 [=====] - ETA: 0s - loss: 75012.9297 - iou_metric: 0.6820 - p_metric: 0.9862 - c_metric: 0.3994



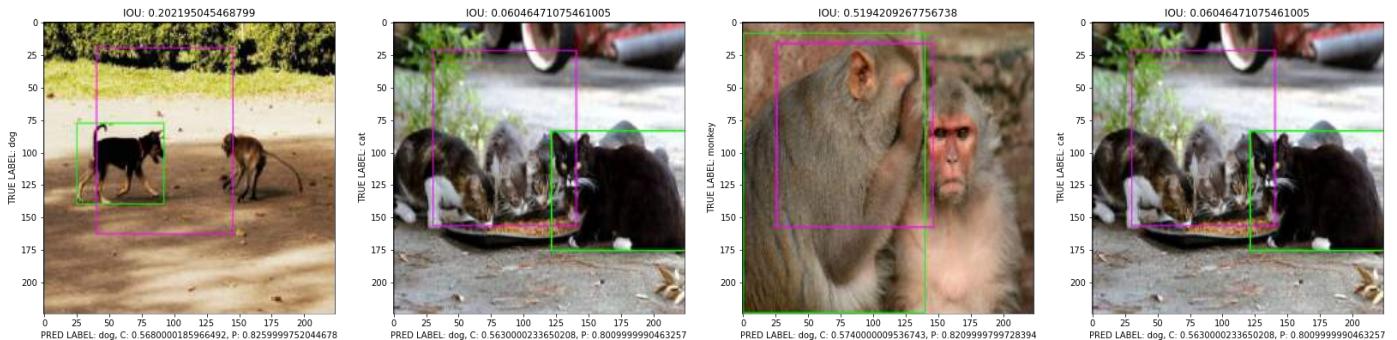
93/93 [=====] - 19s 198ms/step - loss: 75012.9297 - iou_metric: 0.6820 - p_metric: 0.9862 - c_metric: 0.3994 - val_loss: 74663.9844 - val_iou_metric: 0.7439 - val_p_metric: 1.0000 - val_c_metric: 0.4445 - lr: 1.0000e-04

Epoch 81/100



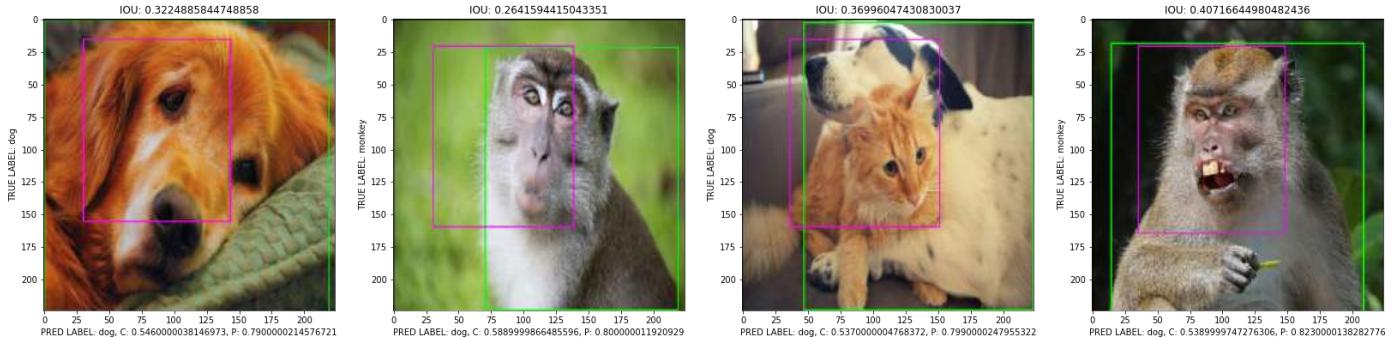
93/93 [=====] - 18s 198ms/step - loss: 74321.8828 - iou_metric: 0.6827 - p_metric: 0.9864 - c_metric: 0.4000 - val_loss: 73971.6719 - val_iou_metric: 0.7442 - val_p_metric: 1.0000 - val_c_metric: 0.4444 - lr: 1.0000e-04
Epoch 82/100

93/93 [=====] - ETA: 0s - loss: 73630.1250 - iou_metric: 0.6833
- p_metric: 0.9866 - c_metric: 0.4005



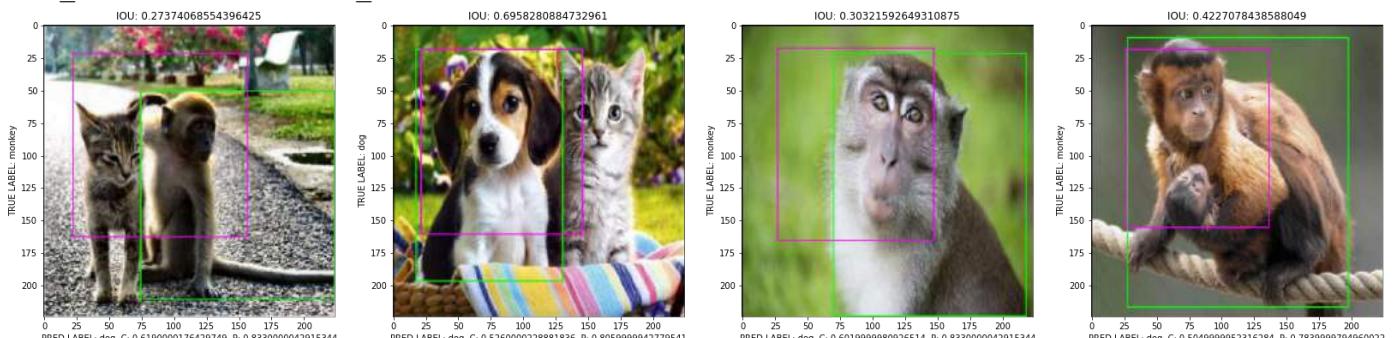
93/93 [=====] - 18s 198ms/step - loss: 73630.1250 - iou_metric: 0.6833 - p_metric: 0.9866 - c_metric: 0.4005 - val_loss: 73283.3125 - val_iou_metric: 0.7444 - val_p_metric: 1.0000 - val_c_metric: 0.4449 - lr: 1.0000e-04
Epoch 83/100

93/93 [=====] - ETA: 0s - loss: 72947.7969 - iou_metric: 0.6839
- p_metric: 0.9867 - c_metric: 0.4010



93/93 [=====] - 18s 198ms/step - loss: 72947.7969 - iou_metric: 0.6839 - p_metric: 0.9867 - c_metric: 0.4010 - val_loss: 72607.6406 - val_iou_metric: 0.7448 - val_p_metric: 1.0000 - val_c_metric: 0.4448 - lr: 1.0000e-04
Epoch 84/100

93/93 [=====] - ETA: 0s - loss: 72273.0078 - iou_metric: 0.6845
- p_metric: 0.9869 - c_metric: 0.4017

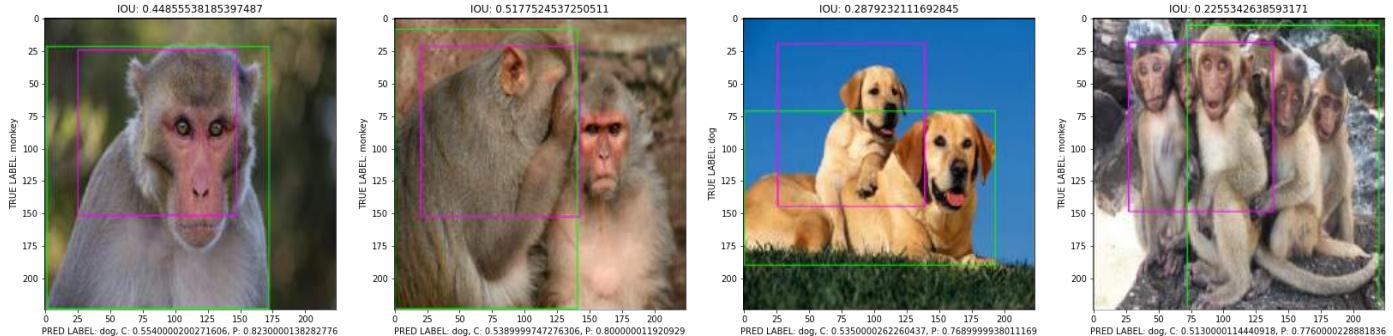


93/93 [=====] - 18s 198ms/step - loss: 72273.0078 - iou_metric: 0.6845 - p_metric: 0.9869 - c_metric: 0.4017 - val_loss: 71931.0781 - val_iou_metric: 0.

7450 - val_p_metric: 1.0000 - val_c_metric: 0.4457 - lr: 1.0000e-04

Epoch 85/100

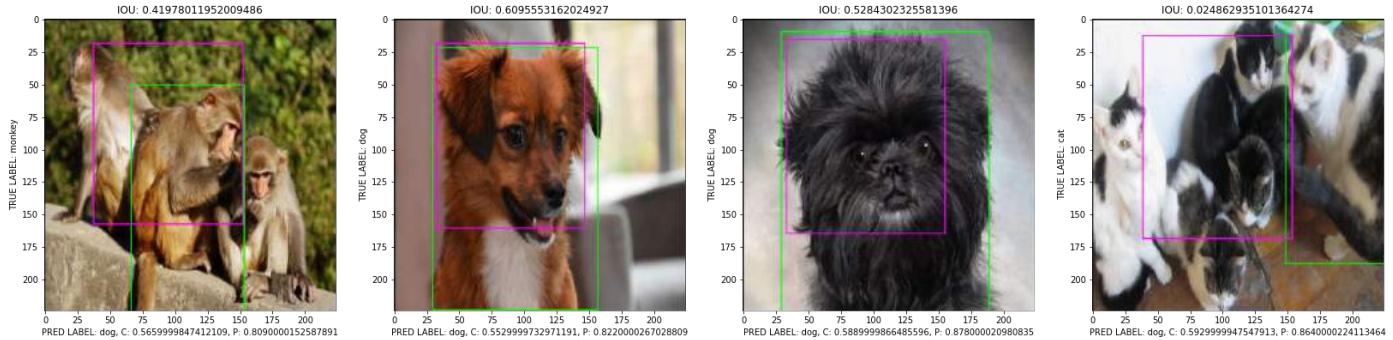
93/93 [=====] - ETA: 0s - loss: 71602.0156 - iou_metric: 0.6851
- p_metric: 0.9870 - c_metric: 0.4019



93/93 [=====] - 19s 198ms/step - loss: 71602.0156 - iou_metric: 0.6851 - p_metric: 0.9870 - c_metric: 0.4019 - val_loss: 71260.3359 - val_iou_metric: 0.7453 - val_p_metric: 1.0000 - val_c_metric: 0.4480 - lr: 1.0000e-04

Epoch 86/100

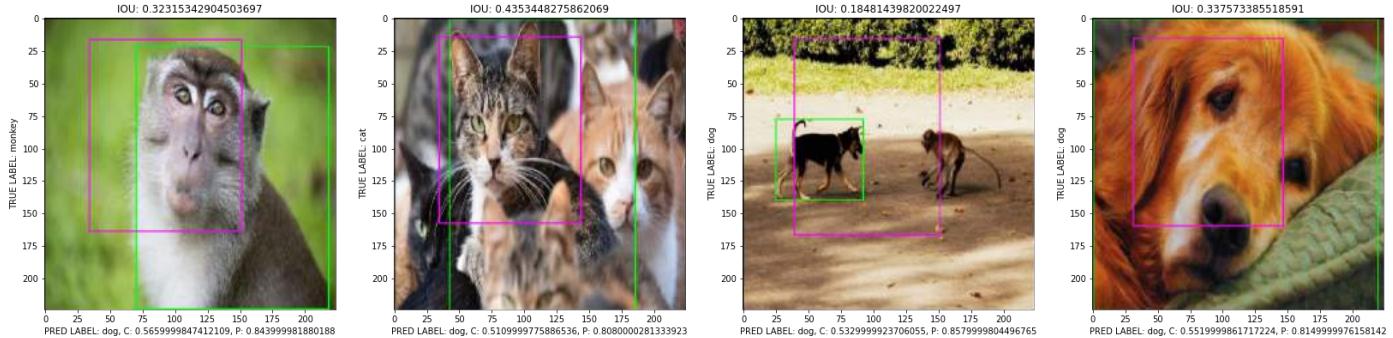
93/93 [=====] - ETA: 0s - loss: 70931.9531 - iou_metric: 0.6857
- p_metric: 0.9872 - c_metric: 0.4024



93/93 [=====] - 18s 198ms/step - loss: 70931.9531 - iou_metric: 0.6857 - p_metric: 0.9872 - c_metric: 0.4024 - val_loss: 70596.6016 - val_iou_metric: 0.7456 - val_p_metric: 1.0000 - val_c_metric: 0.4488 - lr: 1.0000e-04

Epoch 87/100

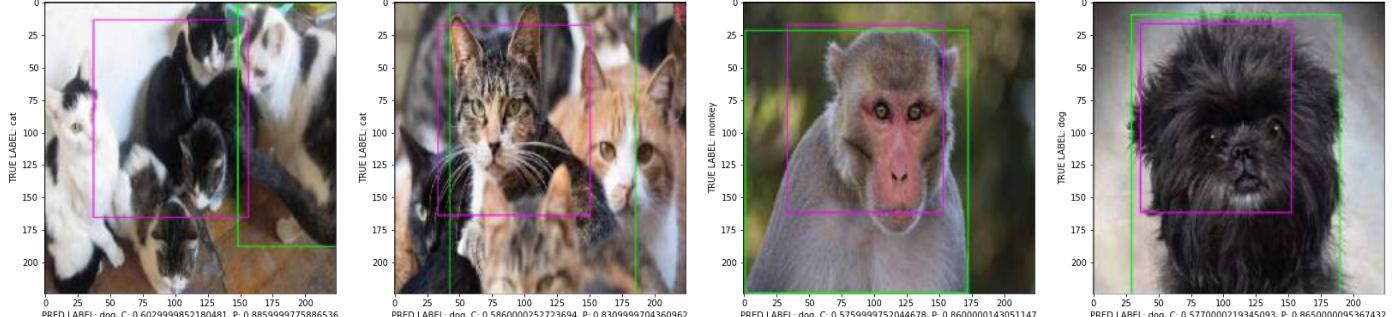
93/93 [=====] - ETA: 0s - loss: 70268.5703 - iou_metric: 0.6863
- p_metric: 0.9873 - c_metric: 0.4028



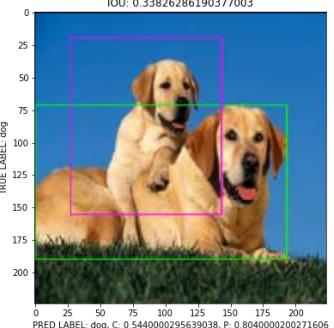
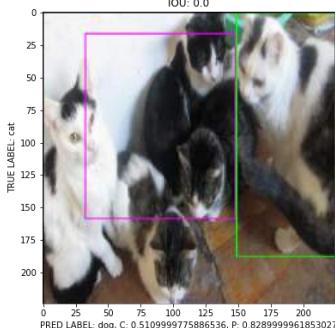
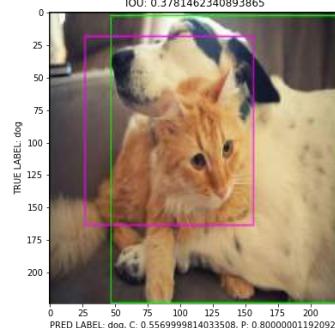
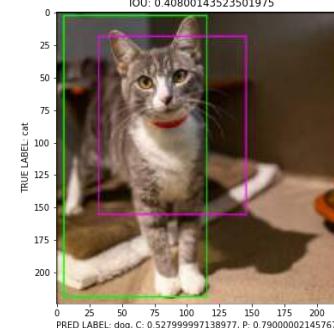
93/93 [=====] - 18s 198ms/step - loss: 70268.5703 - iou_metric: 0.6863 - p_metric: 0.9873 - c_metric: 0.4028 - val_loss: 69933.6406 - val_iou_metric: 0.7457 - val_p_metric: 1.0000 - val_c_metric: 0.4497 - lr: 1.0000e-04

Epoch 88/100

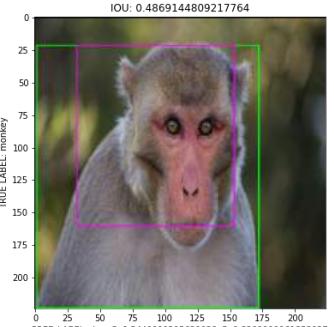
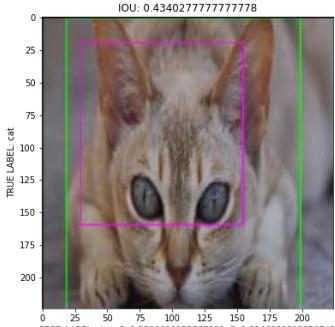
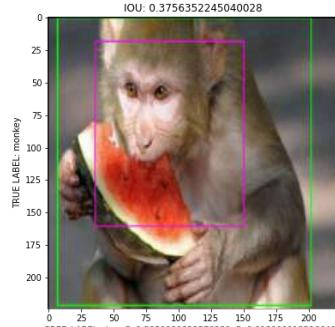
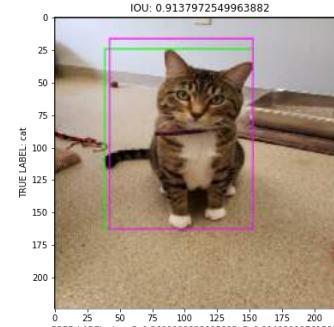
93/93 [=====] - ETA: 0s - loss: 69610.2578 - iou_metric: 0.6868
- p_metric: 0.9875 - c_metric: 0.4033



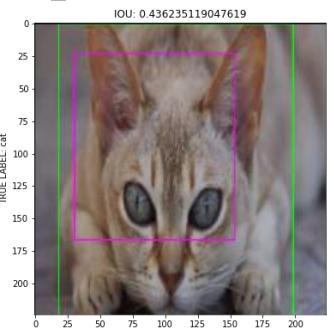
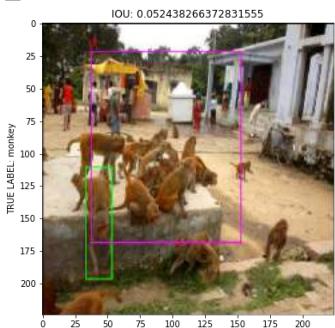
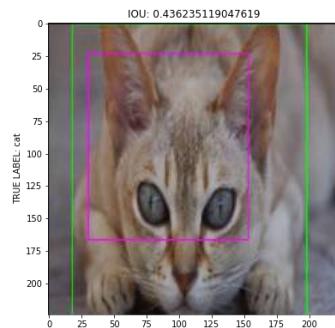
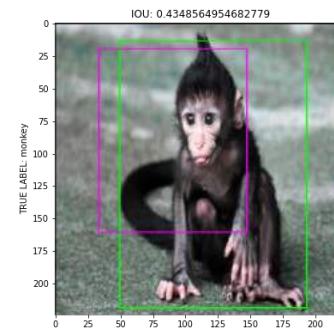
93/93 [=====] - 18s 197ms/step - loss: 69610.2578 - iou_metric: 0.6868 - p_metric: 0.9875 - c_metric: 0.4033 - val_loss: 69282.5938 - val_iou_metric: 0.7461 - val_p_metric: 1.0000 - val_c_metric: 0.4505 - lr: 1.0000e-04
 Epoch 89/100
 93/93 [=====] - ETA: 0s - loss: 68954.8359 - iou_metric: 0.6874 - p_metric: 0.9876 - c_metric: 0.4033

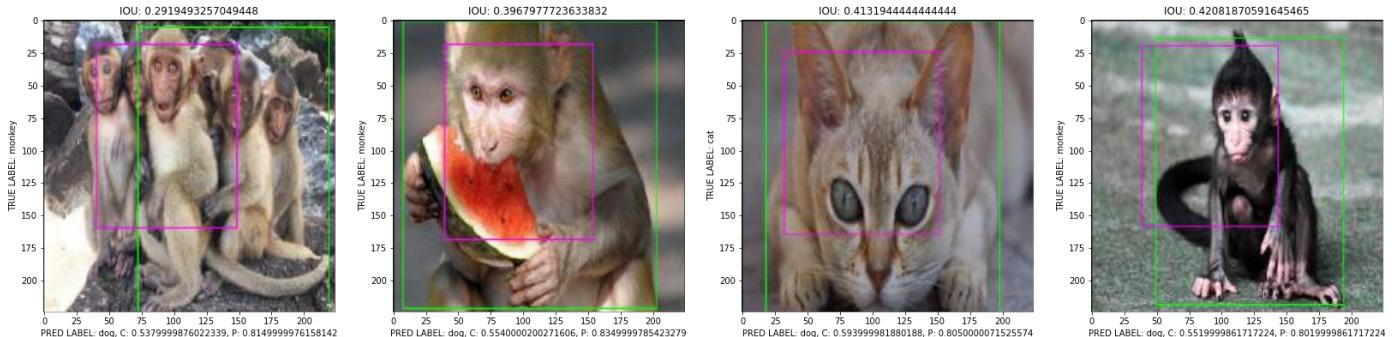
93/93 [=====] - 18s 197ms/step - loss: 68954.8359 - iou_metric: 0.6874 - p_metric: 0.9876 - c_metric: 0.4033 - val_loss: 68625.0391 - val_iou_metric: 0.7464 - val_p_metric: 1.0000 - val_c_metric: 0.4508 - lr: 1.0000e-04
 Epoch 90/100
 93/93 [=====] - ETA: 0s - loss: 68303.5469 - iou_metric: 0.6880 - p_metric: 0.9878 - c_metric: 0.4032

93/93 [=====] - 19s 201ms/step - loss: 68303.5469 - iou_metric: 0.6880 - p_metric: 0.9878 - c_metric: 0.4032 - val_loss: 67970.7500 - val_iou_metric: 0.7467 - val_p_metric: 1.0000 - val_c_metric: 0.4516 - lr: 1.0000e-04
 Epoch 91/100
 93/93 [=====] - ETA: 0s - loss: 67653.4141 - iou_metric: 0.6885 - p_metric: 0.9879 - c_metric: 0.4036

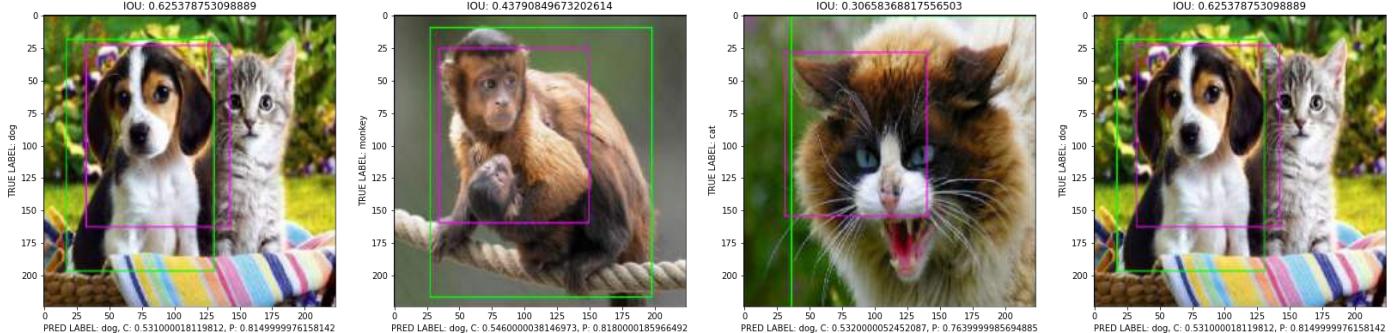
93/93 [=====] - 18s 197ms/step - loss: 67653.4141 - iou_metric: 0.6885 - p_metric: 0.9879 - c_metric: 0.4036 - val_loss: 67329.0938 - val_iou_metric: 0.7469 - val_p_metric: 1.0000 - val_c_metric: 0.4519 - lr: 1.0000e-04
 Epoch 92/100
 93/93 [=====] - ETA: 0s - loss: 67010.4453 - iou_metric: 0.6891 - p_metric: 0.9880 - c_metric: 0.4039



93/93 [=====] - 18s 198ms/step - loss: 67010.4453 - iou_metric: 0.6891 - p_metric: 0.9880 - c_metric: 0.4039 - val_loss: 66683.8359 - val_iou_metric: 0.7471 - val_p_metric: 1.0000 - val_c_metric: 0.4526 - lr: 1.0000e-04

Epoch 93/100

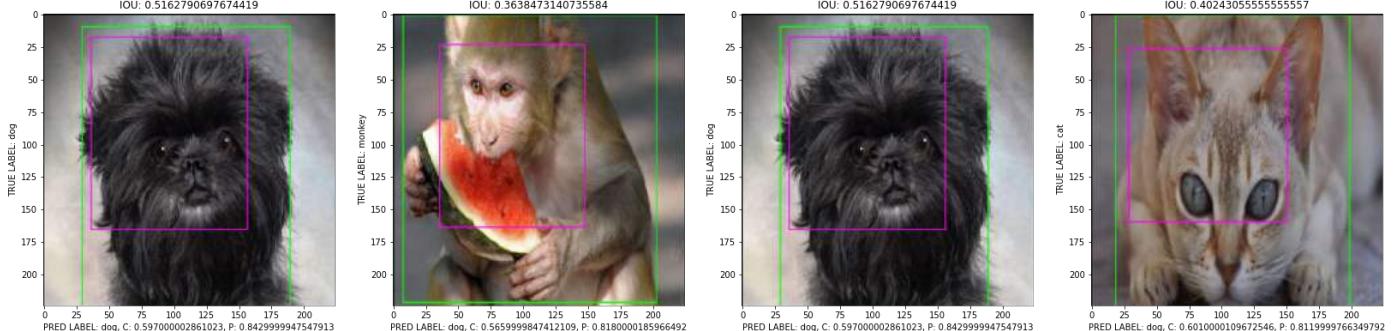
93/93 [=====] - ETA: 0s - loss: 66363.7969 - iou_metric: 0.6895 - p_metric: 0.9882 - c_metric: 0.4039



93/93 [=====] - 18s 197ms/step - loss: 66363.7969 - iou_metric: 0.6895 - p_metric: 0.9882 - c_metric: 0.4039 - val_loss: 66037.7734 - val_iou_metric: 0.7473 - val_p_metric: 1.0000 - val_c_metric: 0.4529 - lr: 1.0000e-04

Epoch 94/100

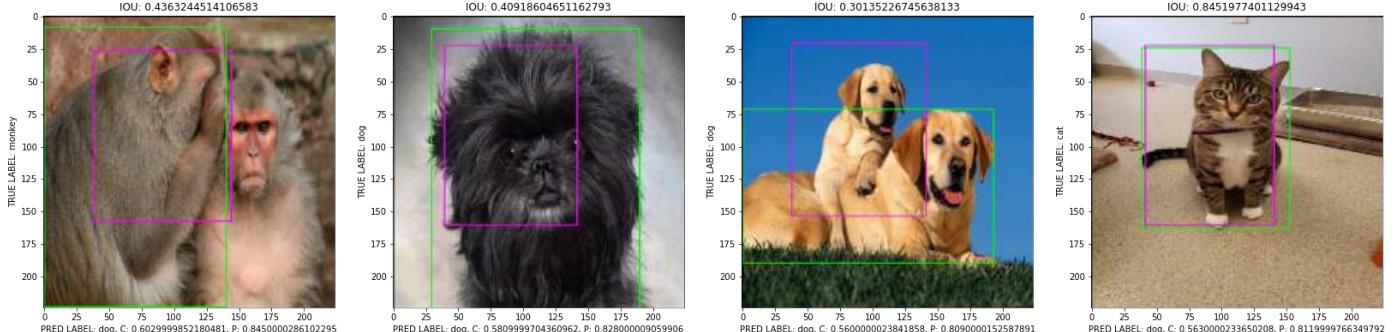
93/93 [=====] - ETA: 0s - loss: 65730.6719 - iou_metric: 0.6900 - p_metric: 0.9883 - c_metric: 0.4043



93/93 [=====] - 18s 197ms/step - loss: 65730.6719 - iou_metric: 0.6900 - p_metric: 0.9883 - c_metric: 0.4043 - val_loss: 65415.1914 - val_iou_metric: 0.7475 - val_p_metric: 1.0000 - val_c_metric: 0.4536 - lr: 1.0000e-04

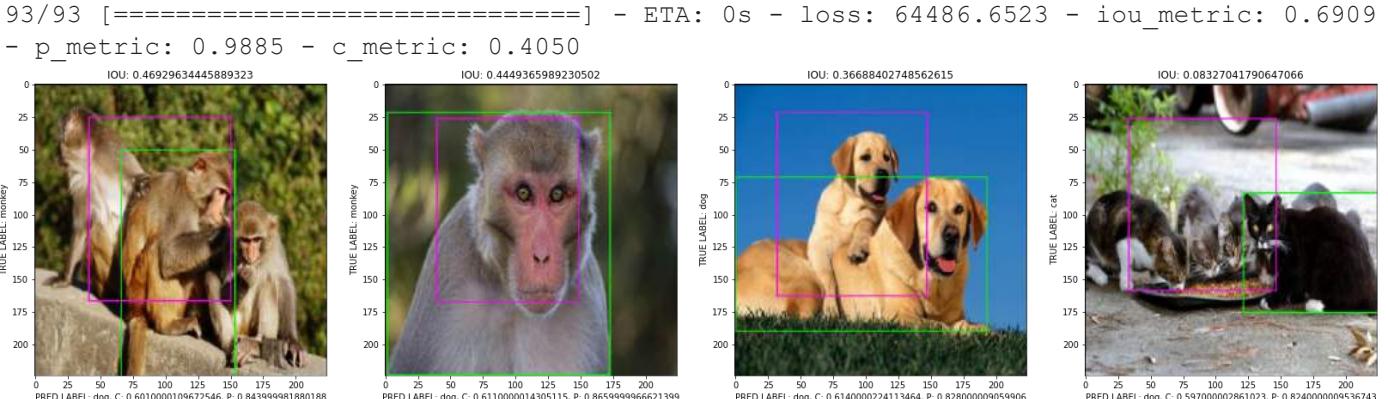
Epoch 95/100

93/93 [=====] - ETA: 0s - loss: 65109.8125 - iou_metric: 0.6904 - p_metric: 0.9884 - c_metric: 0.4045



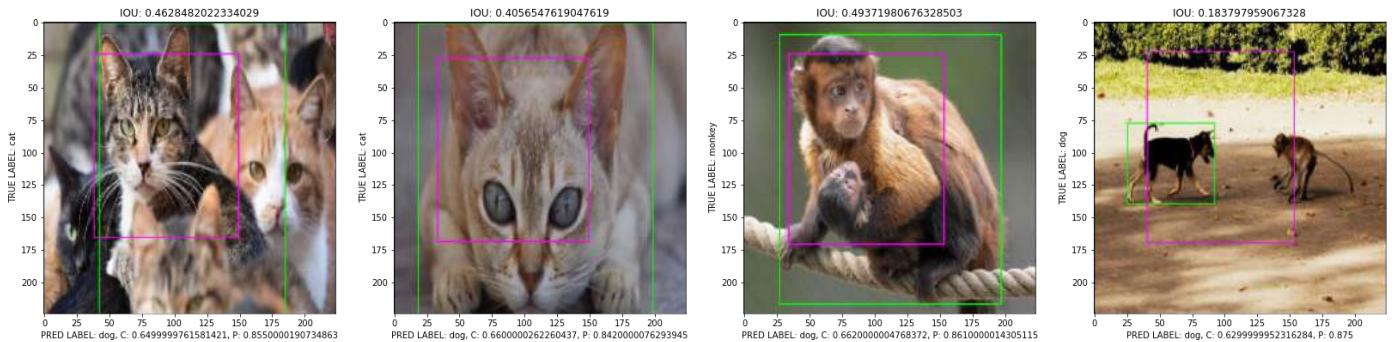
93/93 [=====] - 18s 198ms/step - loss: 65109.8125 - iou_metric: 0.6904 - p_metric: 0.9884 - c_metric: 0.4045 - val_loss: 64794.3047 - val_iou_metric: 0.7478 - val_p_metric: 1.0000 - val_c_metric: 0.4552 - lr: 1.0000e-04

Epoch 96/100



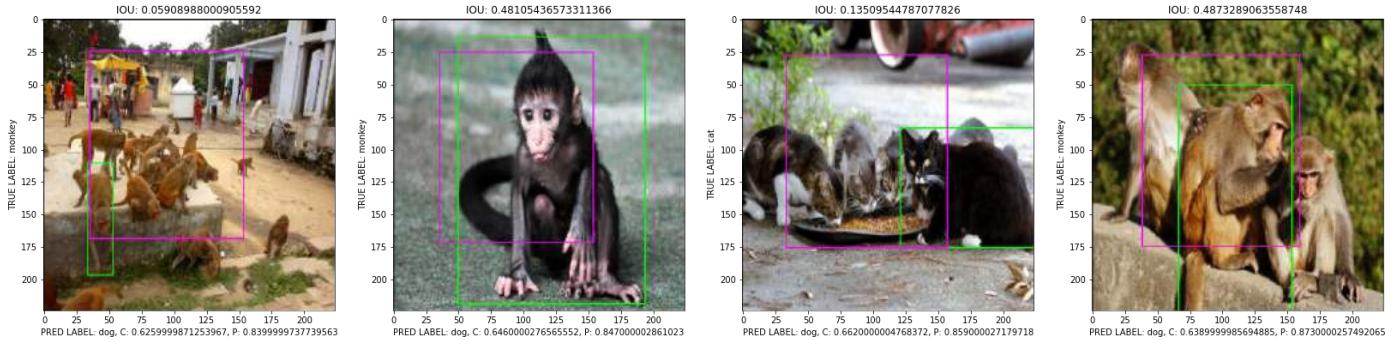
93/93 [=====] - 18s 198ms/step - loss: 64486.6523 - iou_metric: 0.6909 - p_metric: 0.9885 - c_metric: 0.4050 - val_loss: 64172.2109 - val_iou_metric: 0.7480 - val_p_metric: 1.0000 - val_c_metric: 0.4546 - lr: 1.0000e-04
Epoch 97/100

93/93 [=====] - ETA: 0s - loss: 63869.2578 - iou_metric: 0.6914
- p_metric: 0.9886 - c_metric: 0.4050



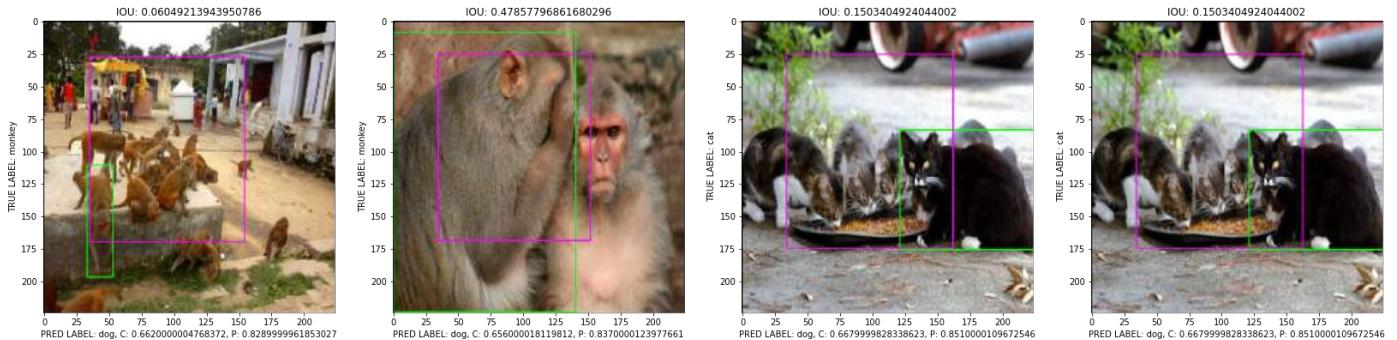
93/93 [=====] - 18s 198ms/step - loss: 63869.2578 - iou_metric: 0.6914 - p_metric: 0.9886 - c_metric: 0.4050 - val_loss: 63564.5469 - val_iou_metric: 0.7483 - val_p_metric: 1.0000 - val_c_metric: 0.4557 - lr: 1.0000e-04
Epoch 98/100

93/93 [=====] - ETA: 0s - loss: 63267.1484 - iou_metric: 0.6919
- p_metric: 0.9888 - c_metric: 0.4053



93/93 [=====] - 18s 198ms/step - loss: 63267.1484 - iou_metric: 0.6919 - p_metric: 0.9888 - c_metric: 0.4053 - val_loss: 62955.3398 - val_iou_metric: 0.7486 - val_p_metric: 1.0000 - val_c_metric: 0.4559 - lr: 1.0000e-04
Epoch 99/100

93/93 [=====] - ETA: 0s - loss: 62652.5781 - iou_metric: 0.6923
- p_metric: 0.9889 - c_metric: 0.4055



93/93 [=====] - 18s 198ms/step - loss: 62652.5781 - iou_metric: 0.6923 - p_metric: 0.9889 - c_metric: 0.4055 - val_loss: 62348.8008 - val_iou_metric: 0.

7487 - val_p_metric: 1.0000 - val_c_metric: 0.4566 - lr: 1.0000e-04
 Epoch 100/100
 93/93 [=====] - ETA: 0s - loss: 62044.0508 - iou_metric: 0.6927
 - p_metric: 0.9890 - c_metric: 0.4060

IOU: 0.46865632738172813
 TRUE LABEL: monkey
 PRED LABEL: dog, C: 0.6909999847412109, P: 0.8679999709129333

IOU: 0.6616512044905278
 TRUE LABEL: dog
 PRED LABEL: dog, C: 0.6579999923706055, P: 0.8180000185966492

IOU: 0.19151923616961528
 TRUE LABEL: dog
 PRED LABEL: dog, C: 0.686999976349792, P: 0.847000002861023

IOU: 0.29574072619702085
 TRUE LABEL: monkey
 PRED LABEL: dog, C: 0.6779999732971191, P: 0.8309999704360969

93/93 [=====] - 18s 198ms/step - loss: 62044.0508 - iou_metric: 0.6927
 - p_metric: 0.9890 - c_metric: 0.4060 - val_loss: 61727.8906 - val_iou_metric: 0.7489 - val_p_metric: 1.0000 - val_c_metric: 0.4568 - lr: 1.0000e-04

EVALUATE MODEL ON TESTING DATA

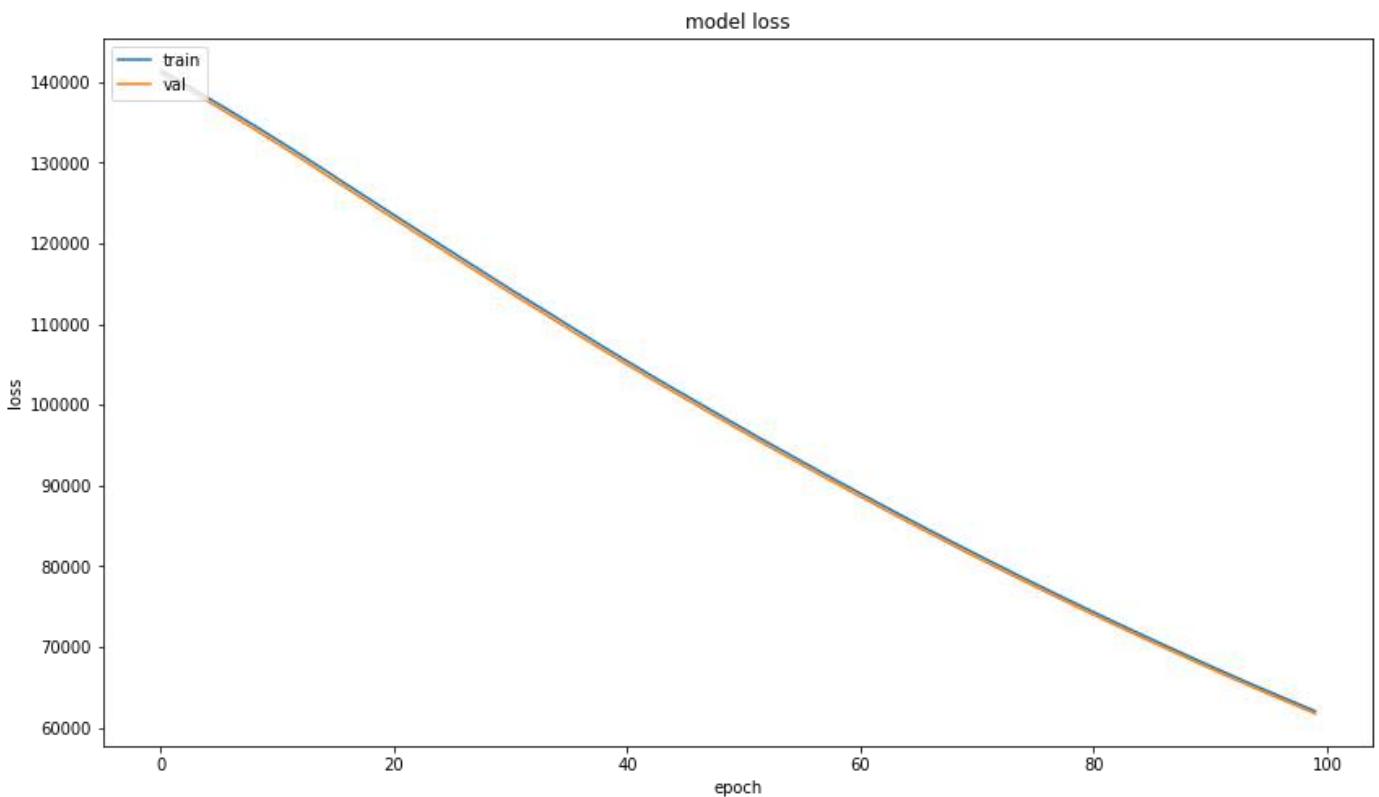
```
In [ ]: model1.save_weights('model1_full.tf')
model1.evaluate(test_gen)
```

5/5 [=====] - 0s 76ms/step - loss: 61727.8906 - iou_metric: 0.7493 - p_metric: 1.0000 - c_metric: 0.4580
[61727.890625, 0.7493180632591248, 1.0, 0.45803919434547424]

Out[]:

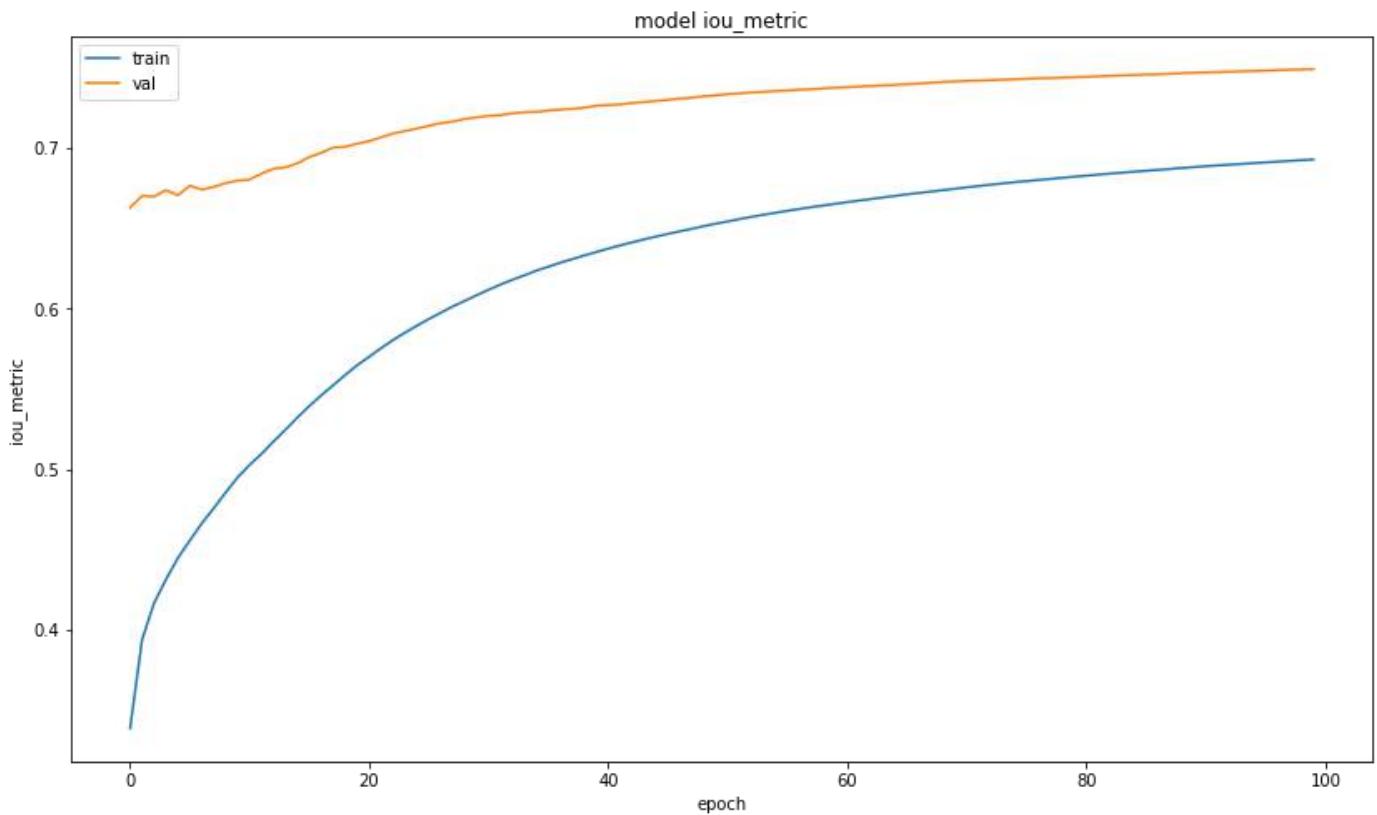
PLOT TRAINING AND VALIDATION LOSS

```
In [ ]: plot_metric(fit1, 'loss')
```

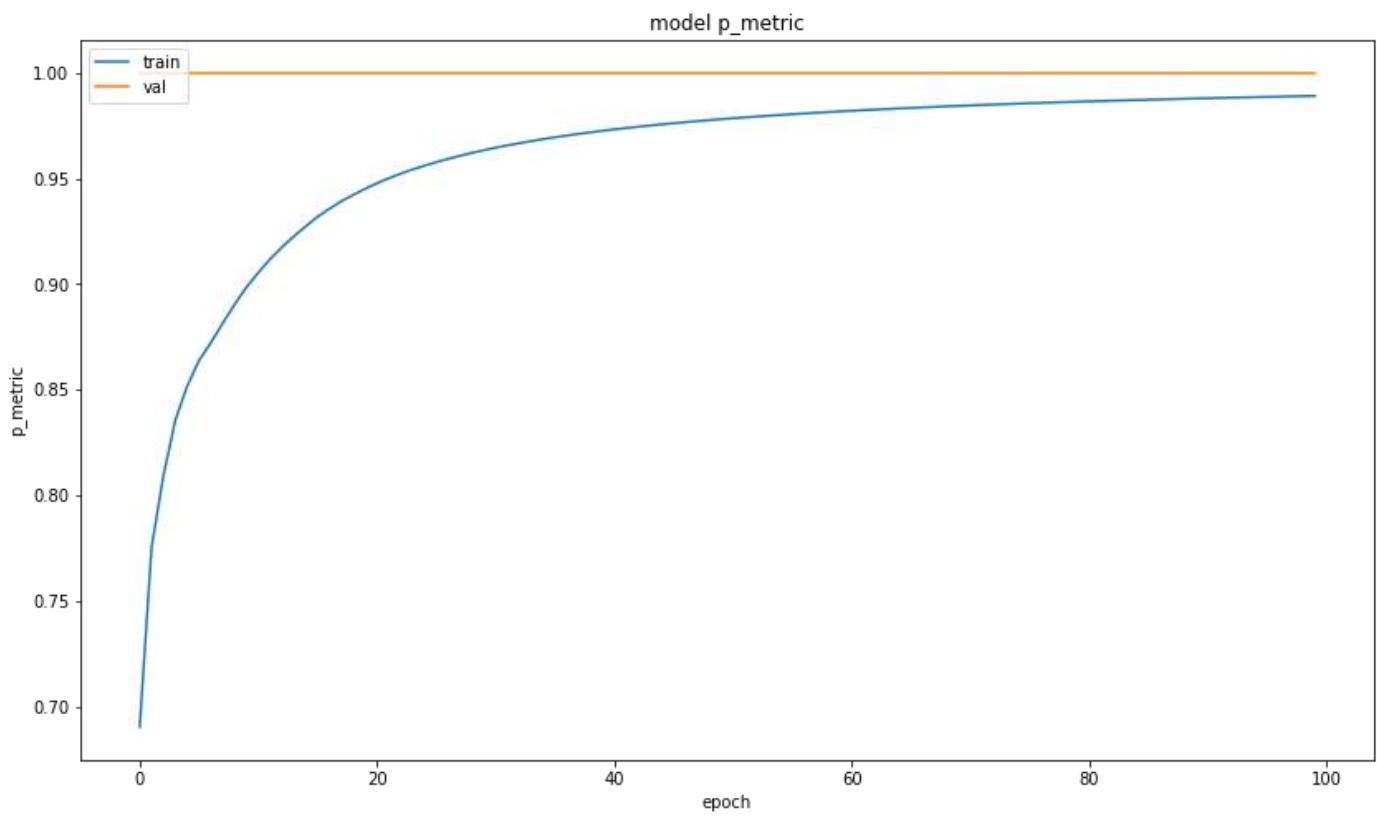


PLOT TRAINING AND VALIDATION METRICS

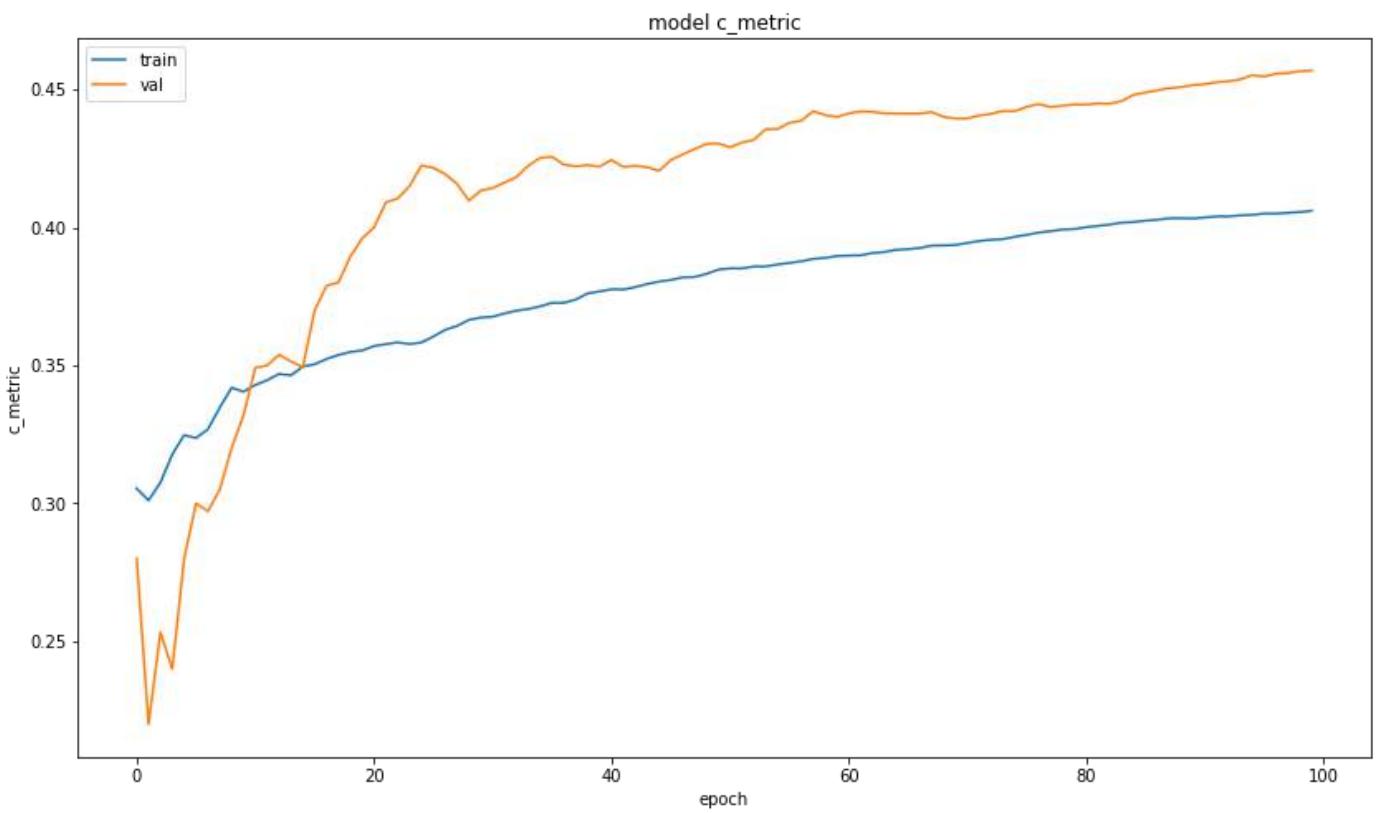
```
In [ ]: plot_metric(fit1, 'iou_metric')
```



```
In [ ]: plot_metric(fit1, 'p_metric')
```



```
In [ ]: plot_metric(fit1, 'c_metric')
```



CREATE MODEL 2 USING TRANSFER LEARNING

```
In [ ]: X_inp = tf.keras.layers.Input(shape=(H, W, C), batch_size=B)
base = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet', input_te
X = tf.keras.layers.Flatten()(base.output)
X = tf.keras.layers.Dense(2048, activation='sigmoid')(X)
X = tf.keras.layers.Dropout(0.3)(X)
X = tf.keras.layers.Dense(1024, activation='tanh')(X)
X = tf.keras.layers.Dropout(0.2)(X)
X = tf.keras.layers.Dense(512, activation='relu')(X)
X = tf.keras.layers.Dropout(0.1)(X)
X = tf.keras.layers.Dense(len(N)+5)(X)
model2 = tf.keras.Model(inputs=X_inp, outputs=X, name='VGG-BasedNet')
model2.summary()
```

Model: "VGG-BasedNet"

Layer (type)	Output Shape	Param #
<hr/>		
input_2 (InputLayer)	[(5, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(5, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(5, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(5, 112, 112, 64)	0
block2_conv1 (Conv2D)	(5, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(5, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(5, 56, 56, 128)	0
block3_conv1 (Conv2D)	(5, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(5, 56, 56, 256)	590080

block3_conv3 (Conv2D)	(5, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(5, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(5, 28, 28, 256)	0
block4_conv1 (Conv2D)	(5, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(5, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(5, 28, 28, 512)	2359808
block4_conv4 (Conv2D)	(5, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(5, 14, 14, 512)	0
block5_conv1 (Conv2D)	(5, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(5, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(5, 14, 14, 512)	2359808
block5_conv4 (Conv2D)	(5, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(5, 7, 7, 512)	0
flatten (Flatten)	(5, 25088)	0
dense (Dense)	(5, 2048)	51382272
dropout (Dropout)	(5, 2048)	0
dense_1 (Dense)	(5, 1024)	2098176
dropout_1 (Dropout)	(5, 1024)	0
dense_2 (Dense)	(5, 512)	524800
dropout_2 (Dropout)	(5, 512)	0
dense_3 (Dense)	(5, 8)	4104

=====

Total params: 74,033,736
Trainable params: 74,033,736
Non-trainable params: 0

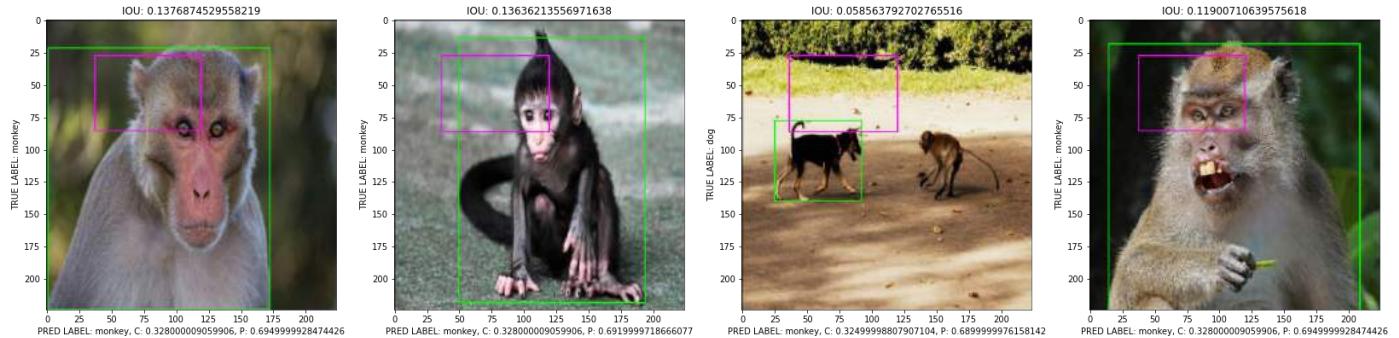
TRAIN MODEL 2 USING PREVIOUS LOSS AND METRIC

```
In [ ]: model2.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4),
                      loss=loss,
                      metrics=metric)
fit2 = model2.fit(train_gen,
                   epochs=e,
                   batch_size=B,
                   callbacks=callback,
                   validation_data=val_gen,
                   validation_batch_size=B)
```

Epoch 1/100
6/93 [>.....] - ETA: 13s - loss: 11.3693 - iou_metric: 0.1753 -
p_metric: 0.7333 - c_metric: 0.3667

WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0434s vs `on_train_batch_end` time: 0.0920s). Check your callbacks.

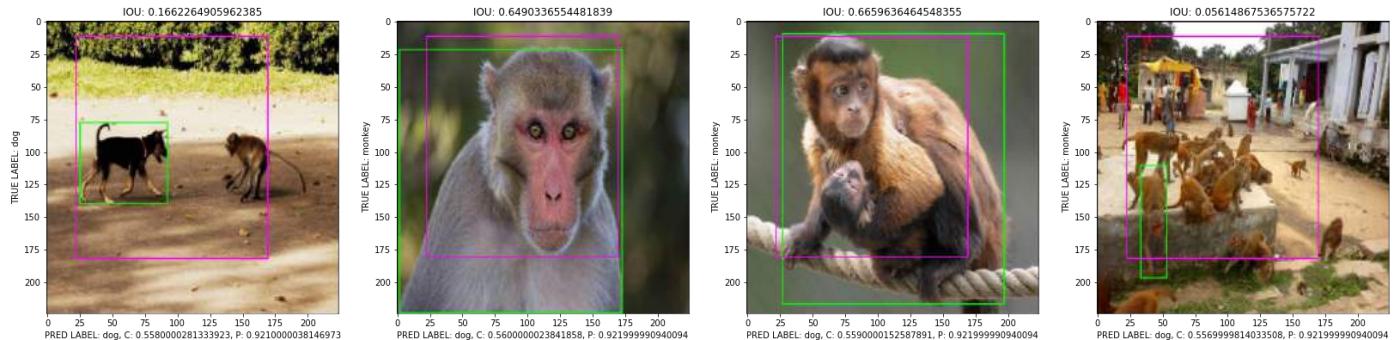
93/93 [=====] - ETA: 0s - loss: 1.6559 - iou_metric: 0.5032 - p_metric: 0.9290 - c_metric: 0.3441



93/93 [=====] - 22s 191ms/step - loss: 1.6559 - iou_metric: 0.5032 - p_metric: 0.9290 - c_metric: 0.3441 - val_loss: 0.8047 - val_iou_metric: 0.4901 - val_p_metric: 1.0000 - val_c_metric: 0.4400 - lr: 1.0000e-04

Epoch 2/100

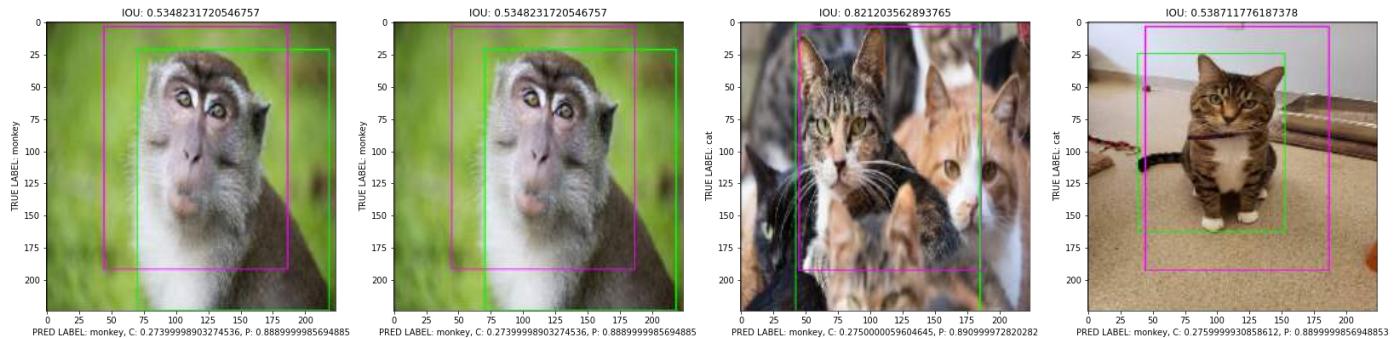
93/93 [=====] - ETA: 0s - loss: 0.5604 - iou_metric: 0.5730 - p_metric: 0.9602 - c_metric: 0.3548



93/93 [=====] - 17s 180ms/step - loss: 0.5604 - iou_metric: 0.5730 - p_metric: 0.9602 - c_metric: 0.3548 - val_loss: 0.3971 - val_iou_metric: 0.6108 - val_p_metric: 1.0000 - val_c_metric: 0.3600 - lr: 1.0000e-04

Epoch 3/100

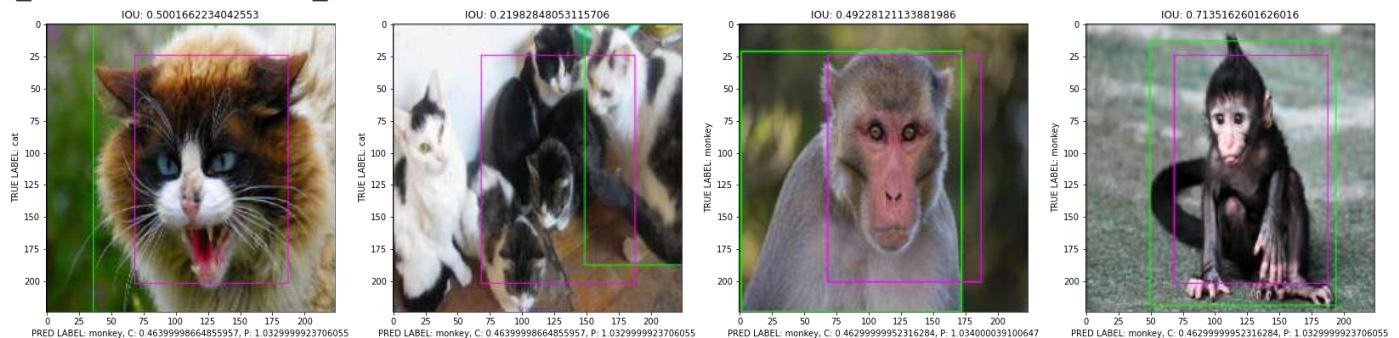
93/93 [=====] - ETA: 0s - loss: 0.4962 - iou_metric: 0.6072 - p_metric: 0.9706 - c_metric: 0.3434



93/93 [=====] - 17s 181ms/step - loss: 0.4962 - iou_metric: 0.6072 - p_metric: 0.9706 - c_metric: 0.3434 - val_loss: 0.3517 - val_iou_metric: 0.6569 - val_p_metric: 1.0000 - val_c_metric: 0.3867 - lr: 1.0000e-04

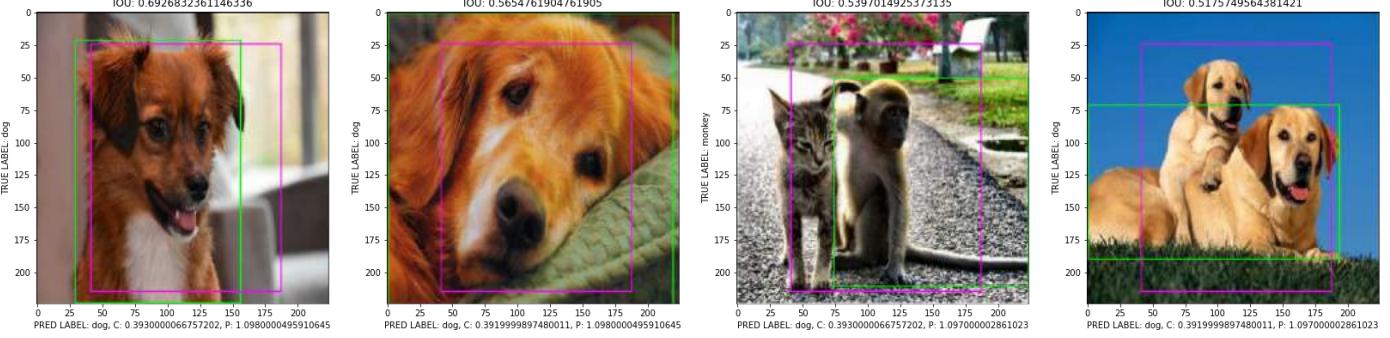
Epoch 4/100

93/93 [=====] - ETA: 0s - loss: 0.4666 - iou_metric: 0.6259 - p_metric: 0.9763 - c_metric: 0.3398



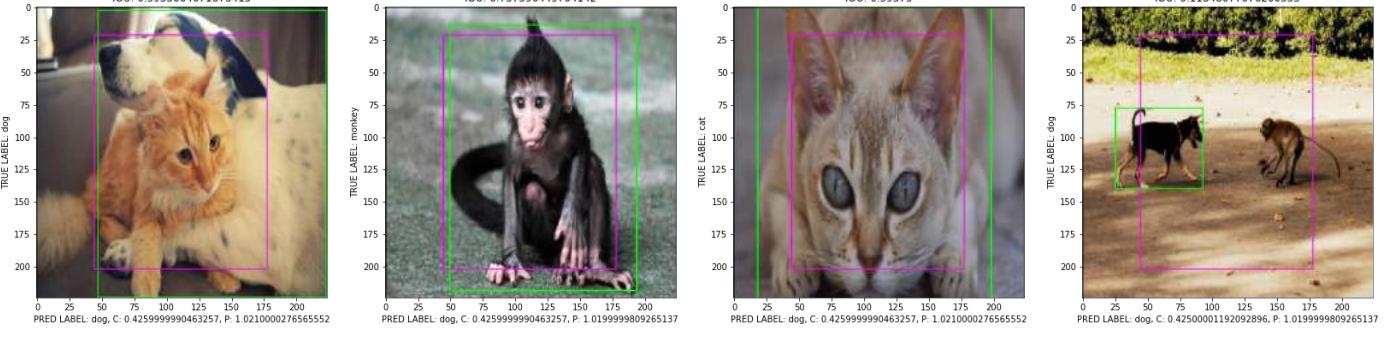
93/93 [=====] - 17s 183ms/step - loss: 0.4666 - iou_metric: 0.6259 - p_metric: 0.9763 - c_metric: 0.3398 - val_loss: 0.3339 - val_iou_metric: 0.6753 - val_p_metric: 1.0000 - val_c_metric: 0.4000 - lr: 1.0000e-04
Epoch 5/100

93/93 [=====] - ETA: 0s - loss: 0.4620 - iou_metric: 0.6378 - p_metric: 0.9802 - c_metric: 0.3376



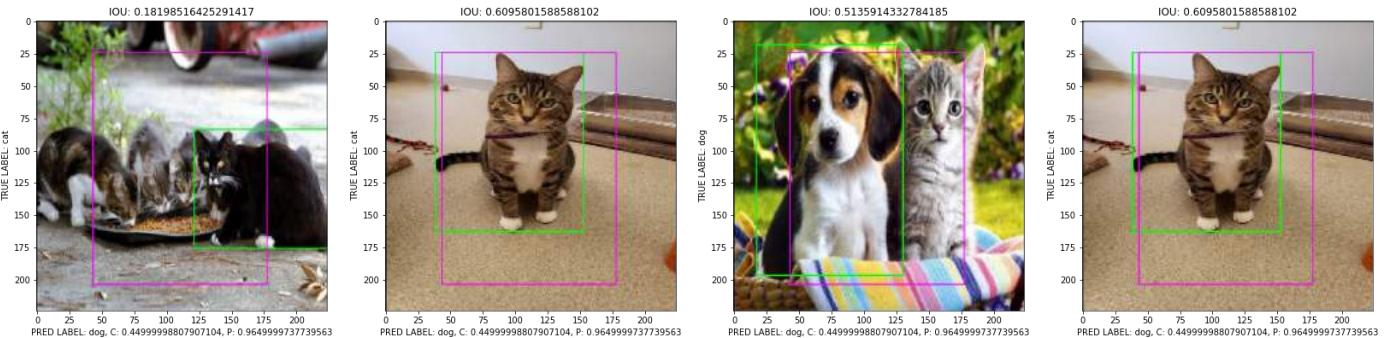
93/93 [=====] - 17s 182ms/step - loss: 0.4620 - iou_metric: 0.6378 - p_metric: 0.9802 - c_metric: 0.3376 - val_loss: 0.3351 - val_iou_metric: 0.6950 - val_p_metric: 1.0000 - val_c_metric: 0.3760 - lr: 1.0000e-04
Epoch 6/100

93/93 [=====] - ETA: 0s - loss: 0.4399 - iou_metric: 0.6472 - p_metric: 0.9835 - c_metric: 0.3355



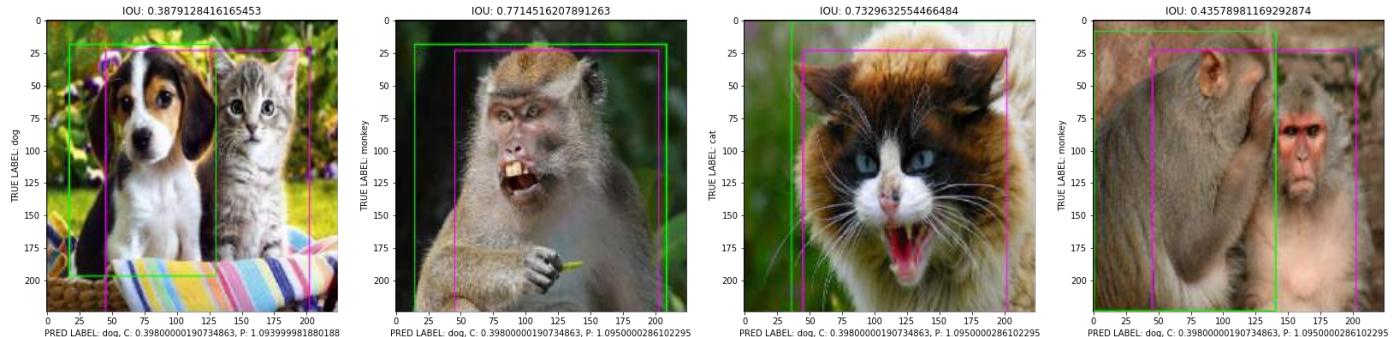
93/93 [=====] - 17s 185ms/step - loss: 0.4399 - iou_metric: 0.6472 - p_metric: 0.9835 - c_metric: 0.3355 - val_loss: 0.3259 - val_iou_metric: 0.7051 - val_p_metric: 1.0000 - val_c_metric: 0.3600 - lr: 1.0000e-04
Epoch 7/100

93/93 [=====] - ETA: 0s - loss: 0.4431 - iou_metric: 0.6544 - p_metric: 0.9853 - c_metric: 0.3367



93/93 [=====] - 17s 186ms/step - loss: 0.4431 - iou_metric: 0.6544 - p_metric: 0.9853 - c_metric: 0.3367 - val_loss: 0.3291 - val_iou_metric: 0.7121 - val_p_metric: 1.0000 - val_c_metric: 0.3486 - lr: 1.0000e-04
Epoch 8/100

93/93 [=====] - ETA: 0s - loss: 0.4267 - iou_metric: 0.6607 - p_metric: 0.9868 - c_metric: 0.3333

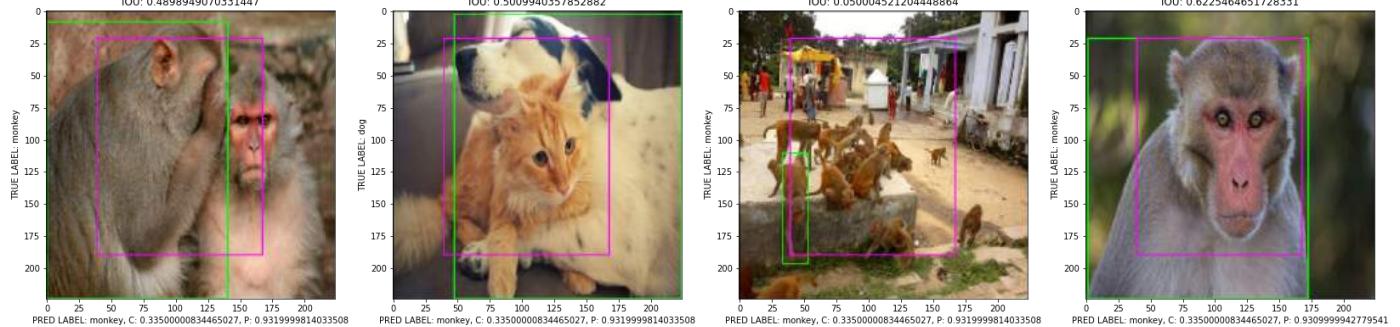


93/93 [=====] - 17s 184ms/step - loss: 0.4267 - iou_metric: 0.6649 - p_metric: 0.9878 - c_metric: 0.3357

607 - p_metric: 0.9868 - c_metric: 0.3333 - val_loss: 0.3508 - val_iou_metric: 0.7189 - val_p_metric: 1.0000 - val_c_metric: 0.3400 - lr: 1.0000e-04

Epoch 9/100

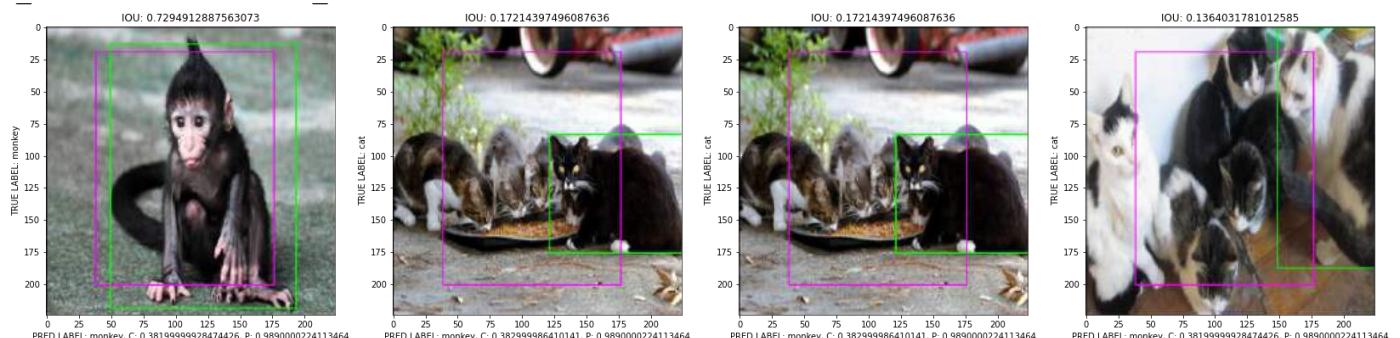
93/93 [=====] - ETA: 0s - loss: 0.4351 - iou_metric: 0.6649 - p_metric: 0.9878 - c_metric: 0.3357



93/93 [=====] - 17s 183ms/step - loss: 0.4351 - iou_metric: 0.6649 - p_metric: 0.9878 - c_metric: 0.3357 - val_loss: 0.3324 - val_iou_metric: 0.7210 - val_p_metric: 1.0000 - val_c_metric: 0.3511 - lr: 1.0000e-04

Epoch 10/100

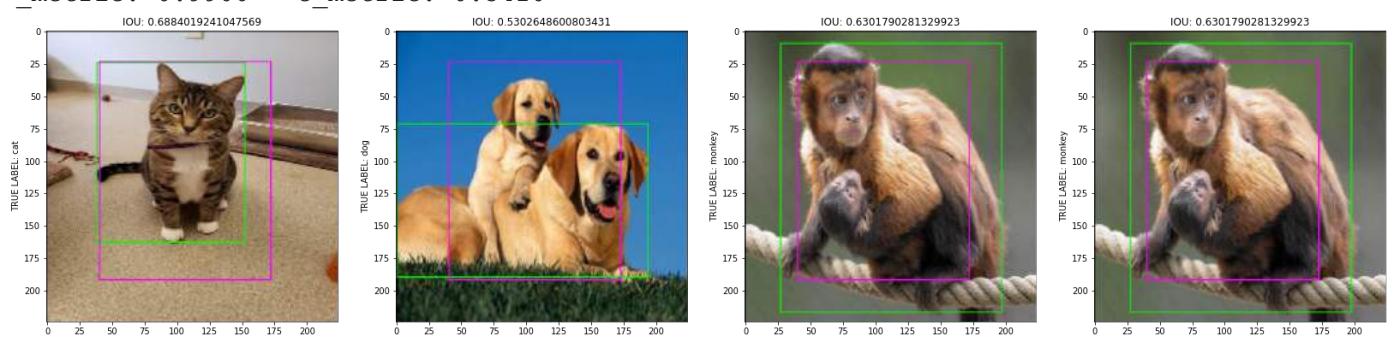
93/93 [=====] - ETA: 0s - loss: 0.3986 - iou_metric: 0.6698 - p_metric: 0.9890 - c_metric: 0.3400



93/93 [=====] - 17s 185ms/step - loss: 0.3986 - iou_metric: 0.6698 - p_metric: 0.9890 - c_metric: 0.3400 - val_loss: 0.3195 - val_iou_metric: 0.7247 - val_p_metric: 1.0000 - val_c_metric: 0.3600 - lr: 3.0000e-05

Epoch 11/100

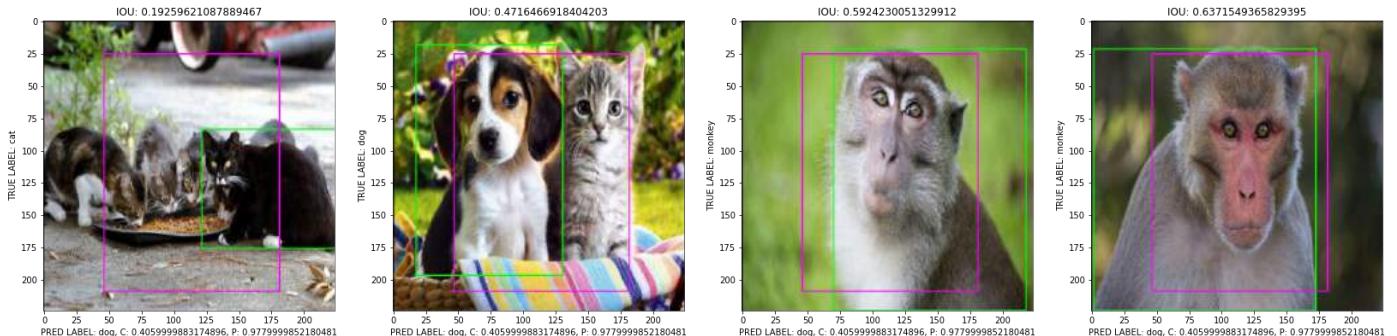
93/93 [=====] - ETA: 0s - loss: 0.4104 - iou_metric: 0.6730 - p_metric: 0.9900 - c_metric: 0.3410



93/93 [=====] - 17s 183ms/step - loss: 0.4104 - iou_metric: 0.6730 - p_metric: 0.9900 - c_metric: 0.3410 - val_loss: 0.3370 - val_iou_metric: 0.7263 - val_p_metric: 1.0000 - val_c_metric: 0.3527 - lr: 3.0000e-05

Epoch 12/100

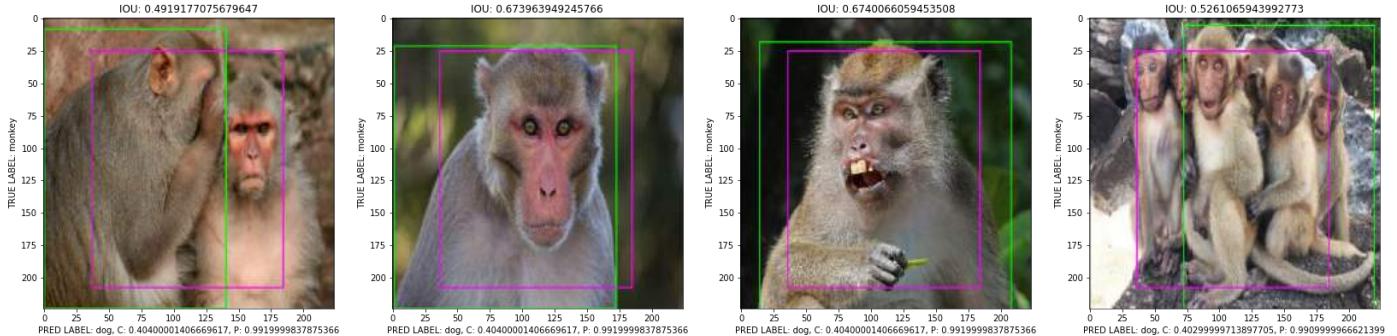
93/93 [=====] - ETA: 0s - loss: 0.4081 - iou_metric: 0.6761 - p_metric: 0.9909 - c_metric: 0.3400



93/93 [=====] - 17s 183ms/step - loss: 0.4081 - iou_metric: 0.6761 - p_metric: 0.9909 - c_metric: 0.3400 - val_loss: 0.3232 - val_iou_metric: 0.7290 - val_p_metric: 1.0000 - val_c_metric: 0.3467 - lr: 3.0000e-05

Epoch 13/100

93/93 [=====] - ETA: 0s - loss: 0.4183 - iou_metric: 0.6784 - p_metric: 0.9912 - c_metric: 0.3385



93/93 [=====] - 17s 183ms/step - loss: 0.4183 - iou_metric: 0.6784 - p_metric: 0.9912 - c_metric: 0.3385 - val_loss: 0.3209 - val_iou_metric: 0.7320 - val_p_metric: 1.0000 - val_c_metric: 0.3415 - lr: 3.0000e-05

Epoch 14/100

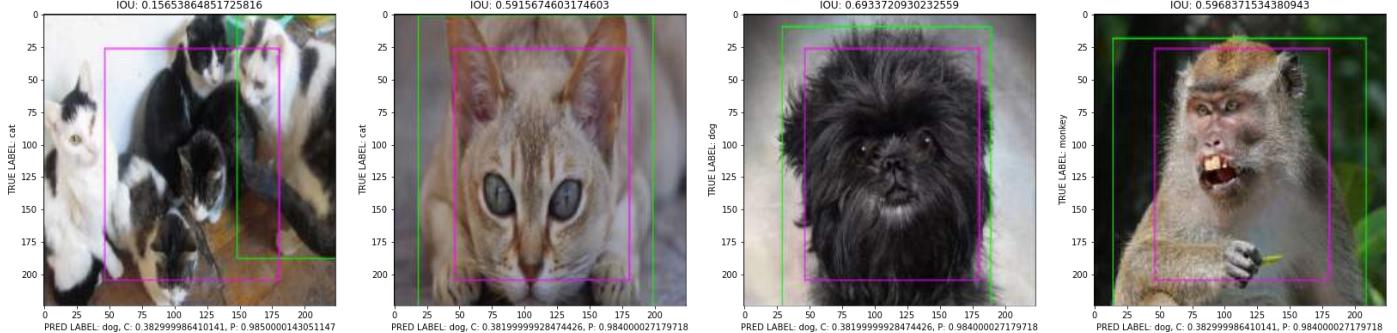
93/93 [=====] - ETA: 0s - loss: 0.3899 - iou_metric: 0.6813 - p_metric: 0.9919 - c_metric: 0.3390



93/93 [=====] - 17s 183ms/step - loss: 0.3899 - iou_metric: 0.6813 - p_metric: 0.9919 - c_metric: 0.3390 - val_loss: 0.3190 - val_iou_metric: 0.7343 - val_p_metric: 1.0000 - val_c_metric: 0.3371 - lr: 9.0000e-06

Epoch 15/100

93/93 [=====] - ETA: 0s - loss: 0.3925 - iou_metric: 0.6837 - p_metric: 0.9921 - c_metric: 0.3418

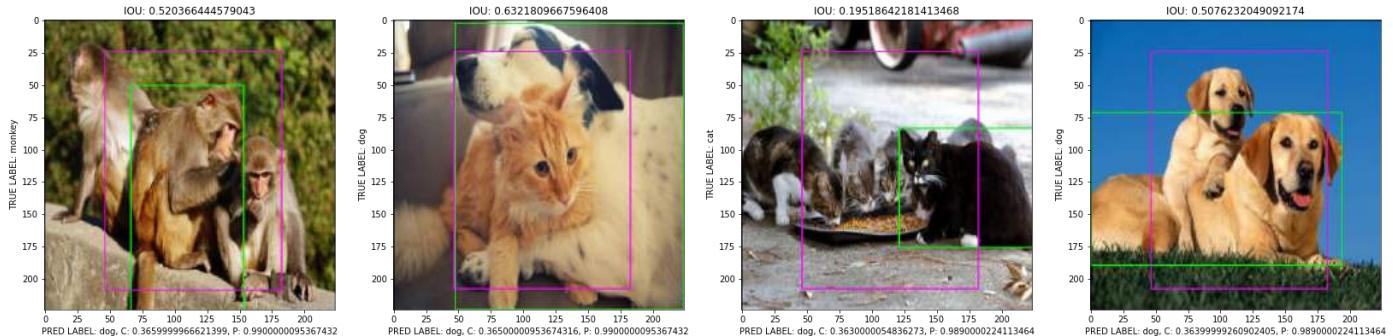


93/93 [=====] - 17s 185ms/step - loss: 0.3925 - iou_metric: 0.6837 - p_metric: 0.9921 - c_metric: 0.3418 - val_loss: 0.3205 - val_iou_metric: 0.7355 -

val_p_metric: 1.0000 - val_c_metric: 0.3333 - lr: 9.0000e-06

Epoch 16/100

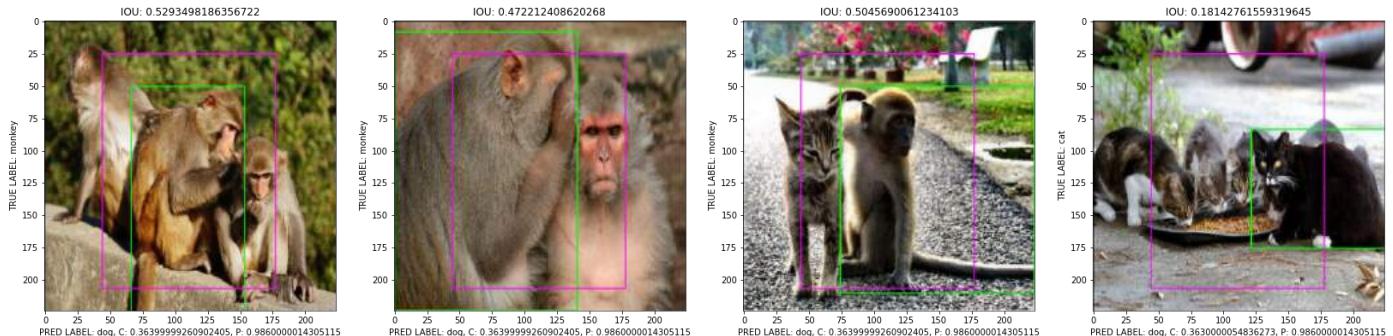
93/93 [=====] - ETA: 0s - loss: 0.4046 - iou_metric: 0.6857 - p_metric: 0.9926 - c_metric: 0.3395



93/93 [=====] - 18s 188ms/step - loss: 0.4046 - iou_metric: 0.6857 - p_metric: 0.9926 - c_metric: 0.3395 - val_loss: 0.3190 - val_iou_metric: 0.7370 - val_p_metric: 1.0000 - val_c_metric: 0.3300 - lr: 9.0000e-06

Epoch 17/100

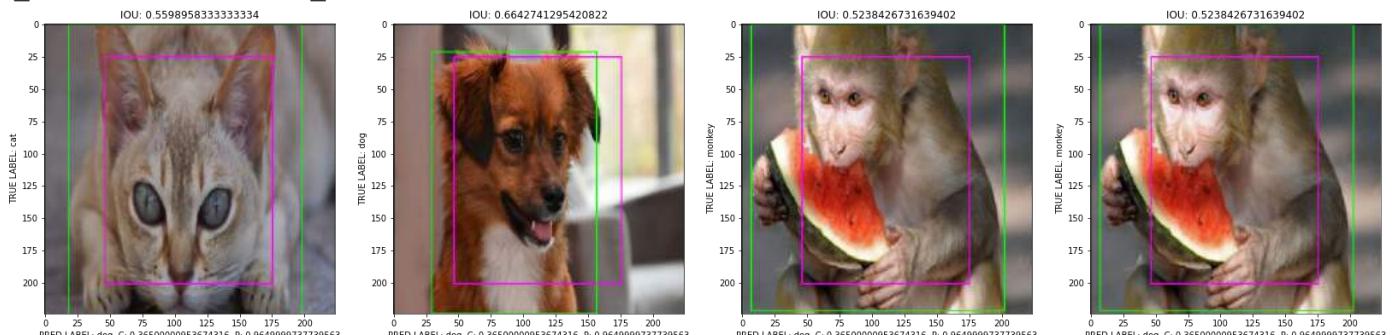
93/93 [=====] - ETA: 0s - loss: 0.4014 - iou_metric: 0.6874 - p_metric: 0.9929 - c_metric: 0.3394



93/93 [=====] - 18s 193ms/step - loss: 0.4014 - iou_metric: 0.6874 - p_metric: 0.9929 - c_metric: 0.3394 - val_loss: 0.3183 - val_iou_metric: 0.7381 - val_p_metric: 1.0000 - val_c_metric: 0.3271 - lr: 9.0000e-06

Epoch 18/100

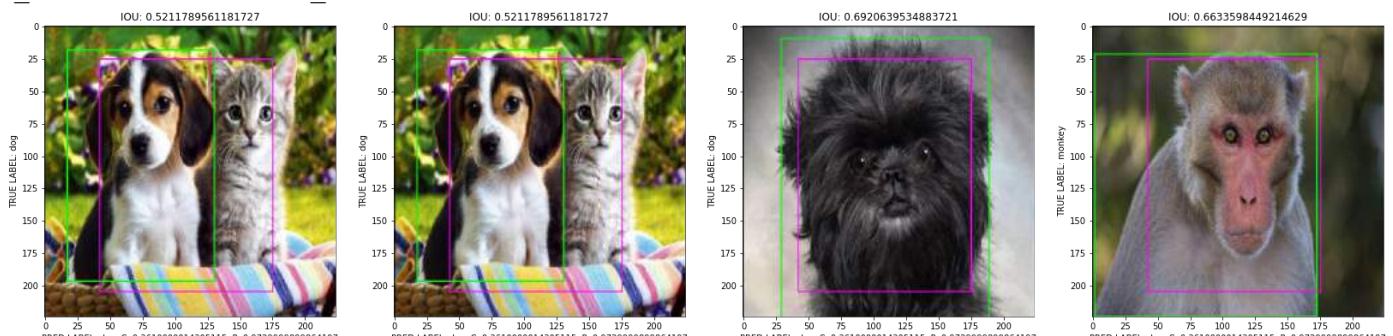
93/93 [=====] - ETA: 0s - loss: 0.3957 - iou_metric: 0.6890 - p_metric: 0.9933 - c_metric: 0.3388



93/93 [=====] - 18s 195ms/step - loss: 0.3957 - iou_metric: 0.6890 - p_metric: 0.9933 - c_metric: 0.3388 - val_loss: 0.3230 - val_iou_metric: 0.7386 - val_p_metric: 1.0000 - val_c_metric: 0.3244 - lr: 9.0000e-06

Epoch 19/100

93/93 [=====] - ETA: 0s - loss: 0.4010 - iou_metric: 0.6907 - p_metric: 0.9937 - c_metric: 0.3372



```
93/93 [=====] - 18s 189ms/step - loss: 0.4010 - iou_metric: 0.6  
907 - p_metric: 0.9937 - c_metric: 0.3372 - val_loss: 0.3190 - val_iou_metric: 0.7394 -  
val_p_metric: 1.0000 - val_c_metric: 0.3221 - lr: 9.0000e-06
```

Epoch 20/100

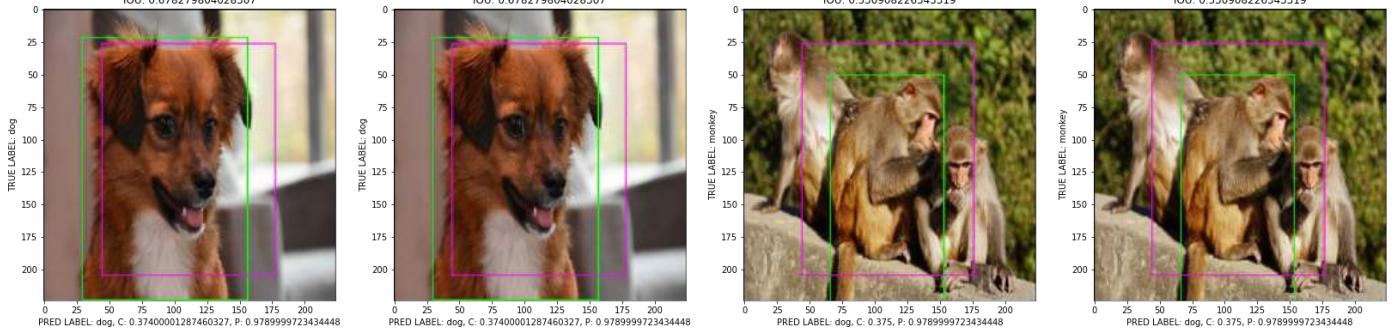
```
93/93 [=====] - ETA: 0s - loss: 0.3912 - iou_metric: 0.6922 - p  
_metric: 0.9940 - c_metric: 0.3362
```



```
93/93 [=====] - 18s 190ms/step - loss: 0.3912 - iou_metric: 0.6  
922 - p_metric: 0.9940 - c_metric: 0.3362 - val_loss: 0.3225 - val_iou_metric: 0.7400 -  
val_p_metric: 1.0000 - val_c_metric: 0.3200 - lr: 9.0000e-06
```

Epoch 21/100

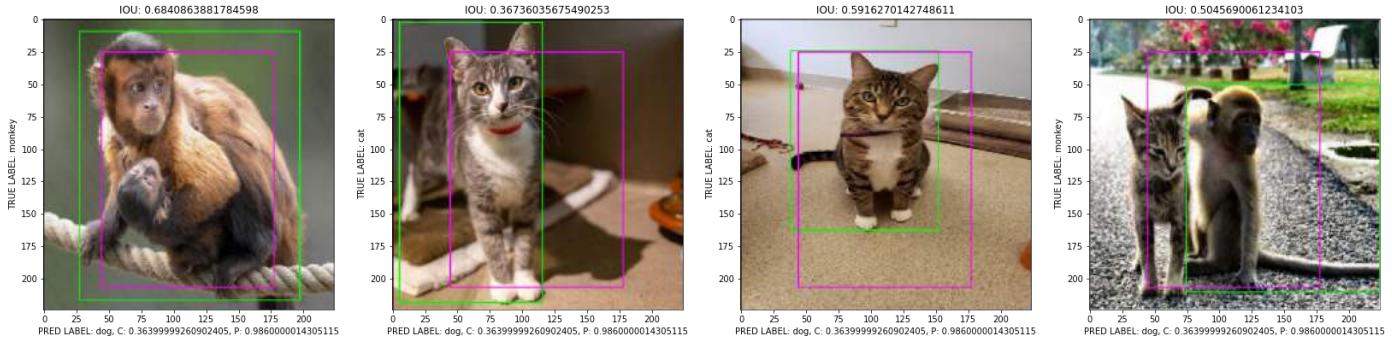
```
93/93 [=====] - ETA: 0s - loss: 0.3981 - iou_metric: 0.6932 - p  
_metric: 0.9942 - c_metric: 0.3367
```



```
93/93 [=====] - 18s 189ms/step - loss: 0.3981 - iou_metric: 0.6  
932 - p_metric: 0.9942 - c_metric: 0.3367 - val_loss: 0.3203 - val_iou_metric: 0.7406 -  
val_p_metric: 1.0000 - val_c_metric: 0.3181 - lr: 2.7000e-06
```

Epoch 22/100

```
93/93 [=====] - ETA: 0s - loss: 0.3864 - iou_metric: 0.6945 - p  
_metric: 0.9943 - c_metric: 0.3383
```



```
93/93 [=====] - 18s 193ms/step - loss: 0.3864 - iou_metric: 0.6  
945 - p_metric: 0.9943 - c_metric: 0.3383 - val_loss: 0.3205 - val_iou_metric: 0.7412 -  
val_p_metric: 1.0000 - val_c_metric: 0.3164 - lr: 2.7000e-06
```

SAVE AND EVALUATE MODEL ON TESTING DATA

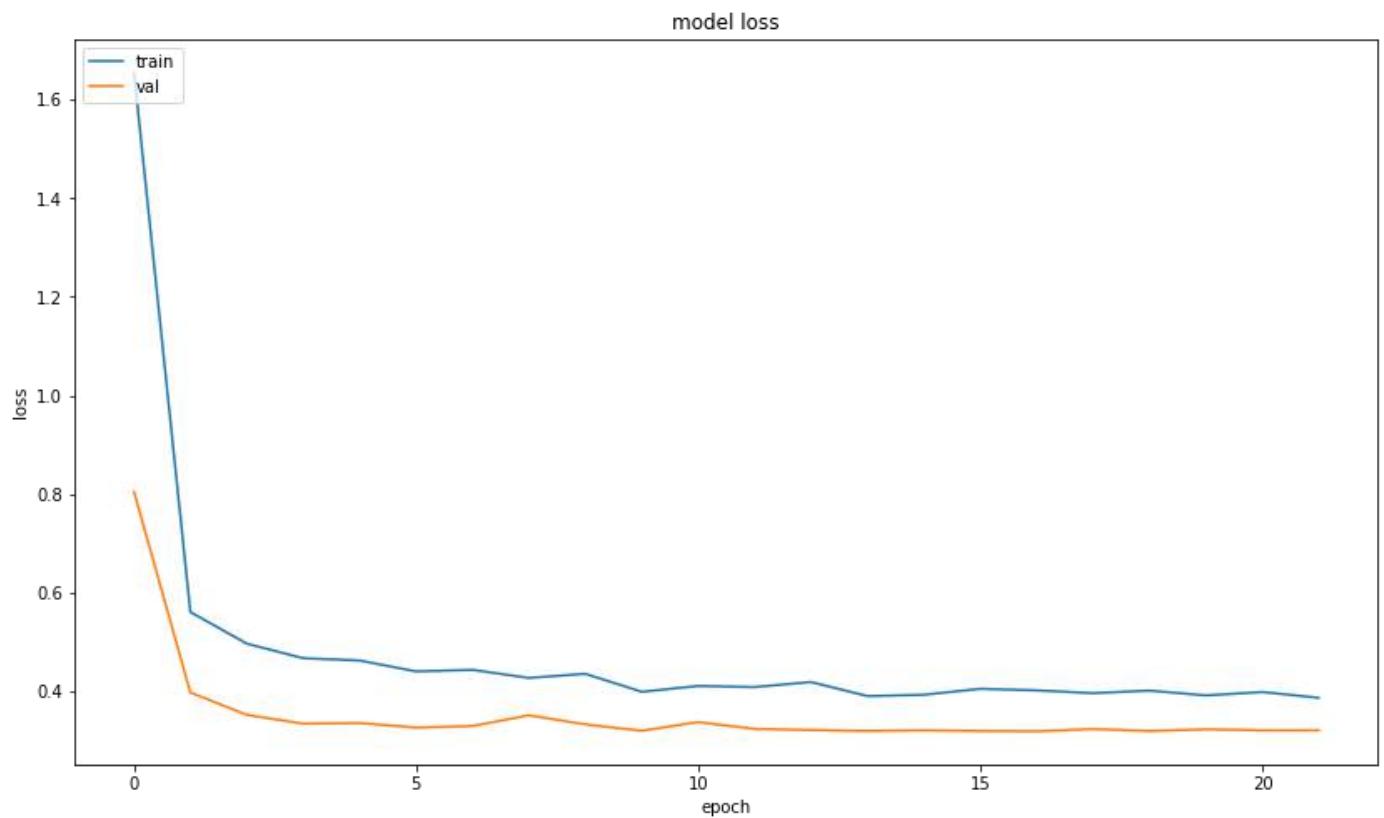
```
In [ ]: model2.save('model2_full.tf')  
model2.evaluate(test_gen)
```

```
5/5 [=====] - 1s 91ms/step - loss: 0.3059 - iou_metric: 0.7421  
- p_metric: 1.0000 - c_metric: 0.3235
```

```
Out[ ]: [0.3059229552745819, 0.7421384453773499, 1.0, 0.3234783709049225]
```

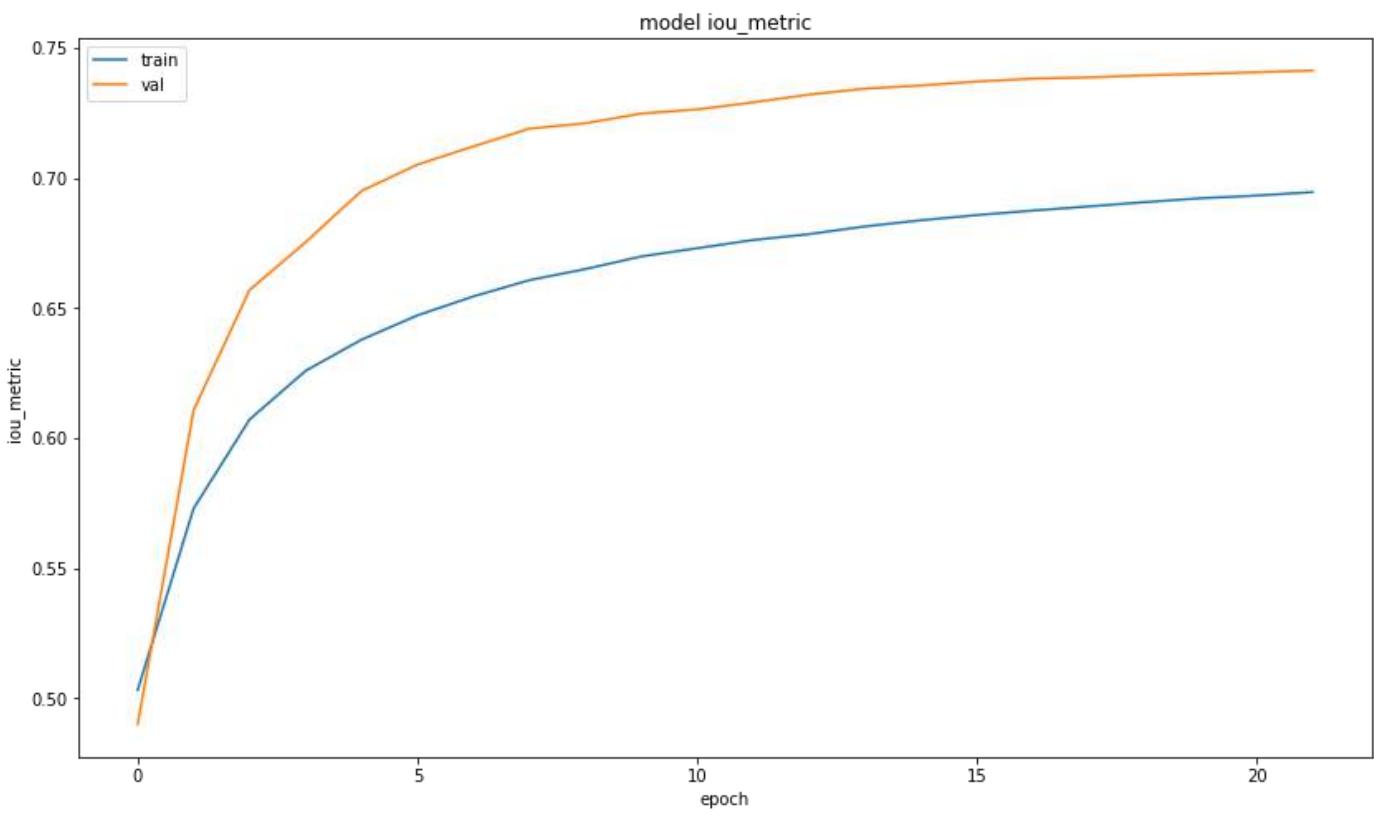
PLOT MODEL LOSS

```
In [ ]: plot_metric(fit2, 'loss')
```

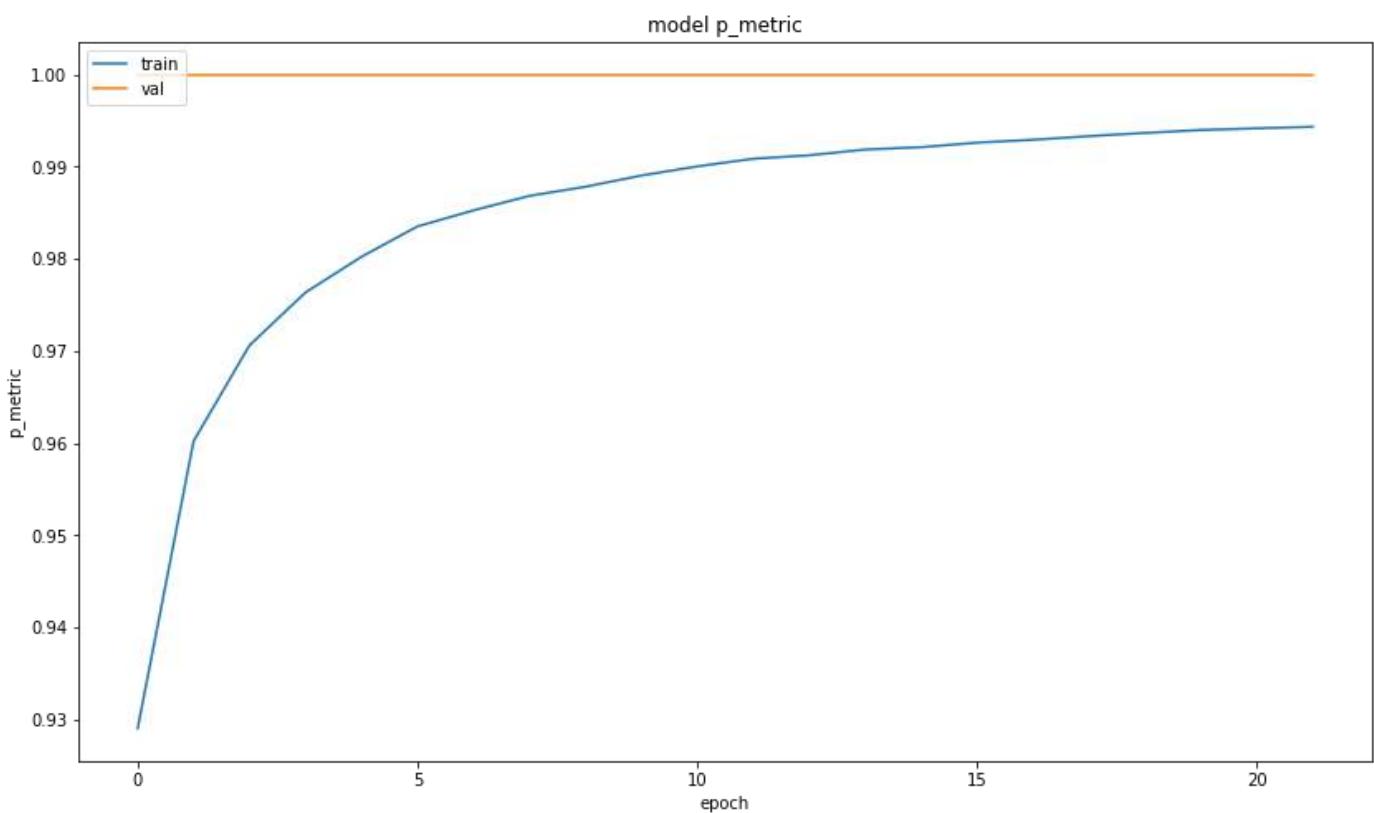


PLOT MODEL METRIC

```
In [ ]: plot_metric(fit2, 'iou_metric')
```



```
In [ ]: plot_metric(fit2, 'p_metric')
```



```
In [ ]: plot_metric(fit2, 'c_metric')
```

